INSTITUT TUTTE INSTITUTE

Clustering Data

A Guide for the Perplexed

Leland McInnes
John Healy

Tutte Institute
for Mathematics and Computing

leland.mcinnes@gmail.com
jchealy@gmail.com

INSTITUT TUTTE INSTITUTE

What is clustering?

INSTITUT TUTTE INSTITUTE

Find groups of data that are all similar

INSTITUT TUTTE INSTITUTE

What that means depends on your application

INSTITUT TUTTE INSTITUTE

Goals

- Partition your data
- Summarize your data
- Explore your data
- Embed in a vector space
- Find patterns in your data
- …

INSTITUT TUTTE INSTITUTE

Find groups of data that are all similar

INSTITUT TUTTE INSTITUTE

Easy in theory

Hard in practice

INSTITUT TUTTE INSTITUTE

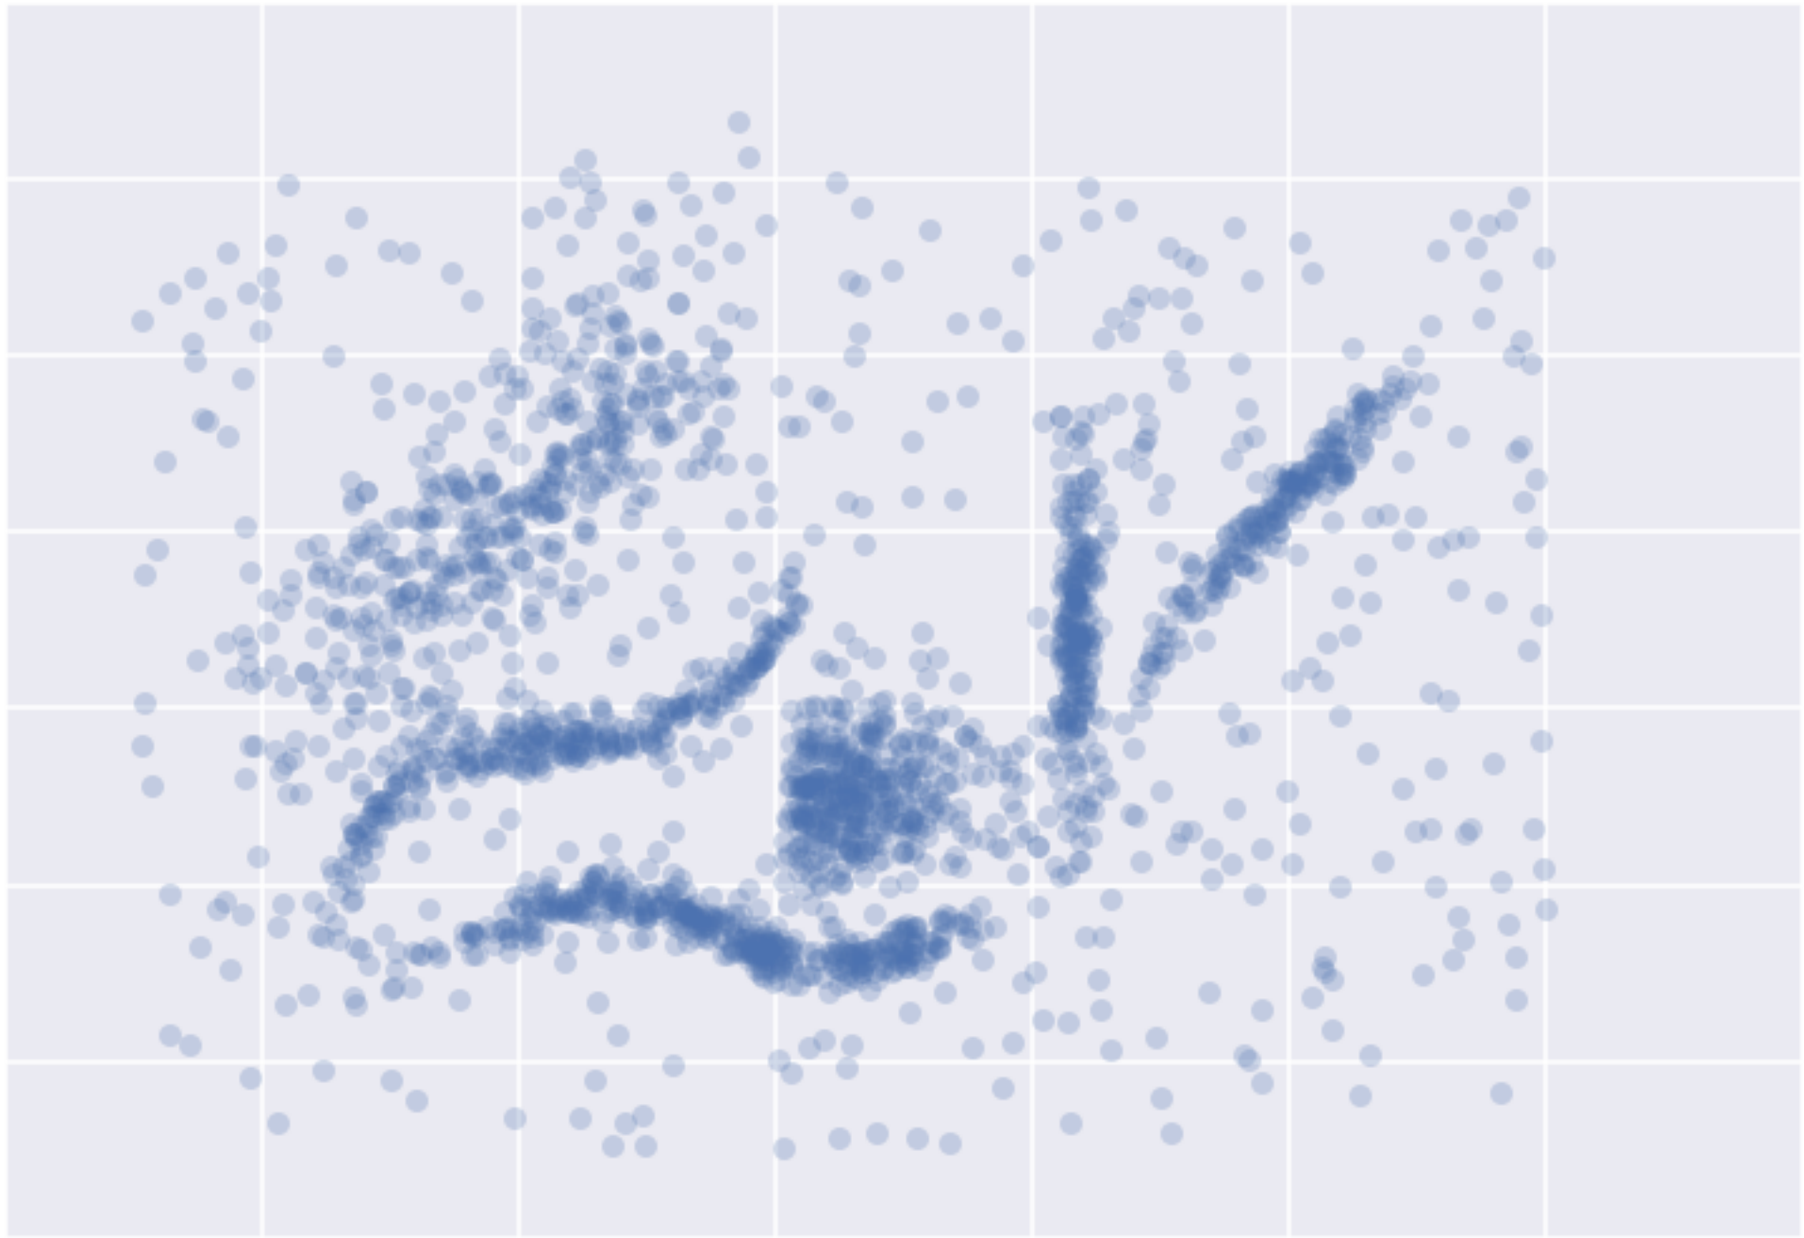

INSTITUT TUTTE INSTITUTE

Clusters found by KMeans

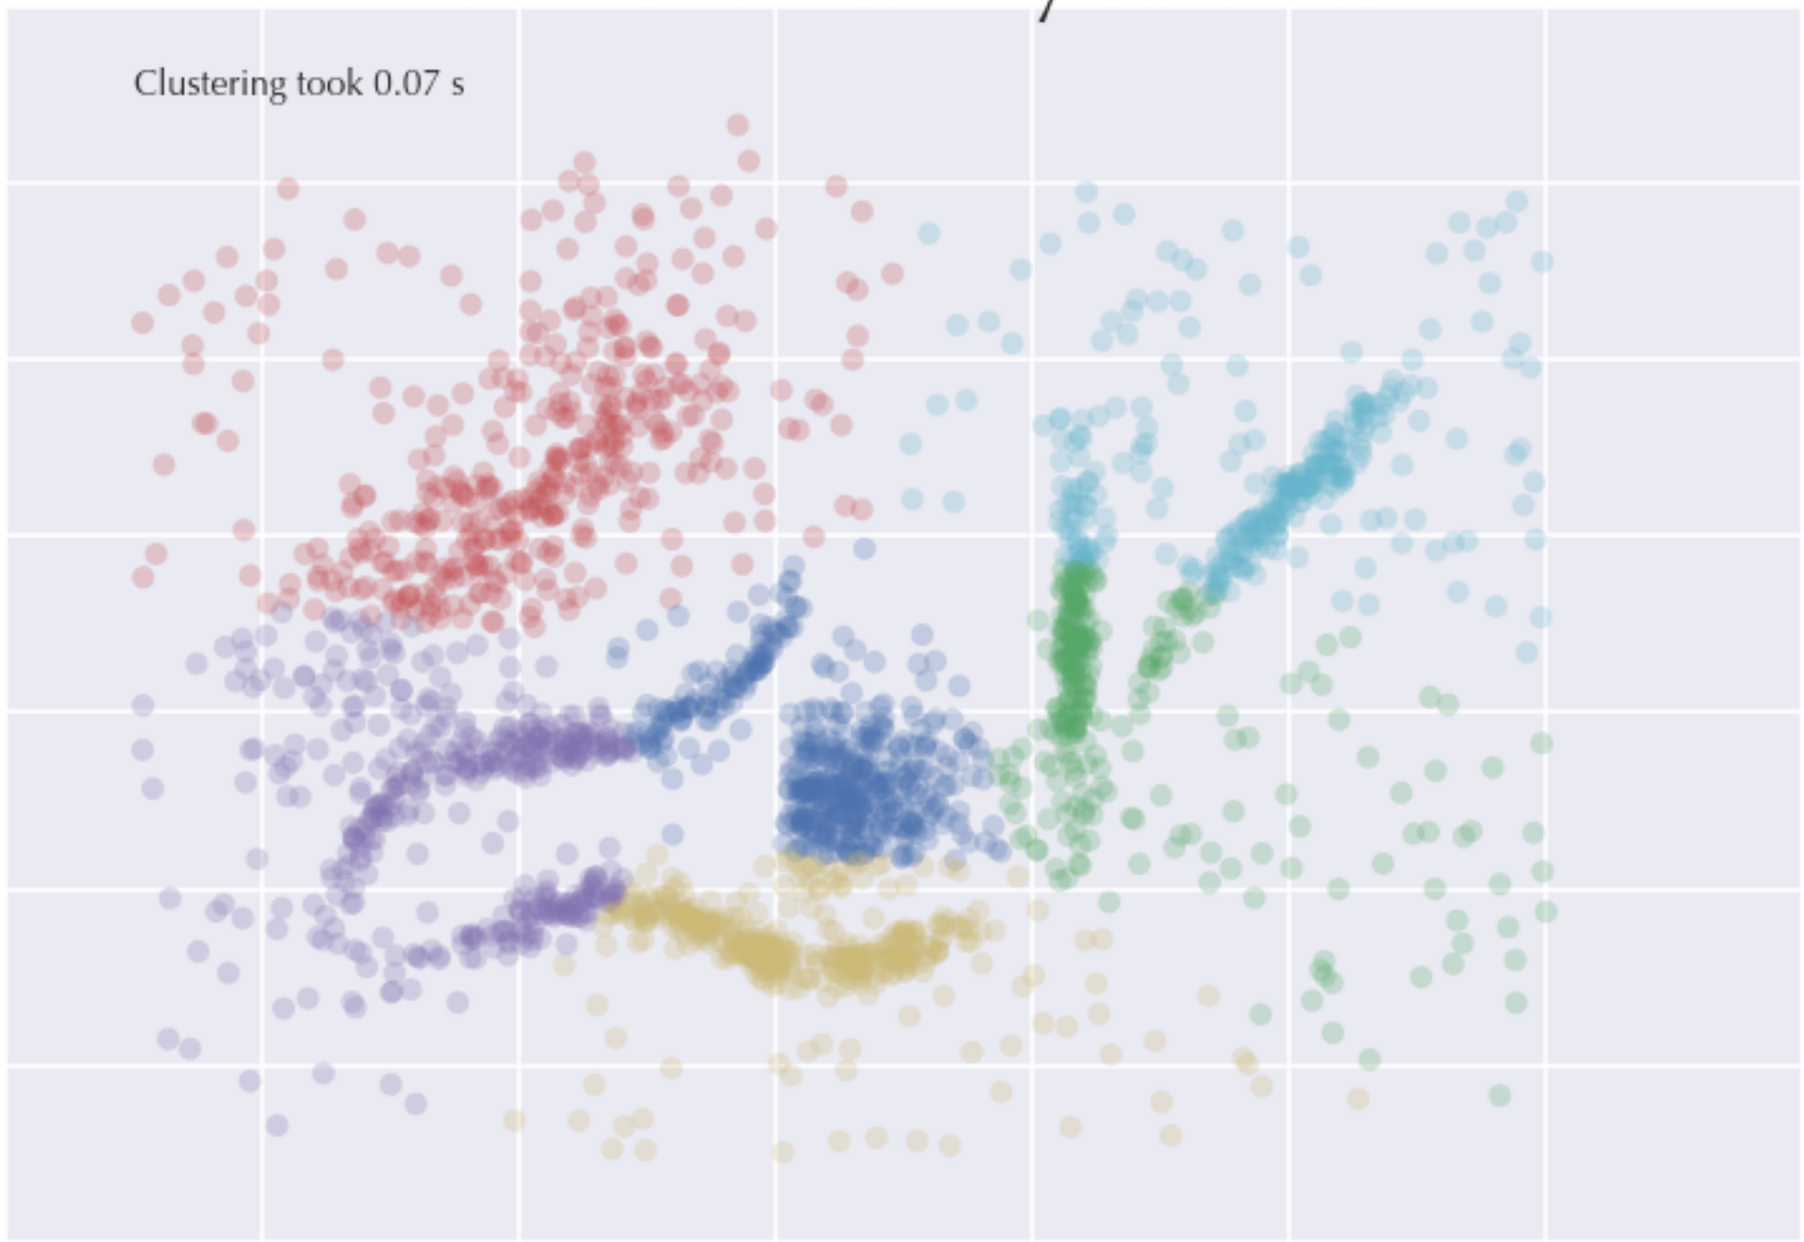

INSTITUT TUTTE INSTITUTE

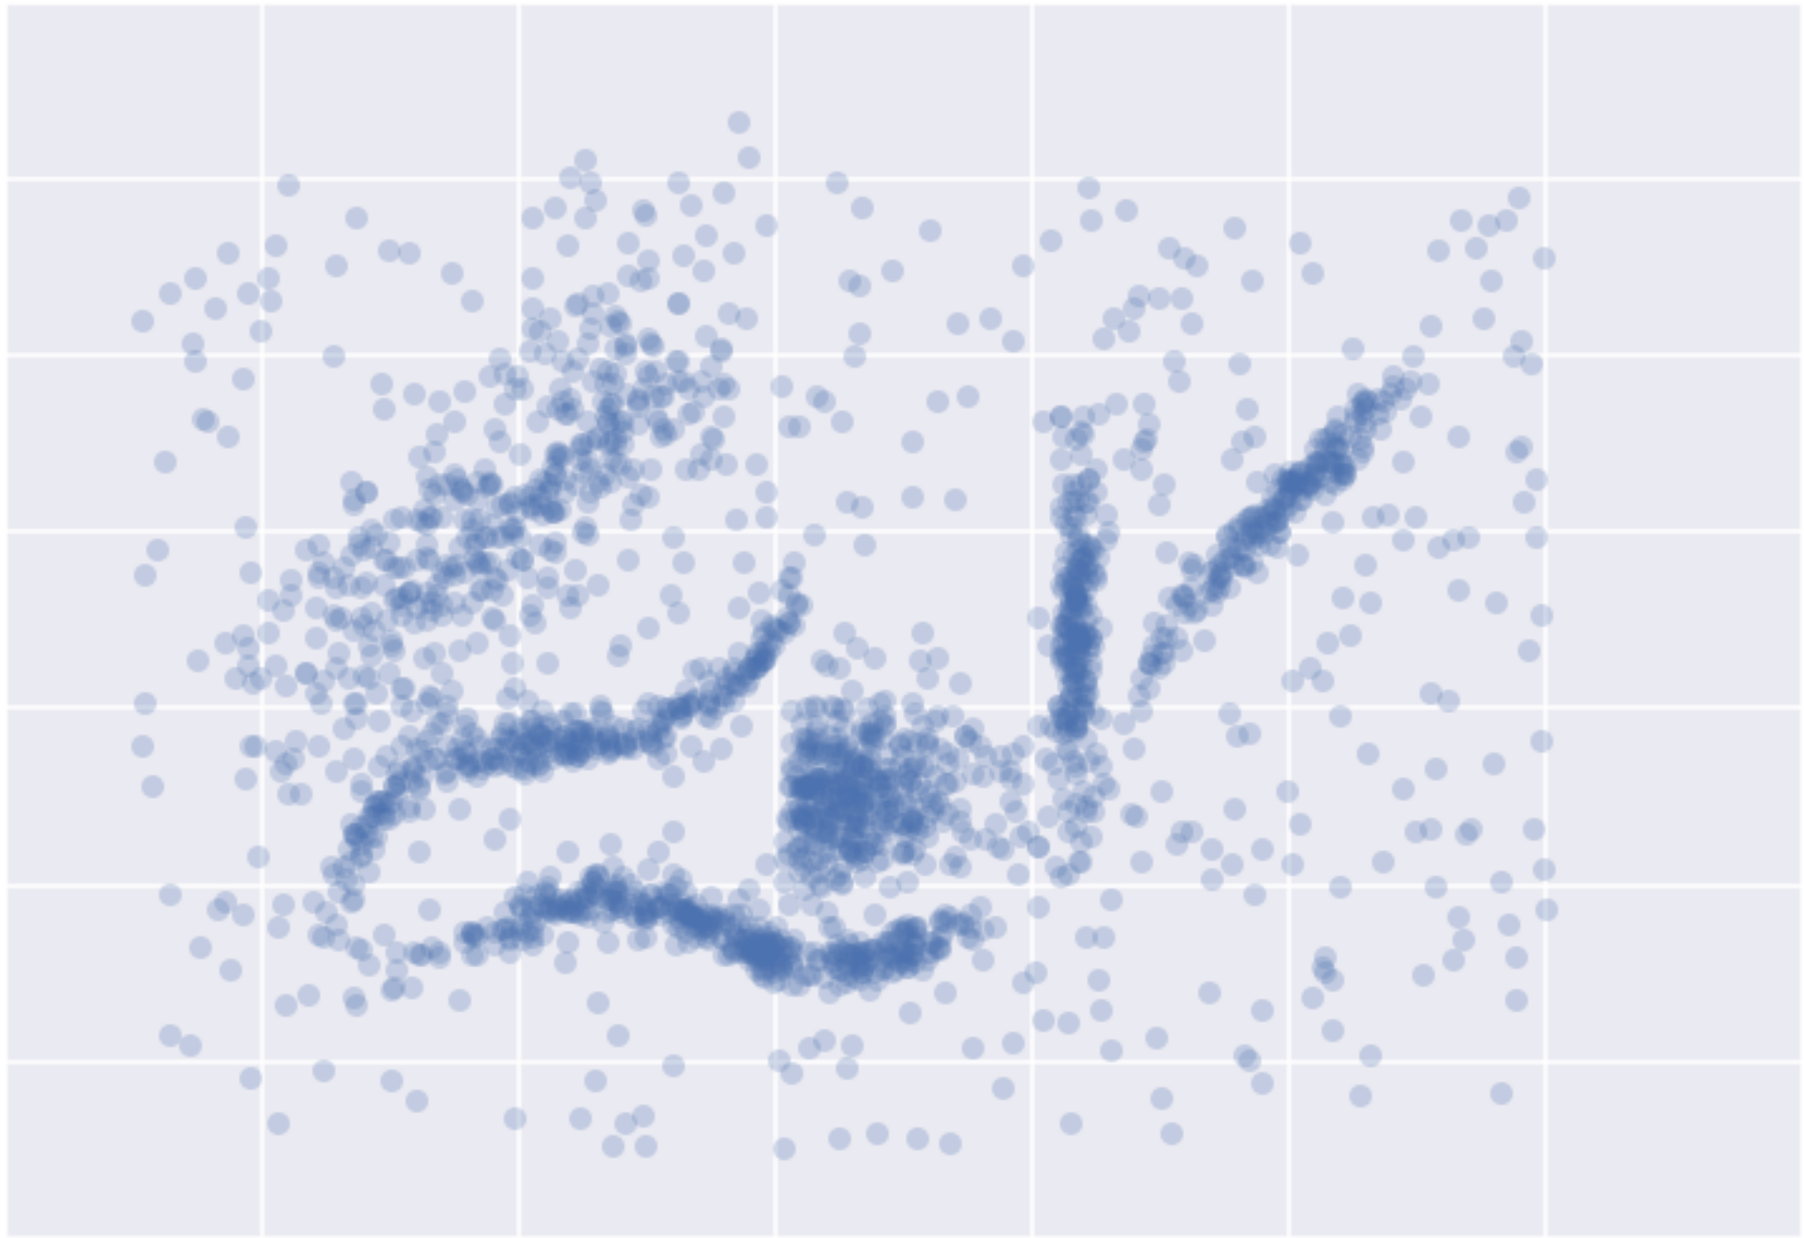

INSTITUT TUTTE INSTITUTE

Clusters found by AffinityPropagation

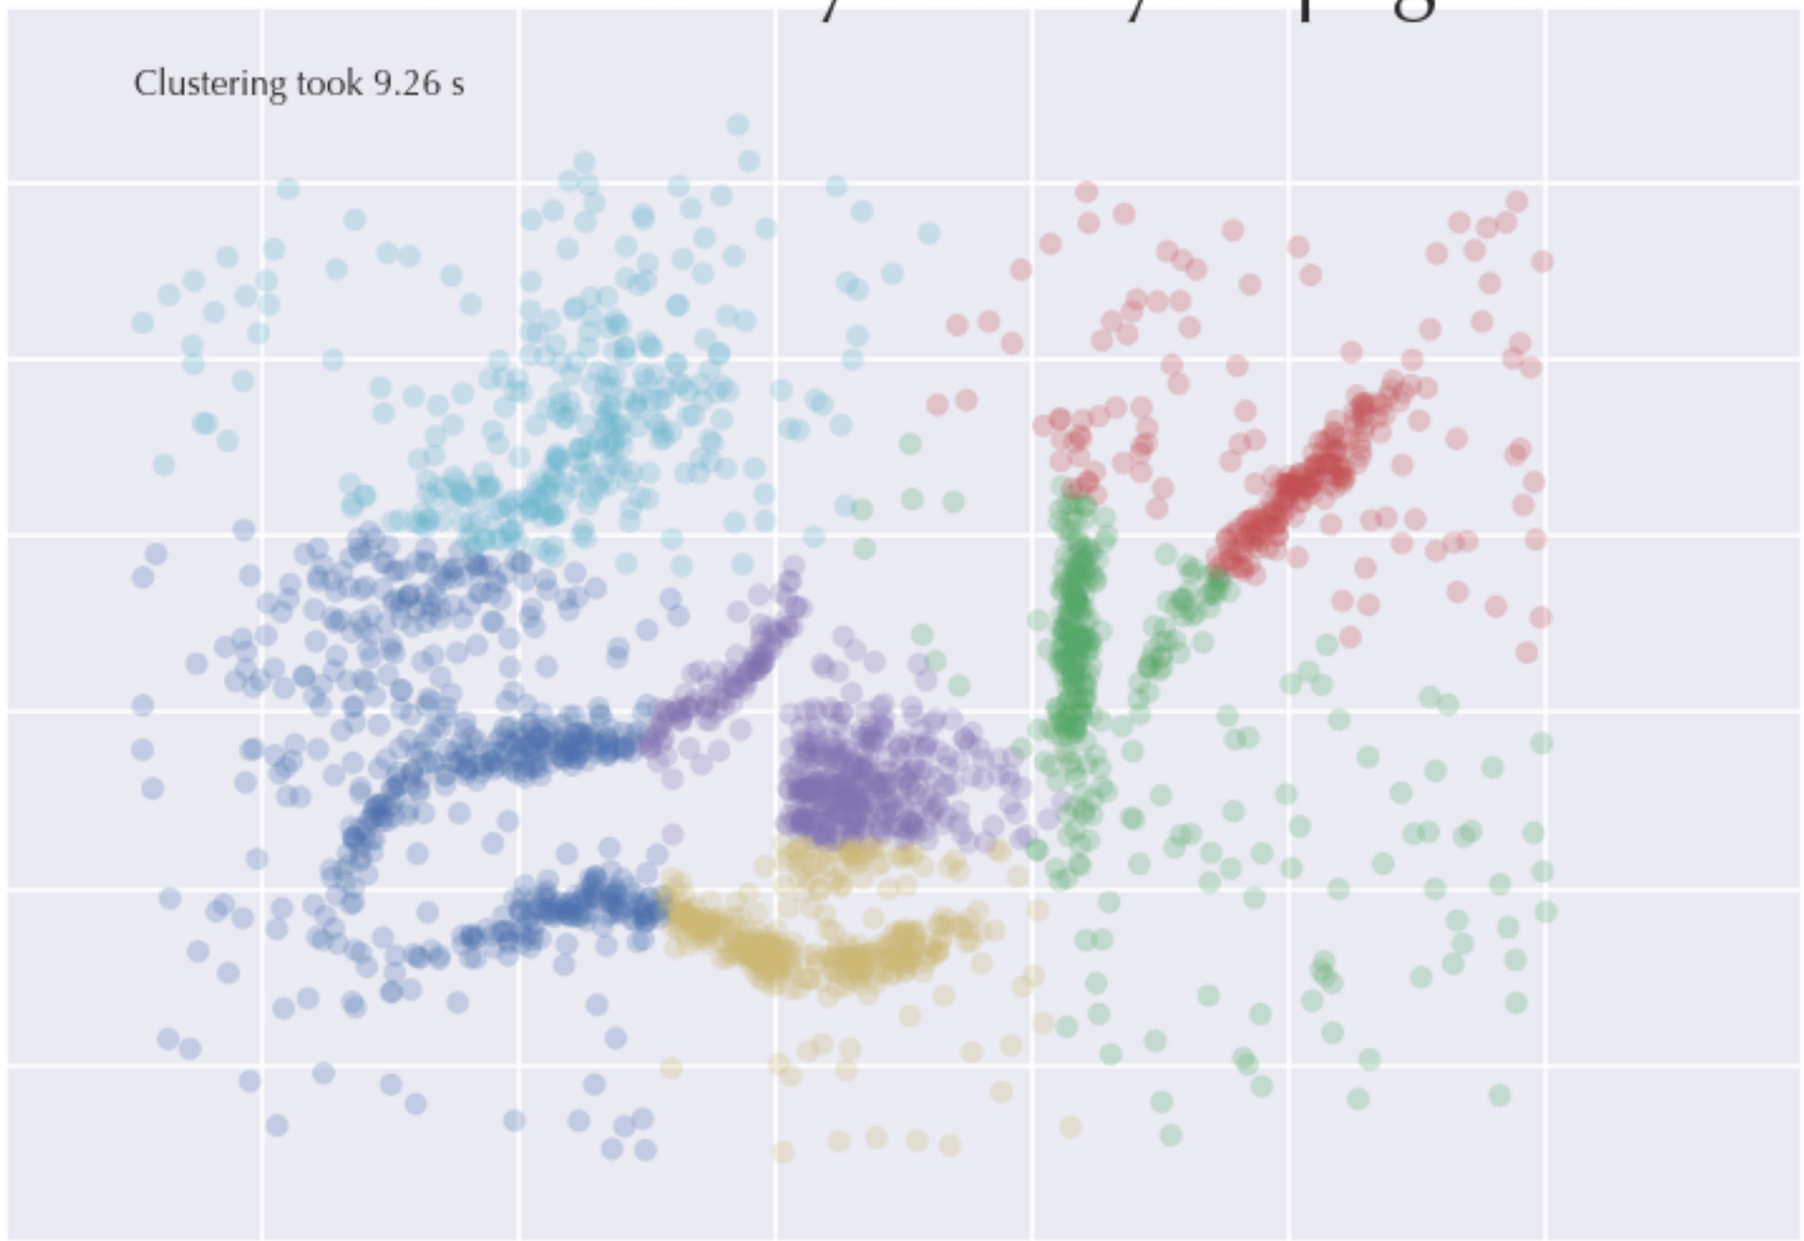

INSTITUT TUTTE INSTITUTE

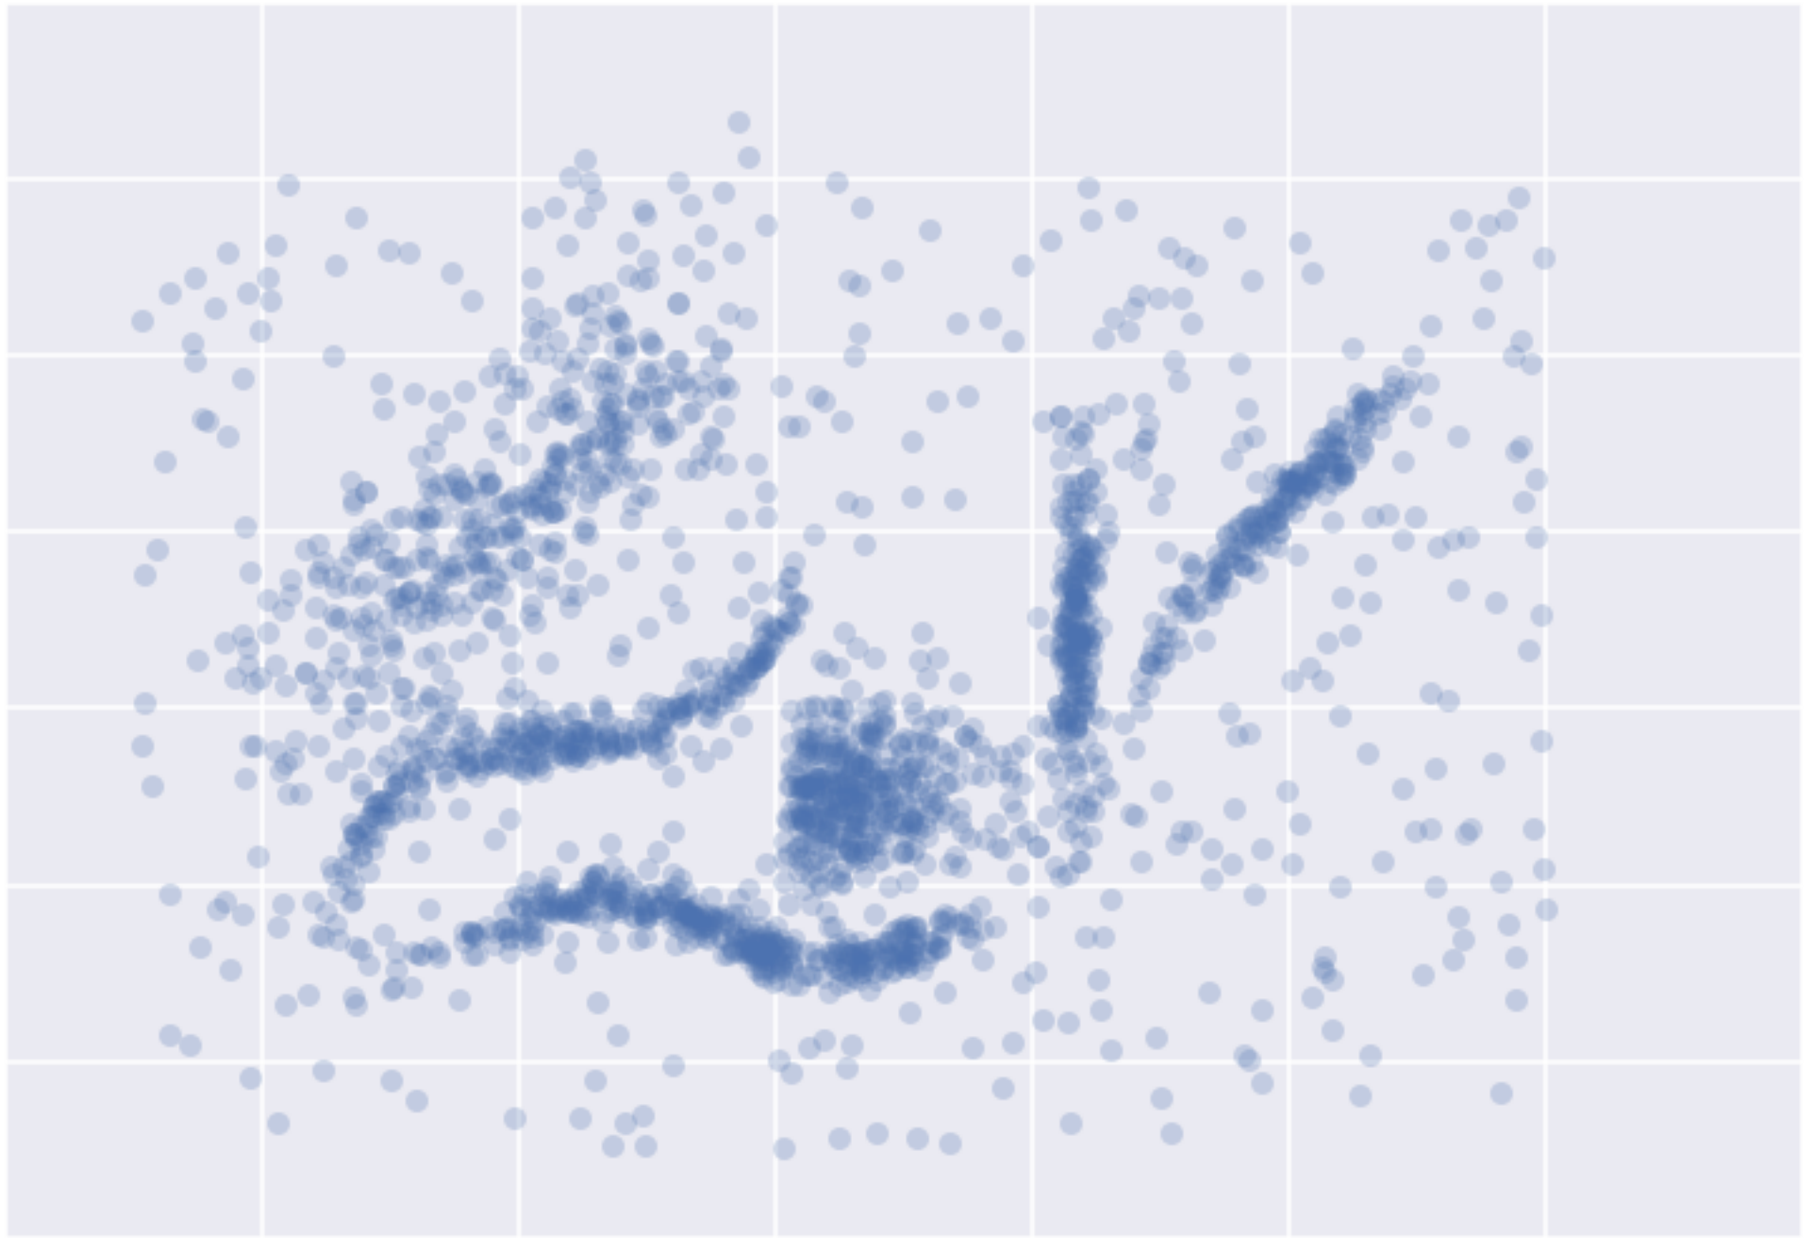

INSTITUT TUTTE INSTITUTE

Clusters found by MeanShift

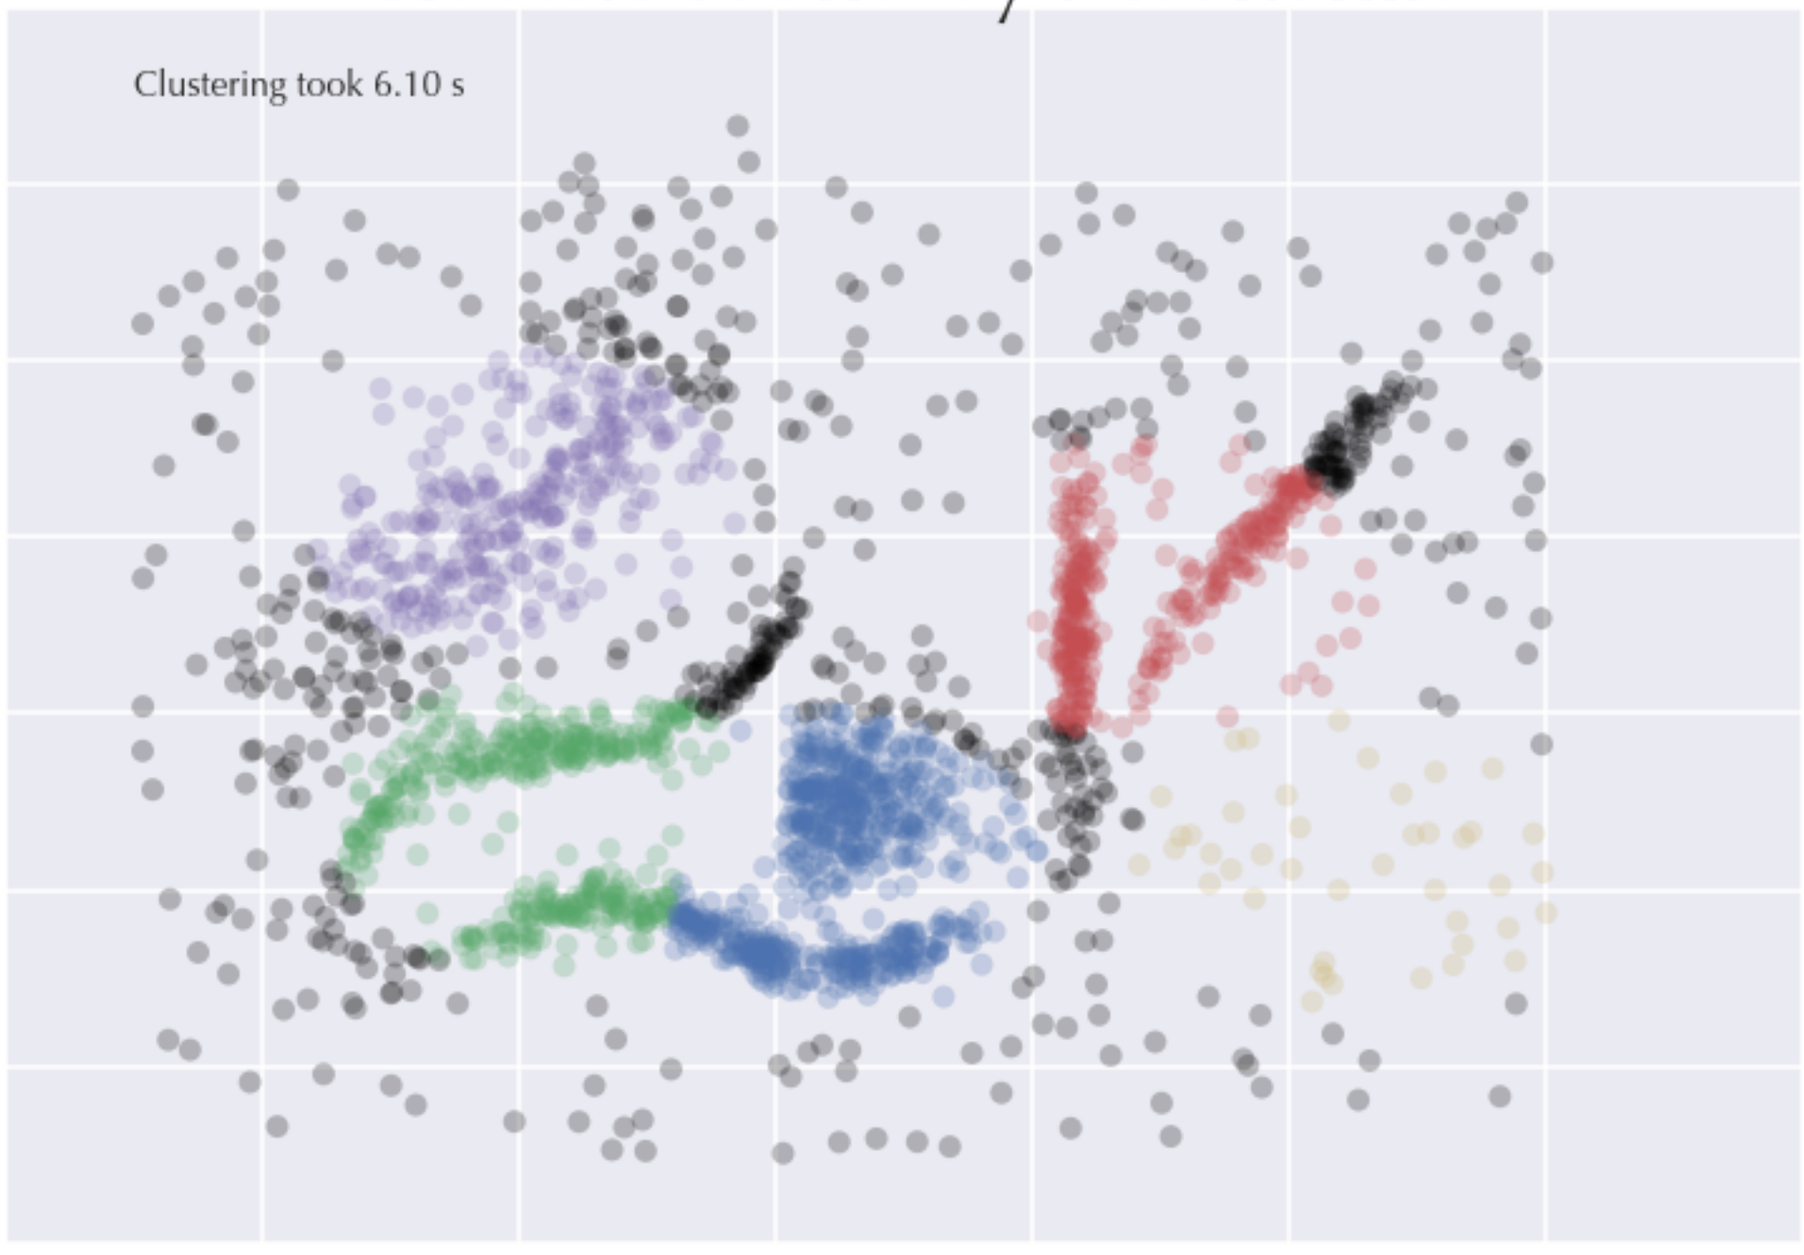

INSTITUT TUTTE INSTITUTE

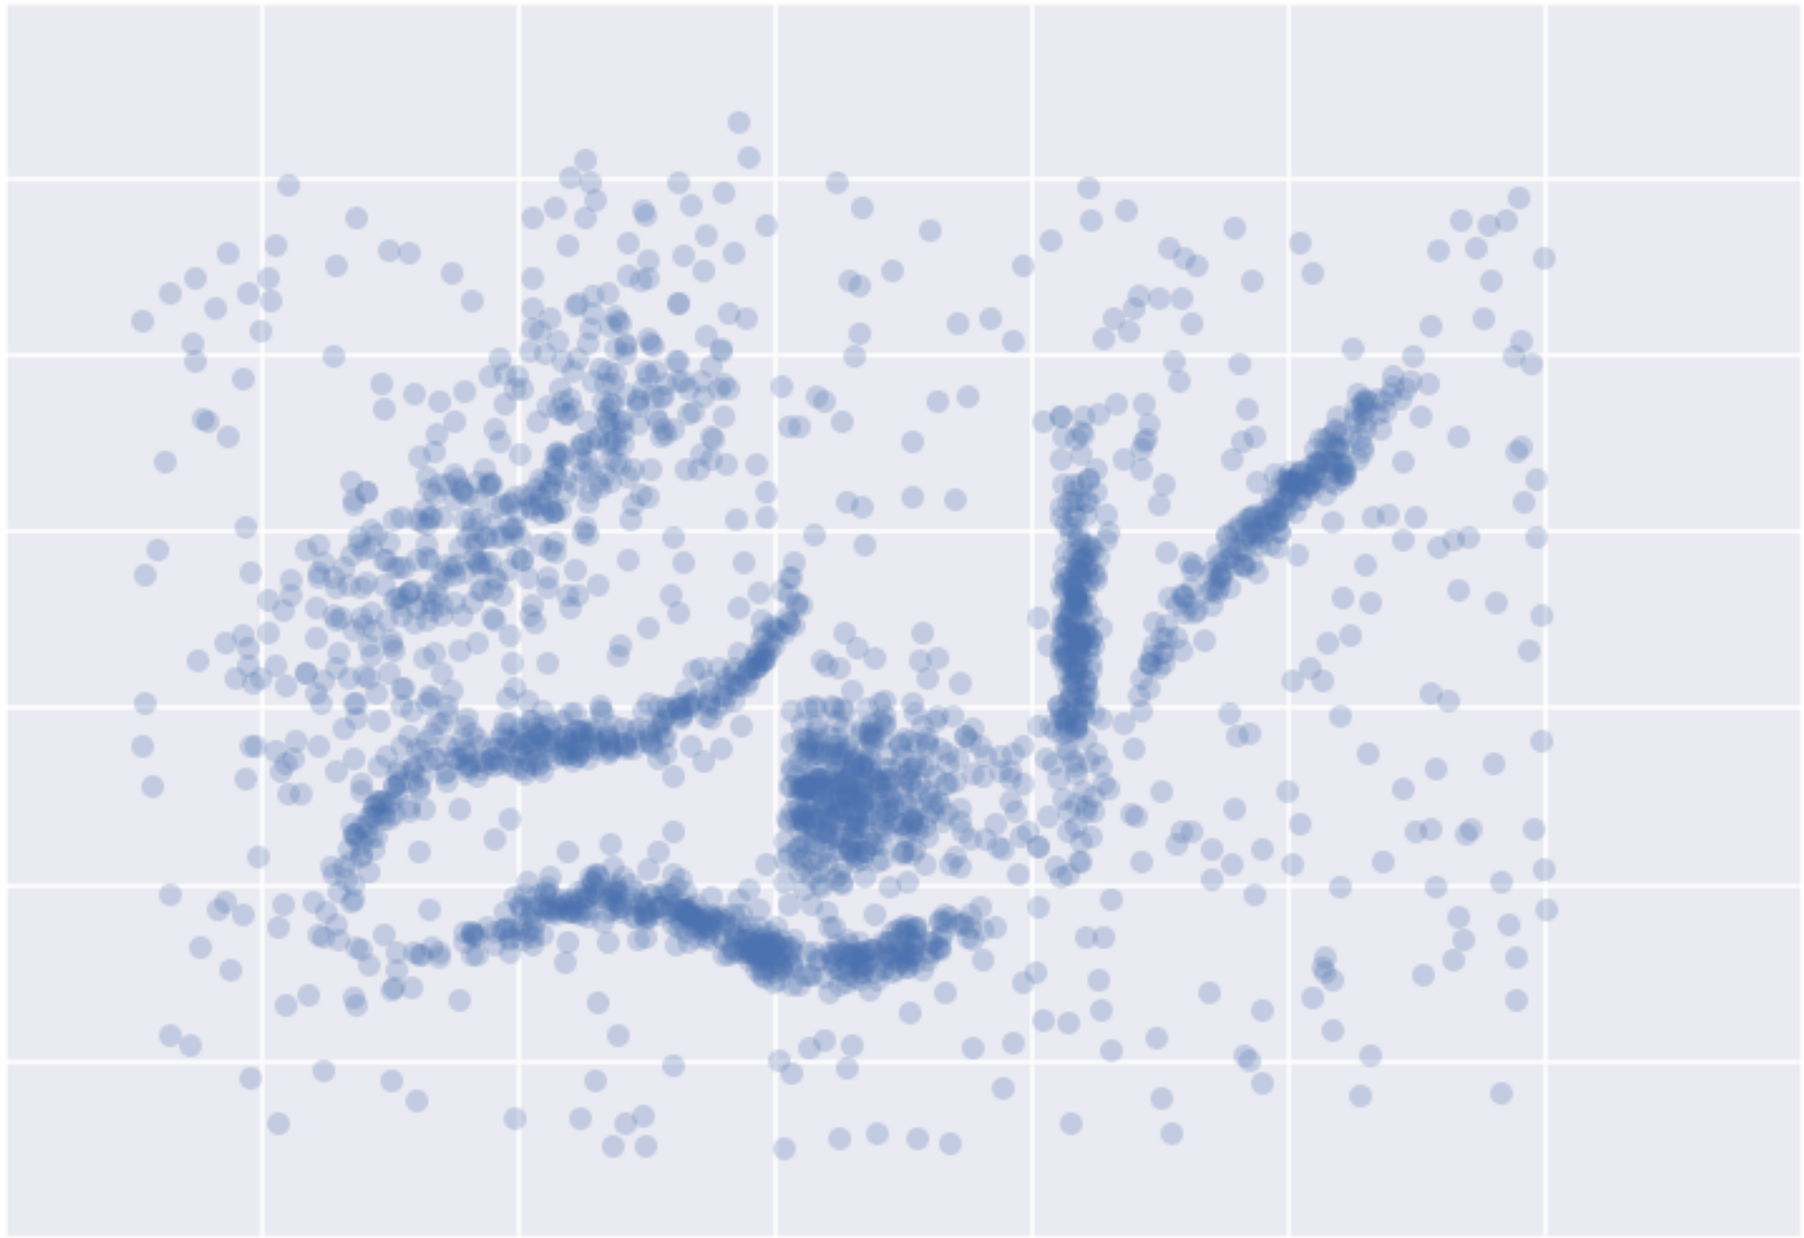

INSTITUT TUTTE INSTITUTE

Clusters found by SpectralClustering

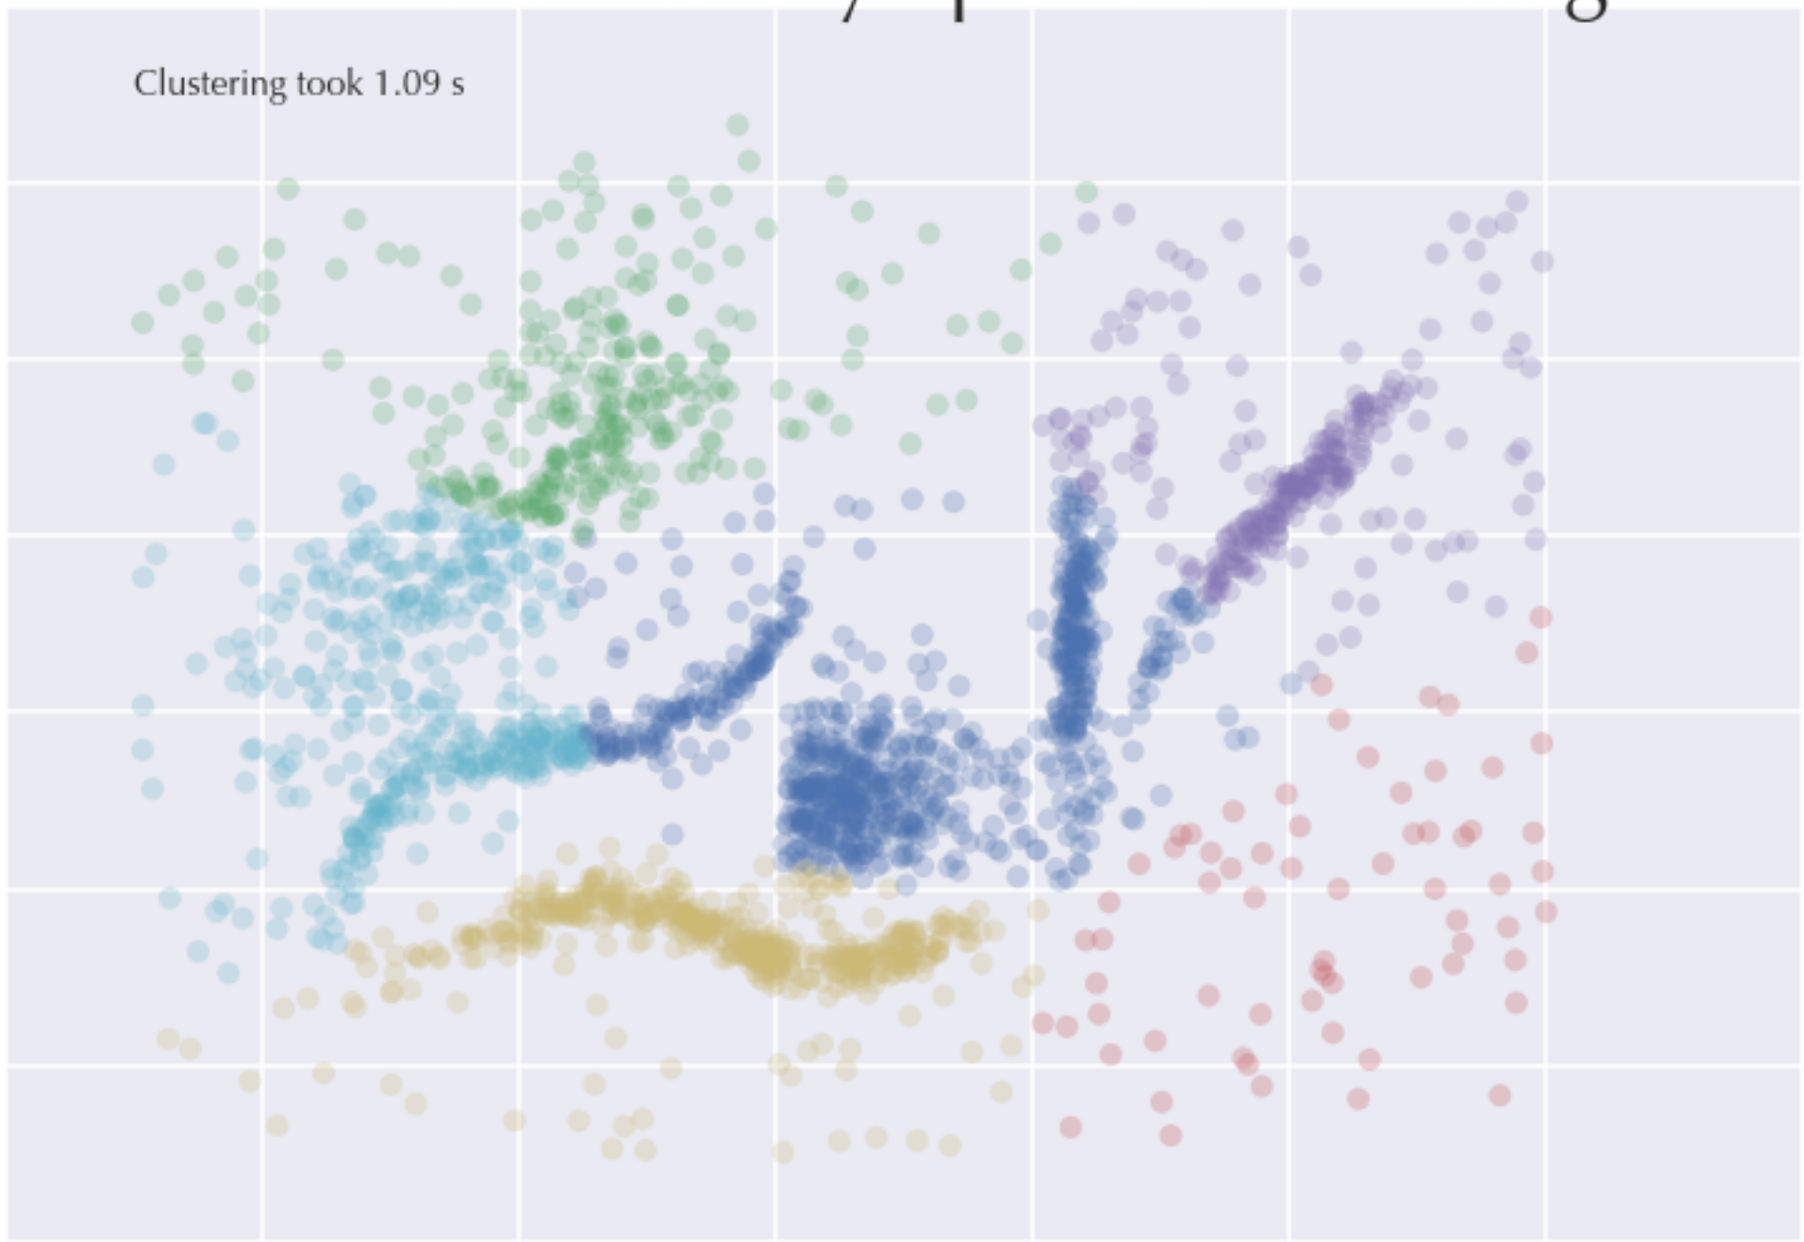

INSTITUT TUTTE INSTITUTE

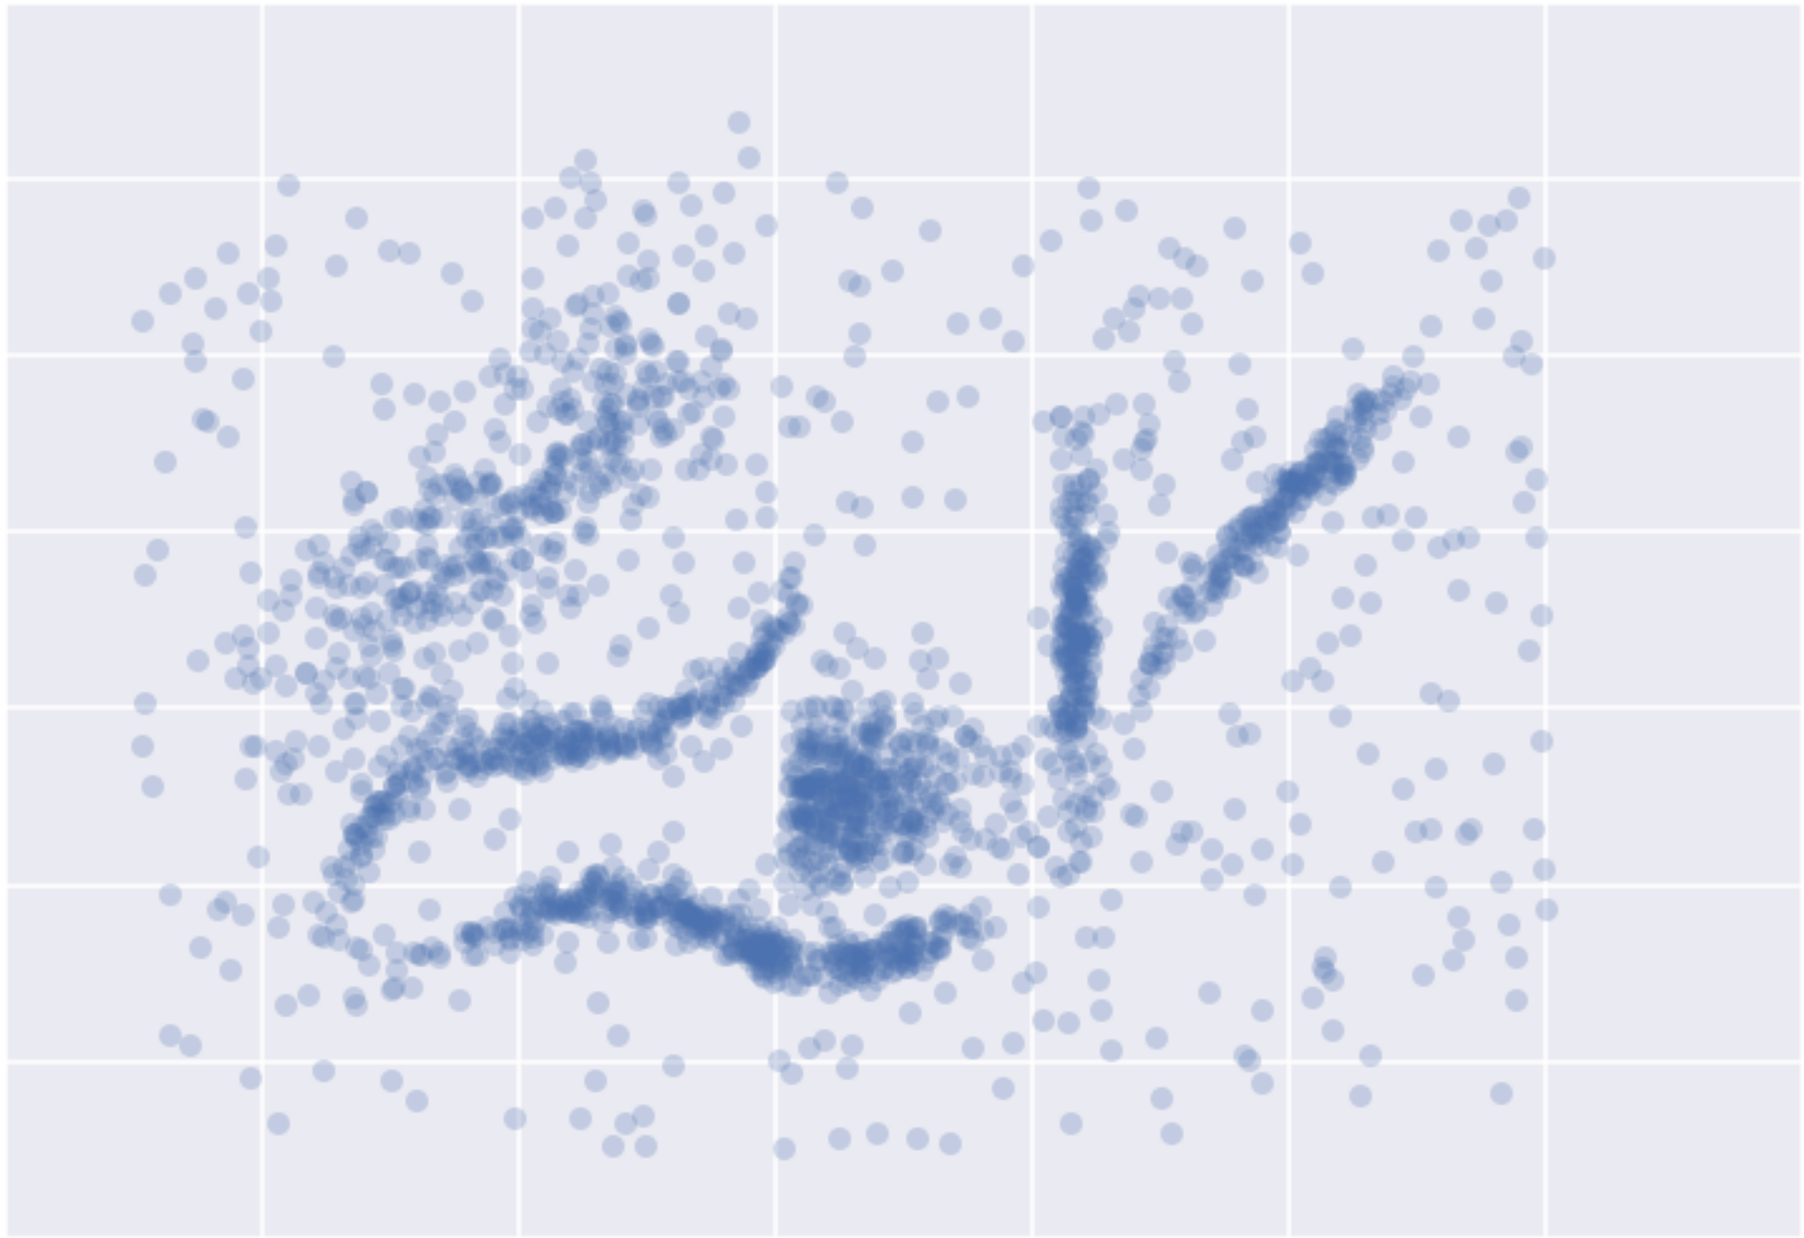

INSTITUT TUTTE INSTITUTE

Clusters found by Birch

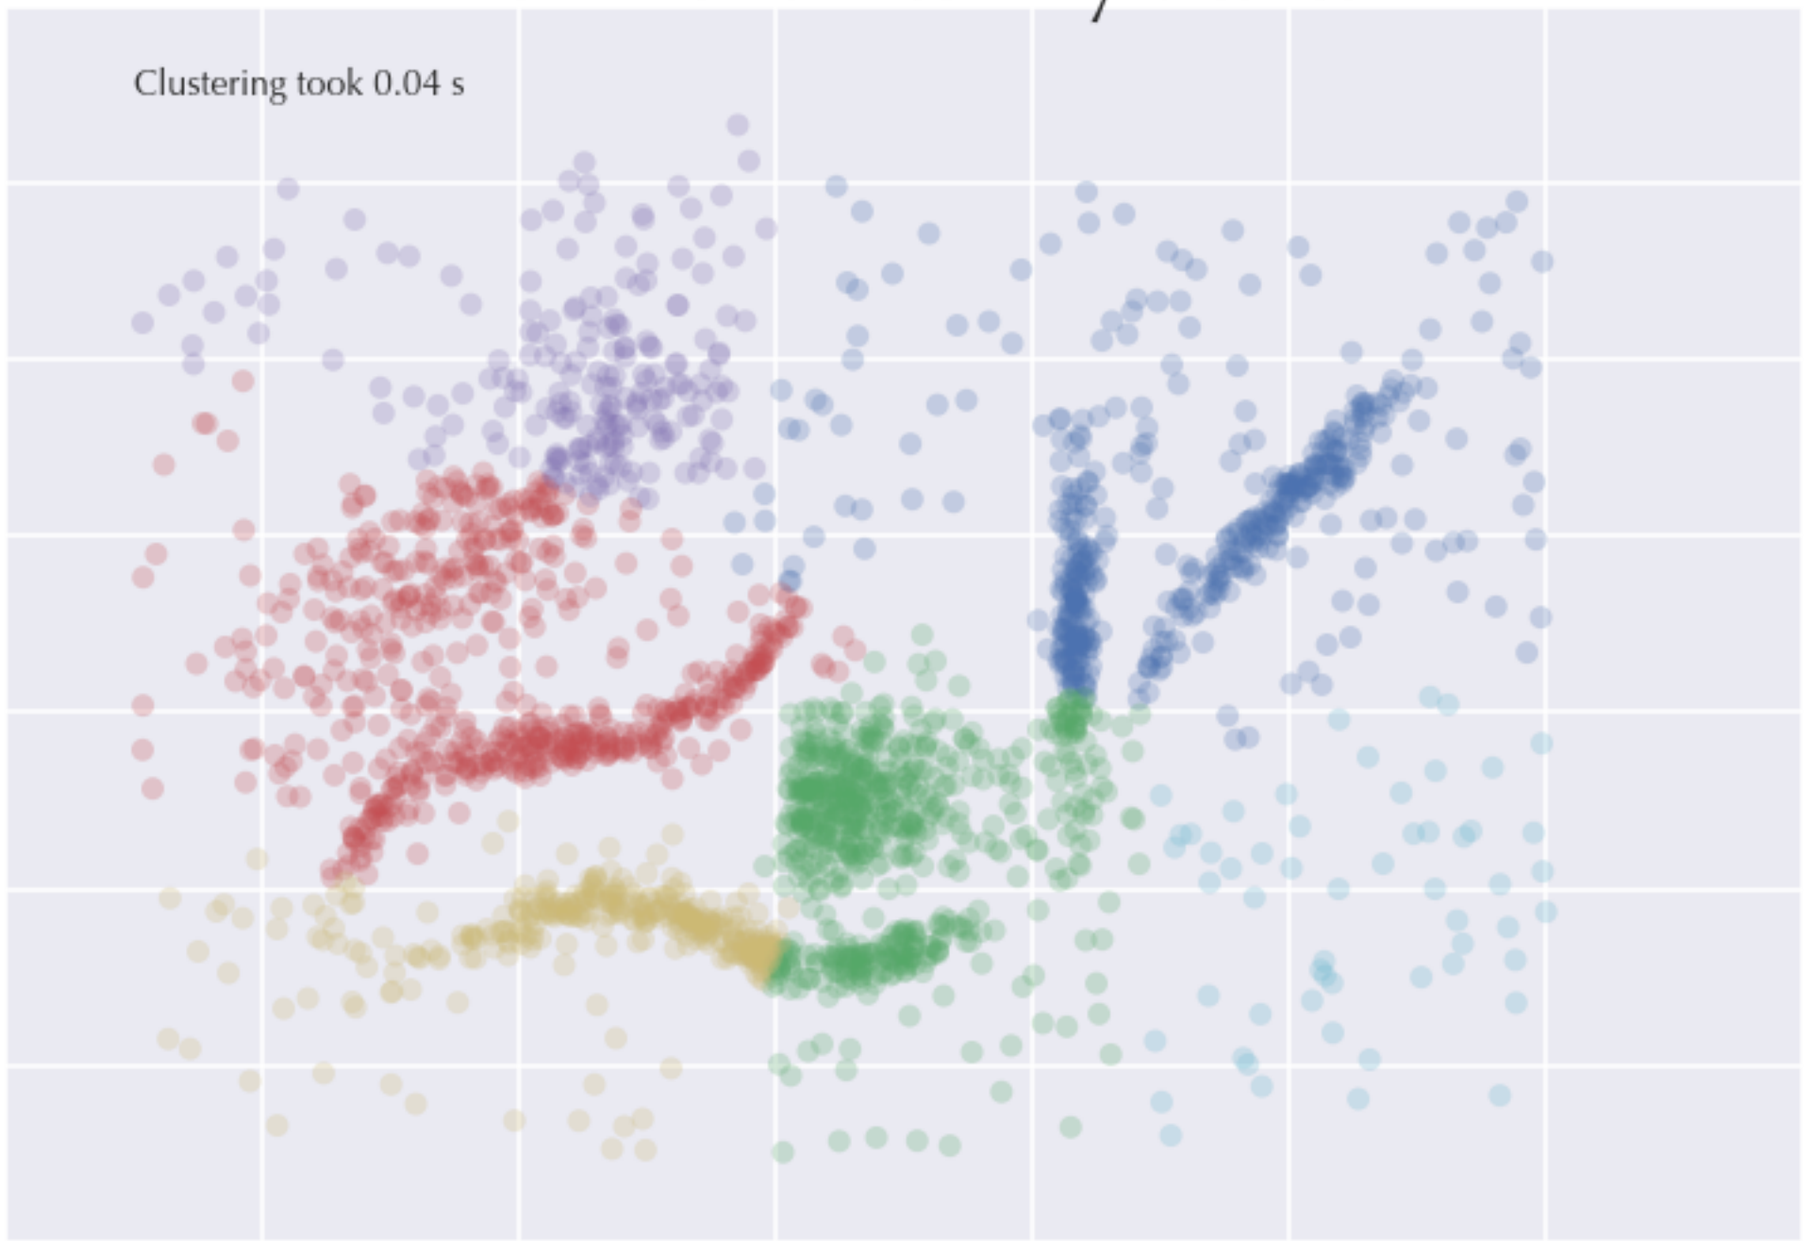

INSTITUT TUTTE INSTITUTE

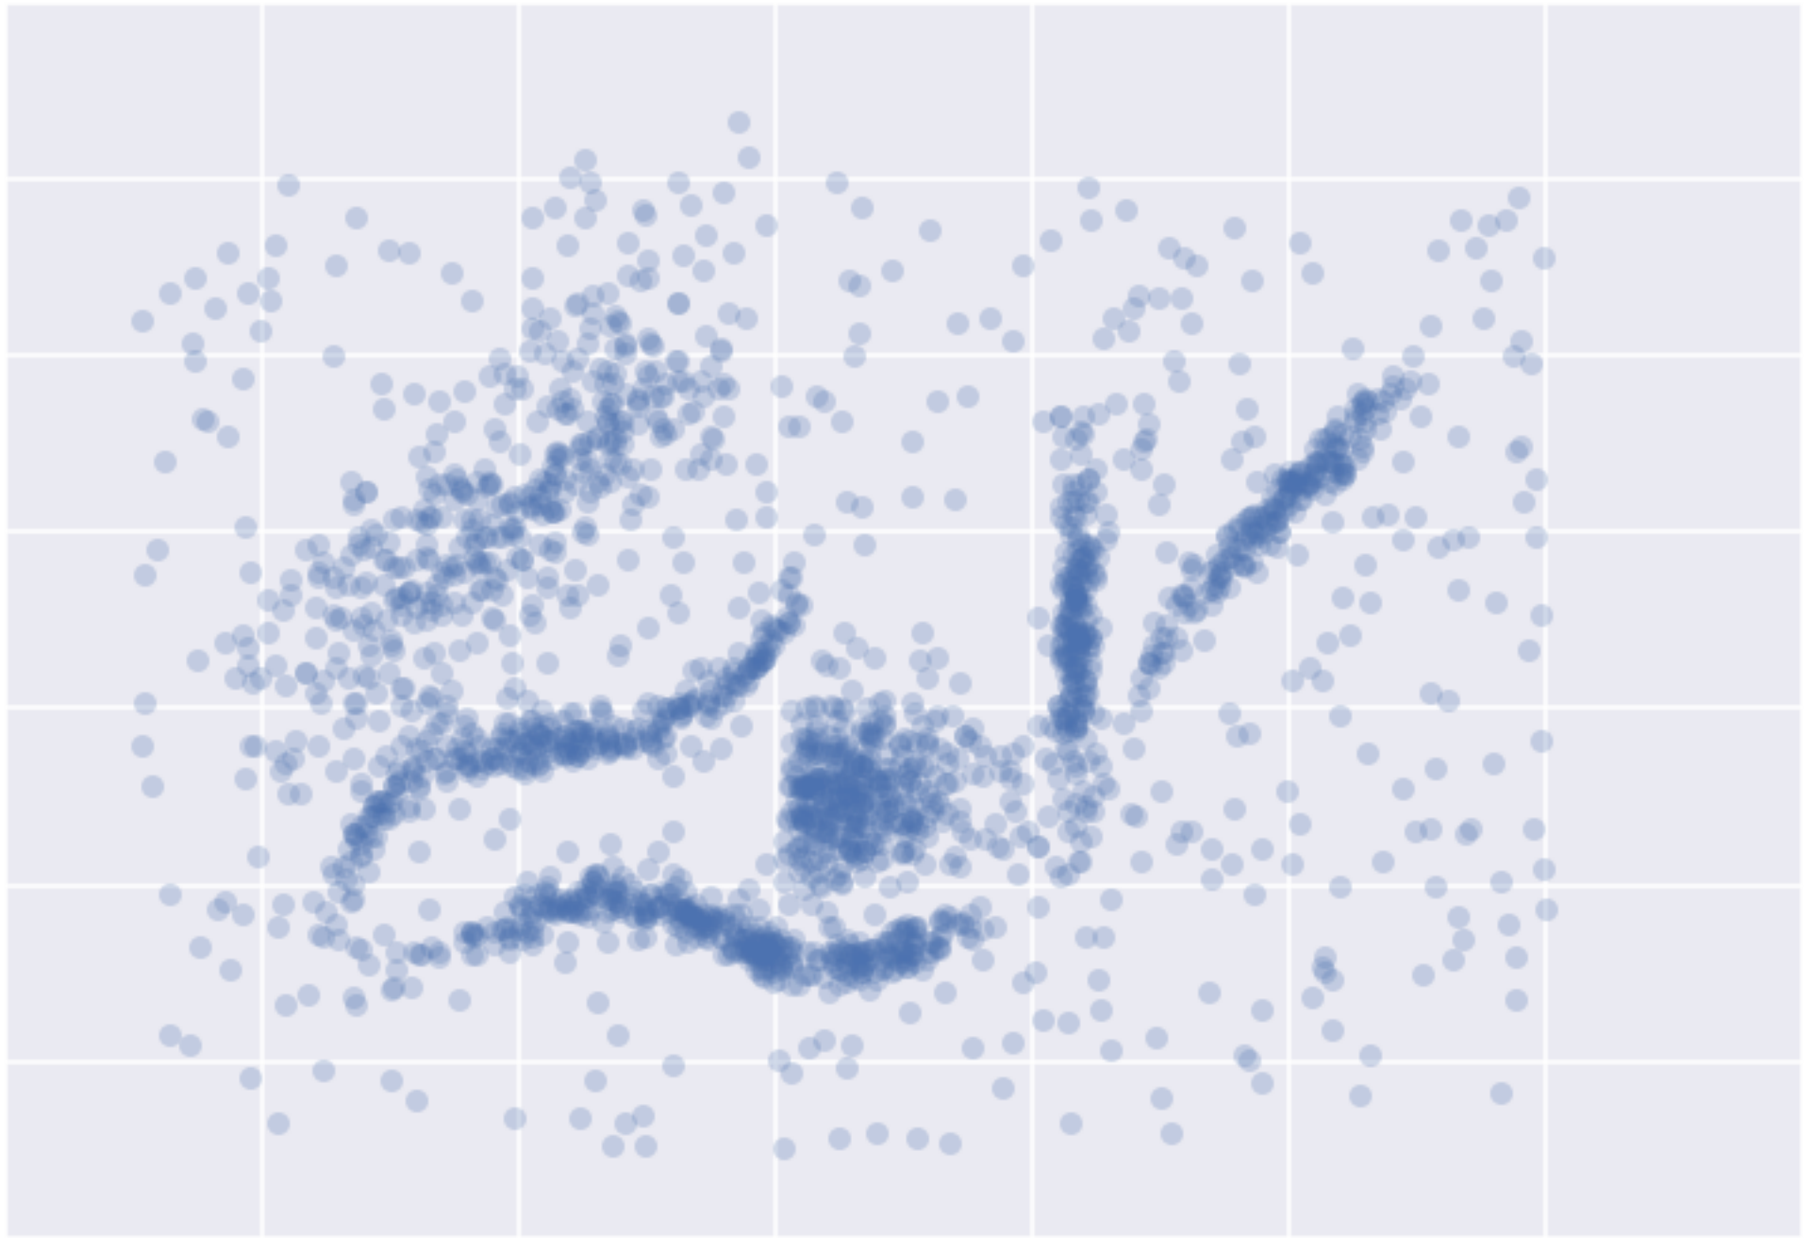

INSTITUT TUTTE INSTITUTE

Clusters found by AgglomerativeClustering

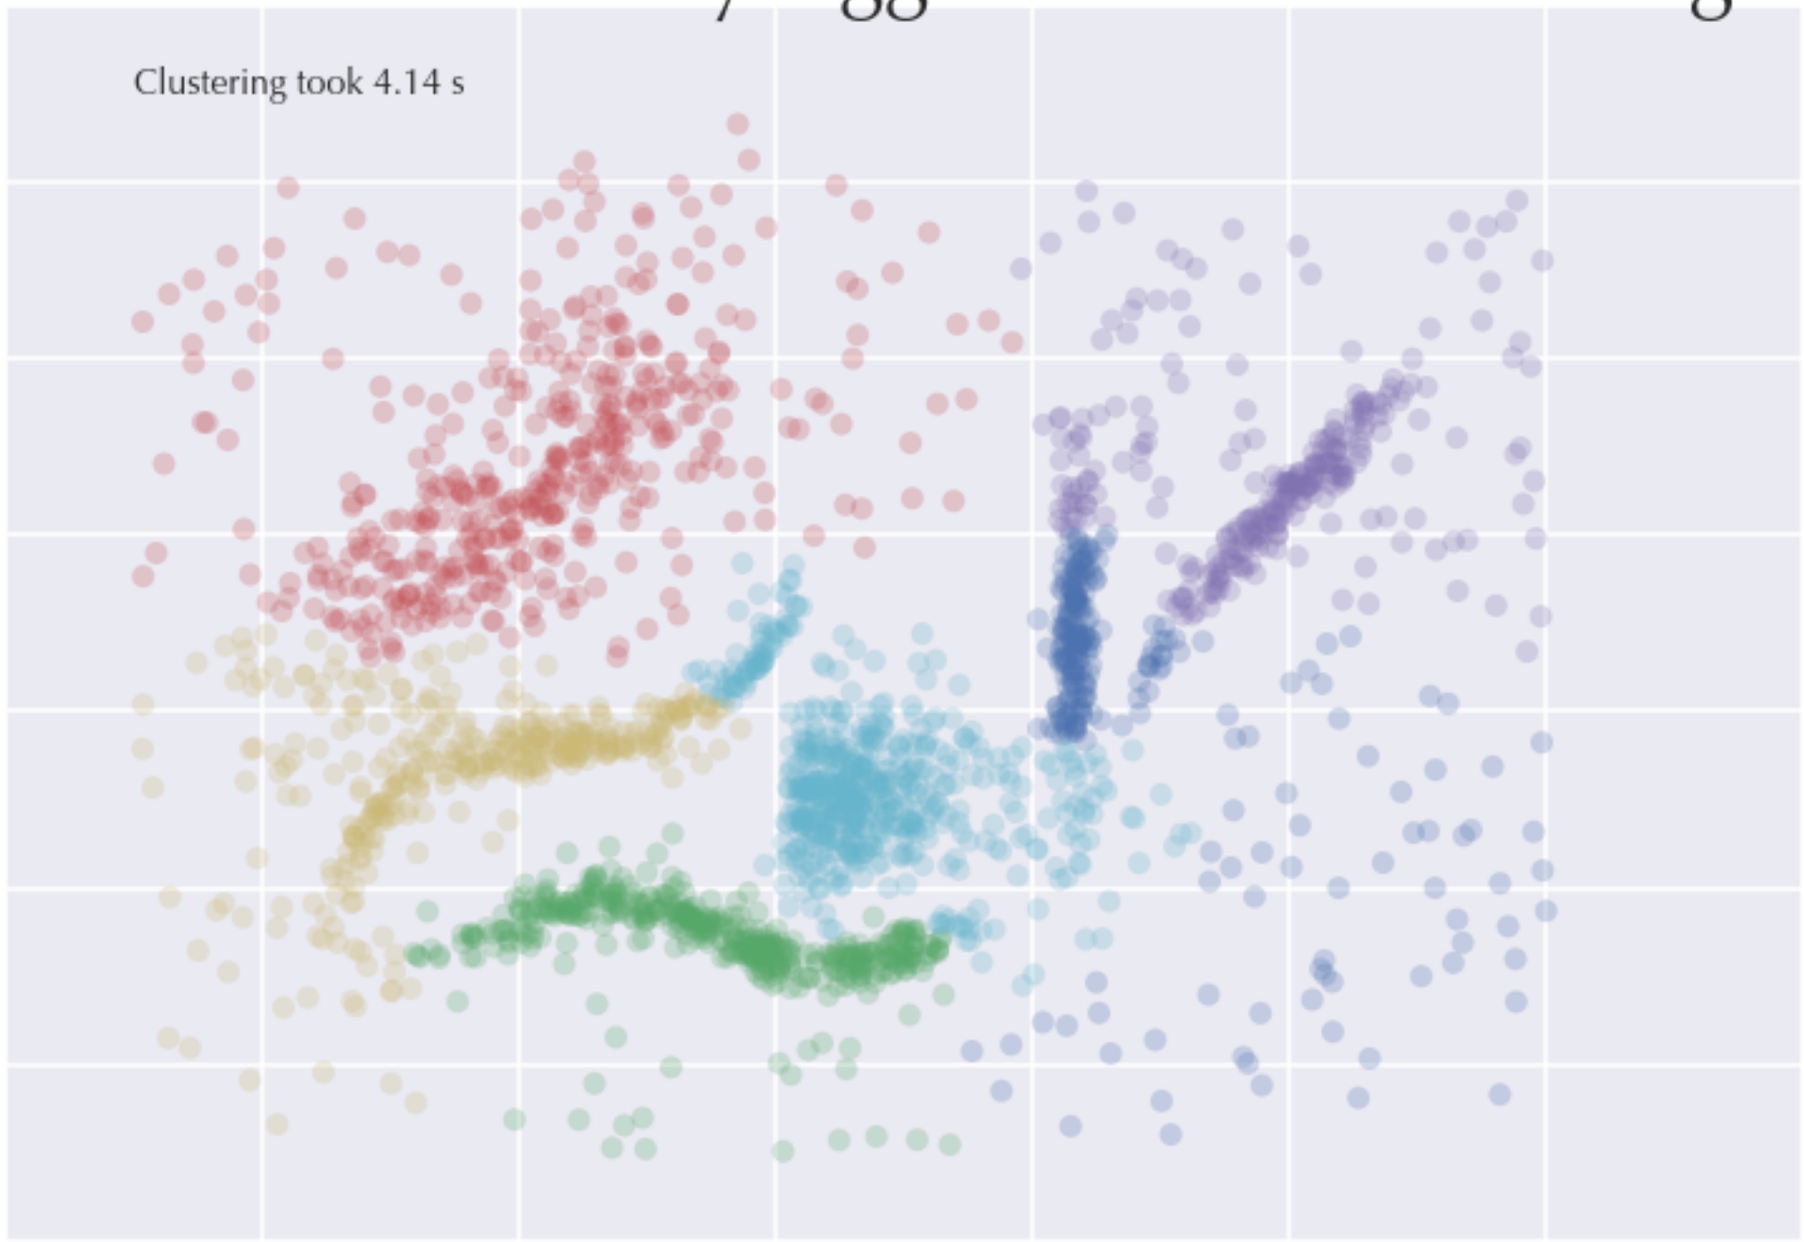

INSTITUT TUTTE INSTITUTE

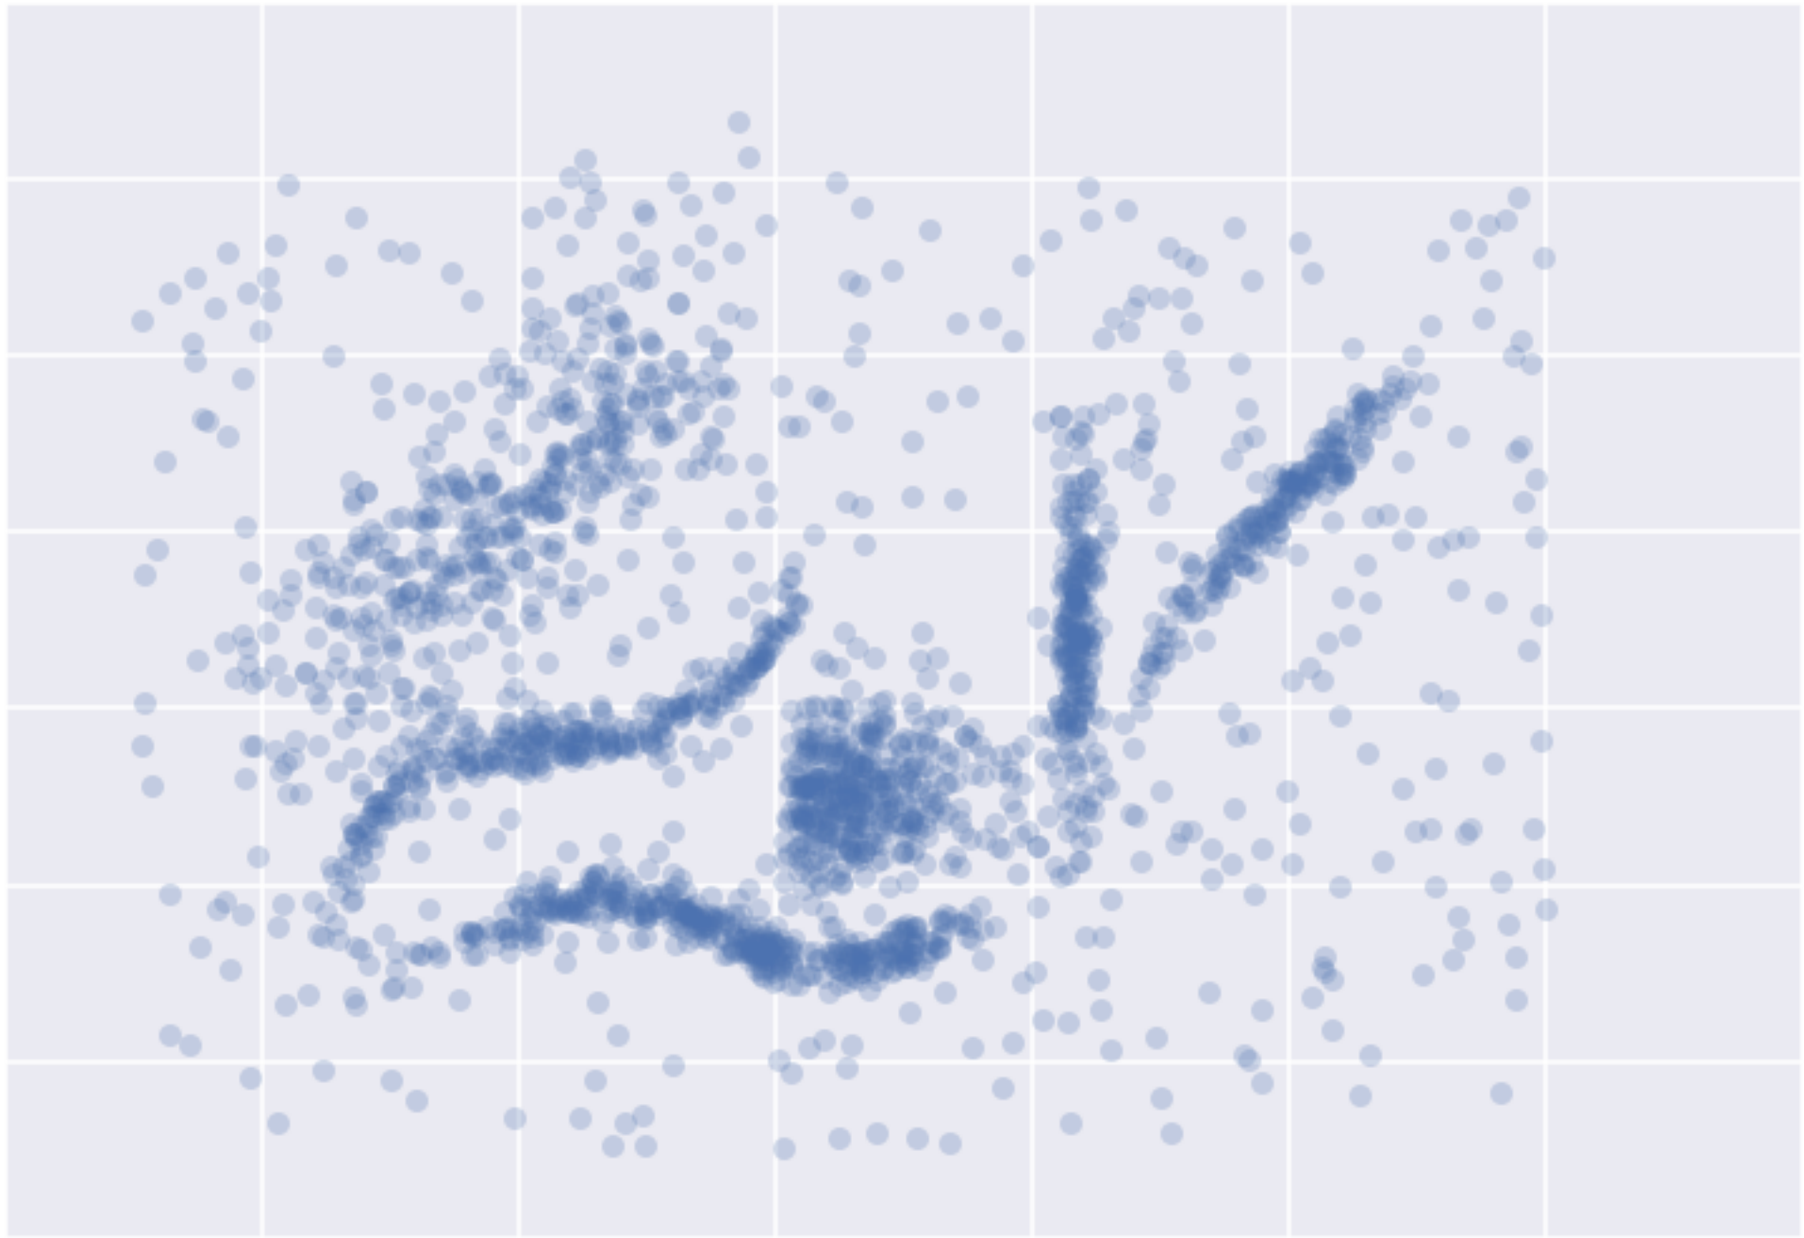

INSTITUT TUTTE INSTITUTE

Clusters found by DBSCAN

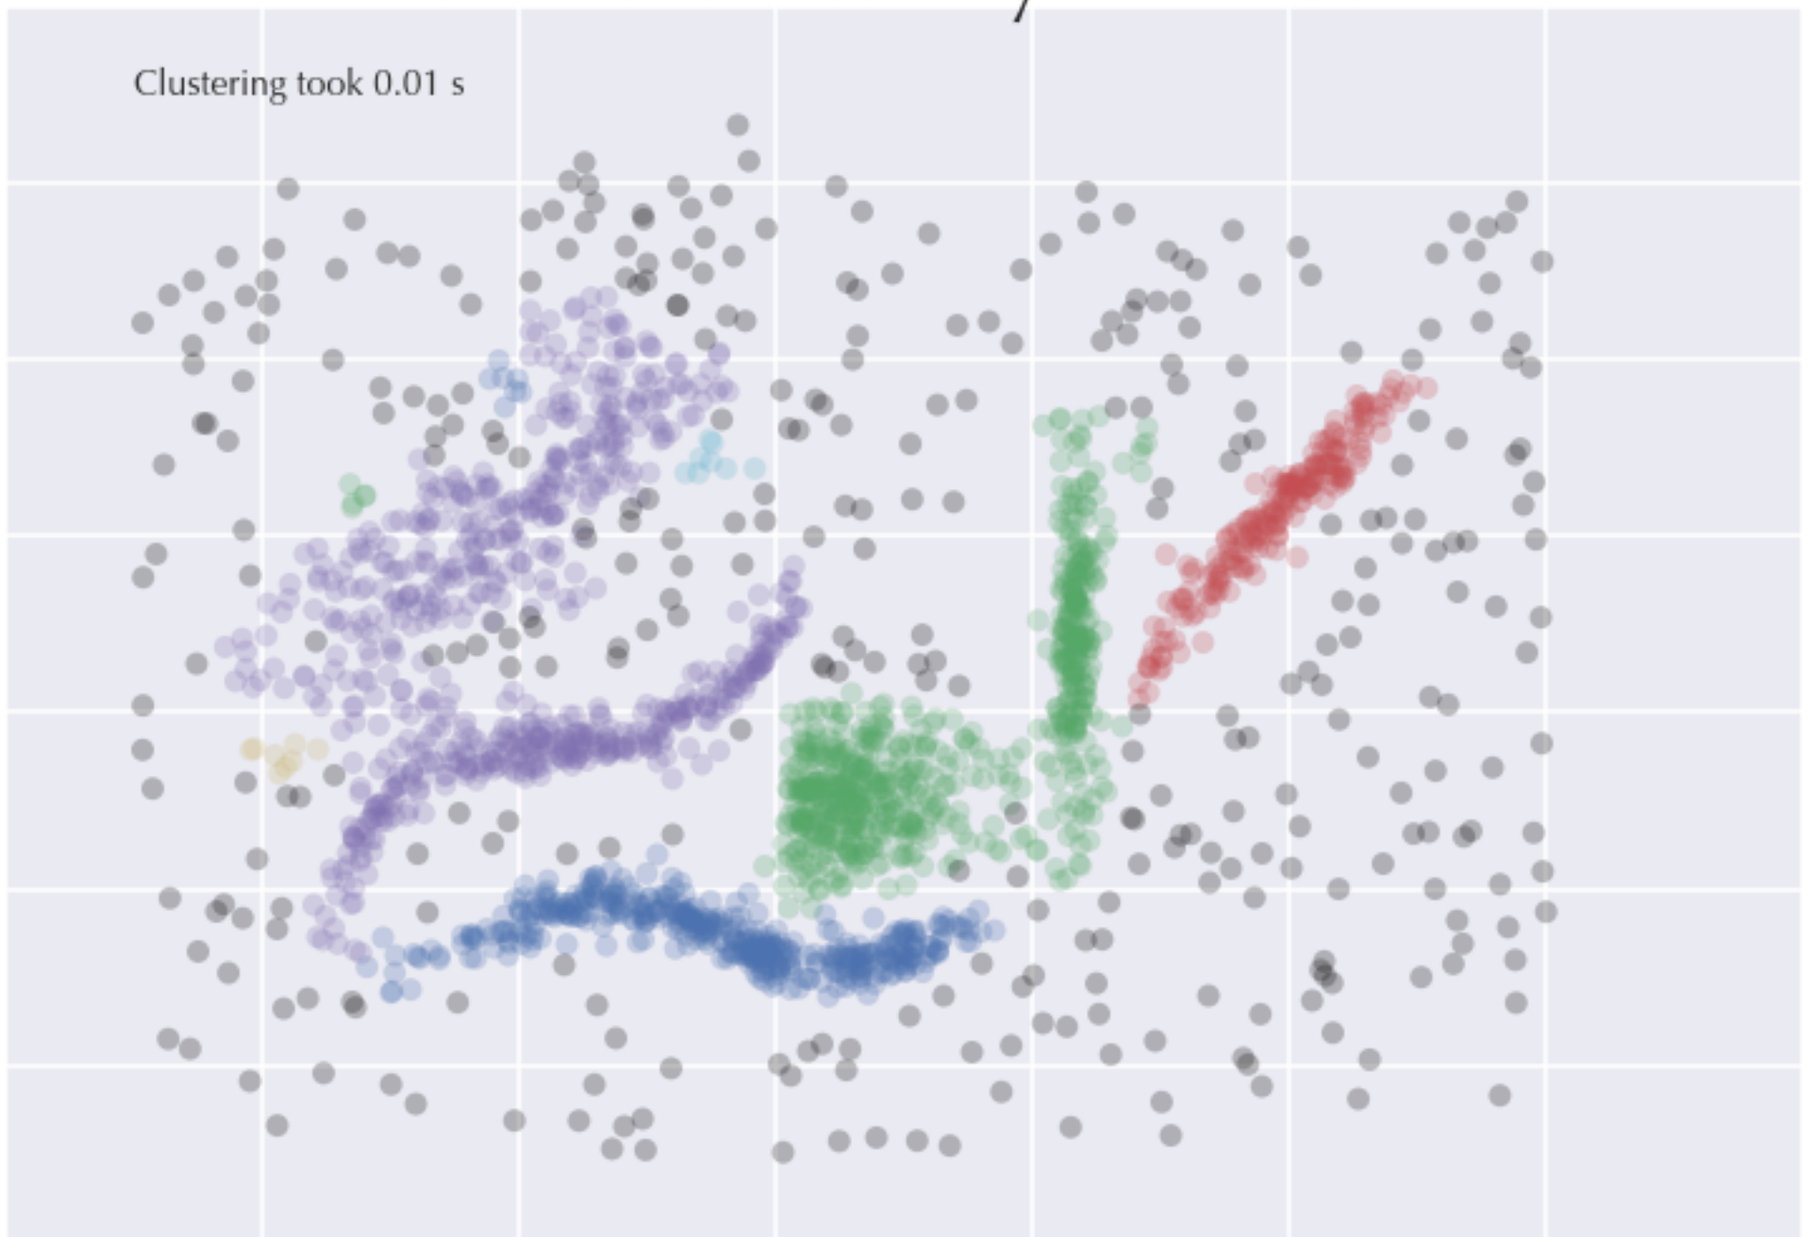

INSTITUT TUTTE INSTITUTE

What makes it so hard?

INSTITUT TUTTE INSTITUTE

What makes a clustering good? How do you tell?

INSTITUT TUTTE INSTITUTE

Which clustering is better?

INSTITUT TUTTE INSTITUTE

A Possible Clustering

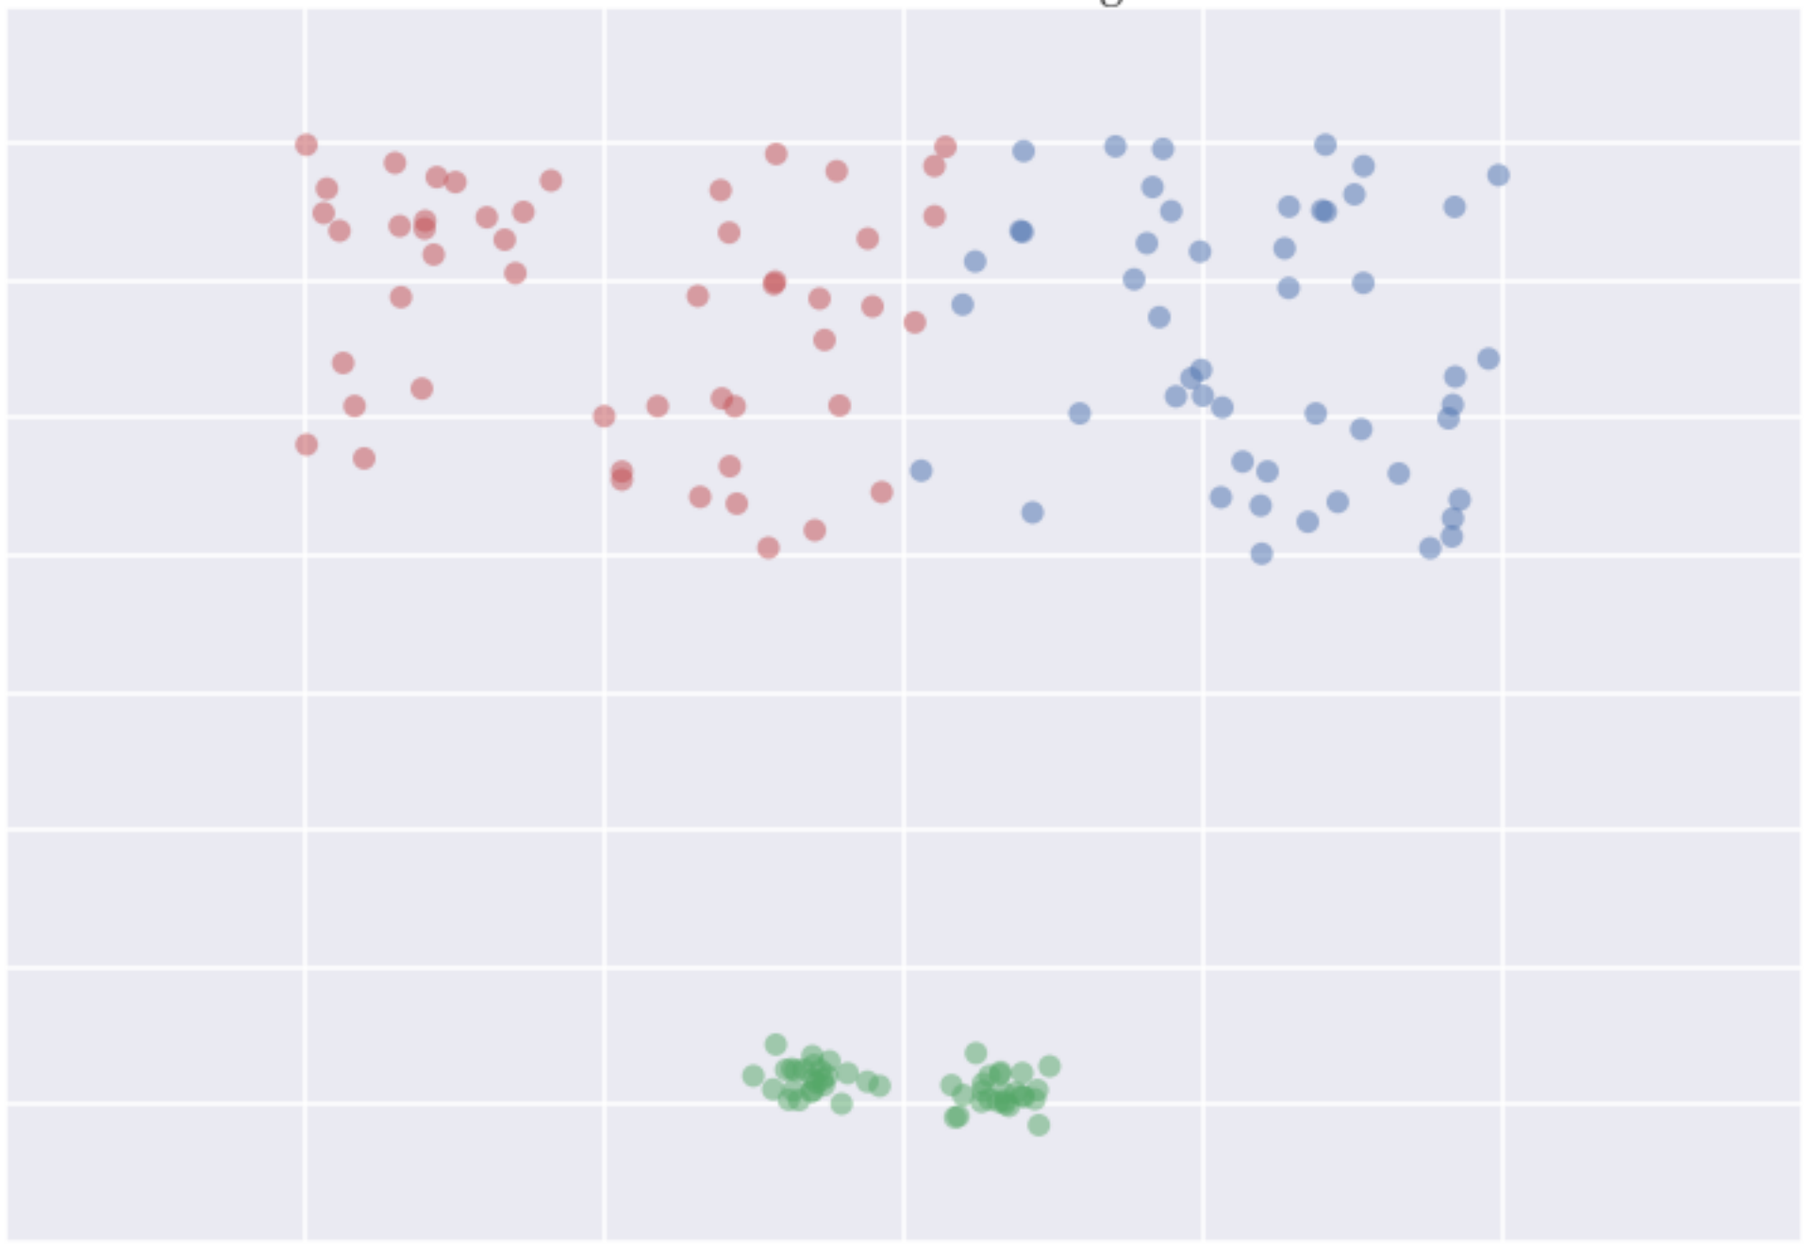

INSTITUT TUTTE INSTITUTE

INSTITUT TUTTE INSTITUTE

What is a cluster?

INSTITUT TUTTE INSTITUTE

What do I mean by a cluster?

INSTITUT TUTTE INSTITUTE

Let's start by looking at what not to do

INSTITUT TUTTE INSTITUTE

Clusters found by KMeans

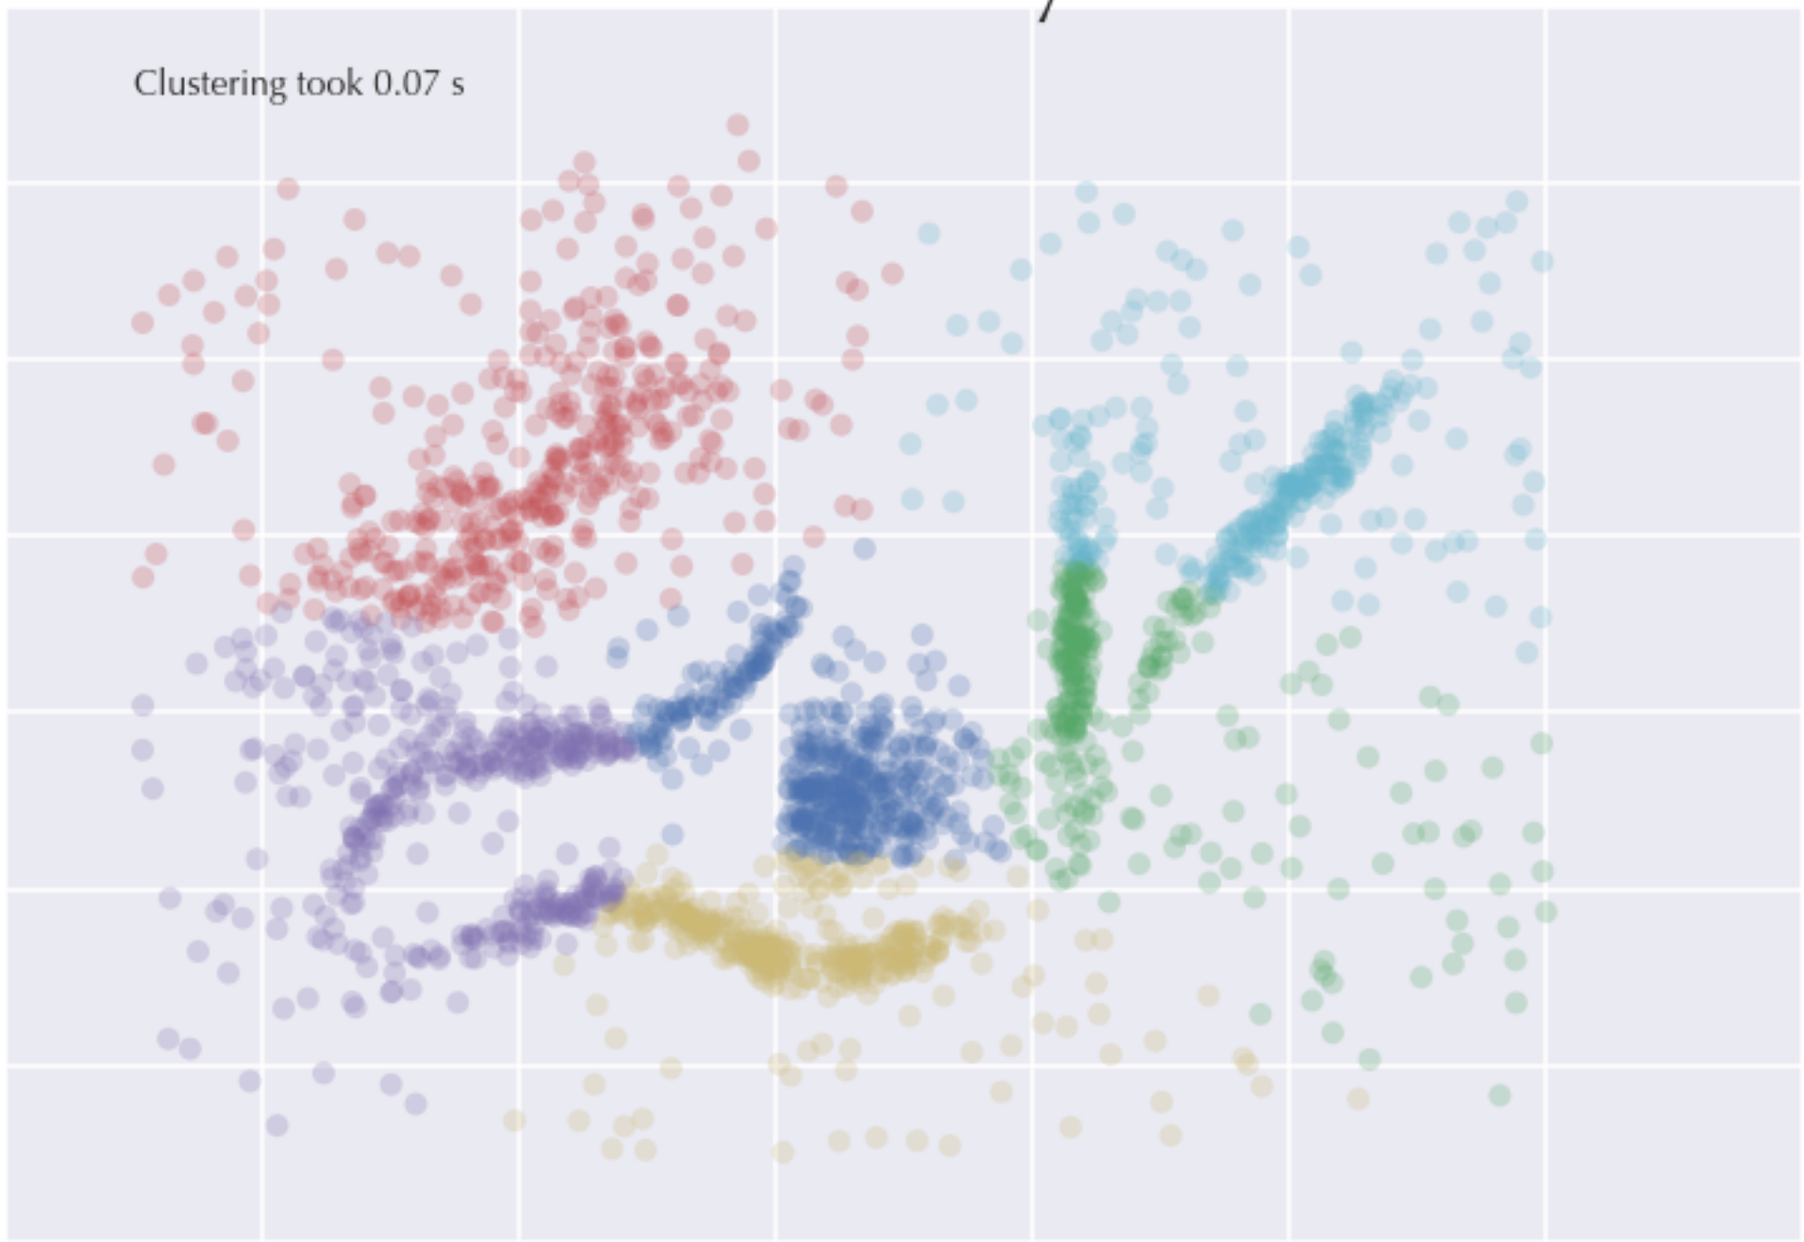

INSTITUT TUTTE INSTITUTE

Clusters need not be balls!

INSTITUT TUTTE INSTITUTE

Clusters found by SpectralClustering

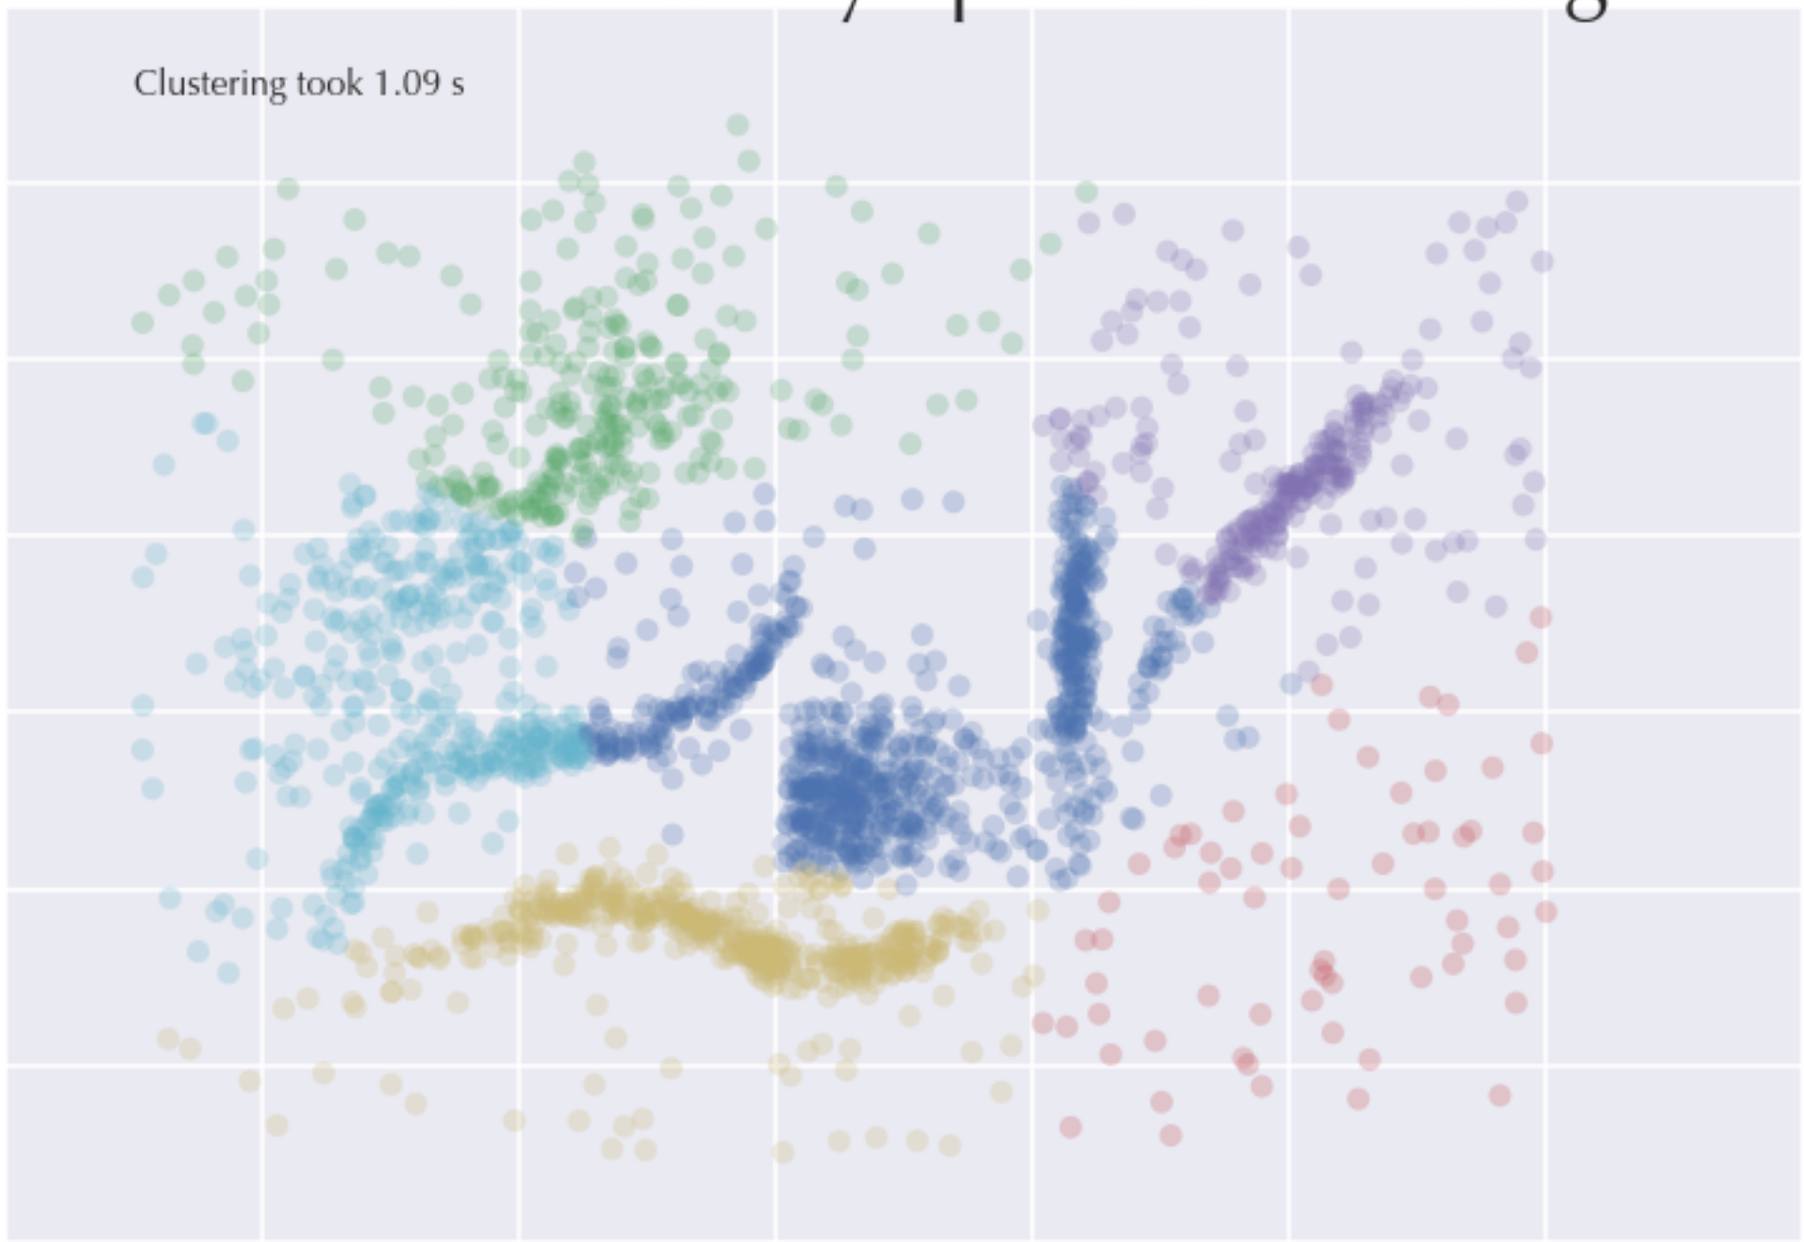

Not every point is in a cluster

Real data has noise!

INSTITUT TUTTE INSTITUTE

Clusters found by DBSCAN

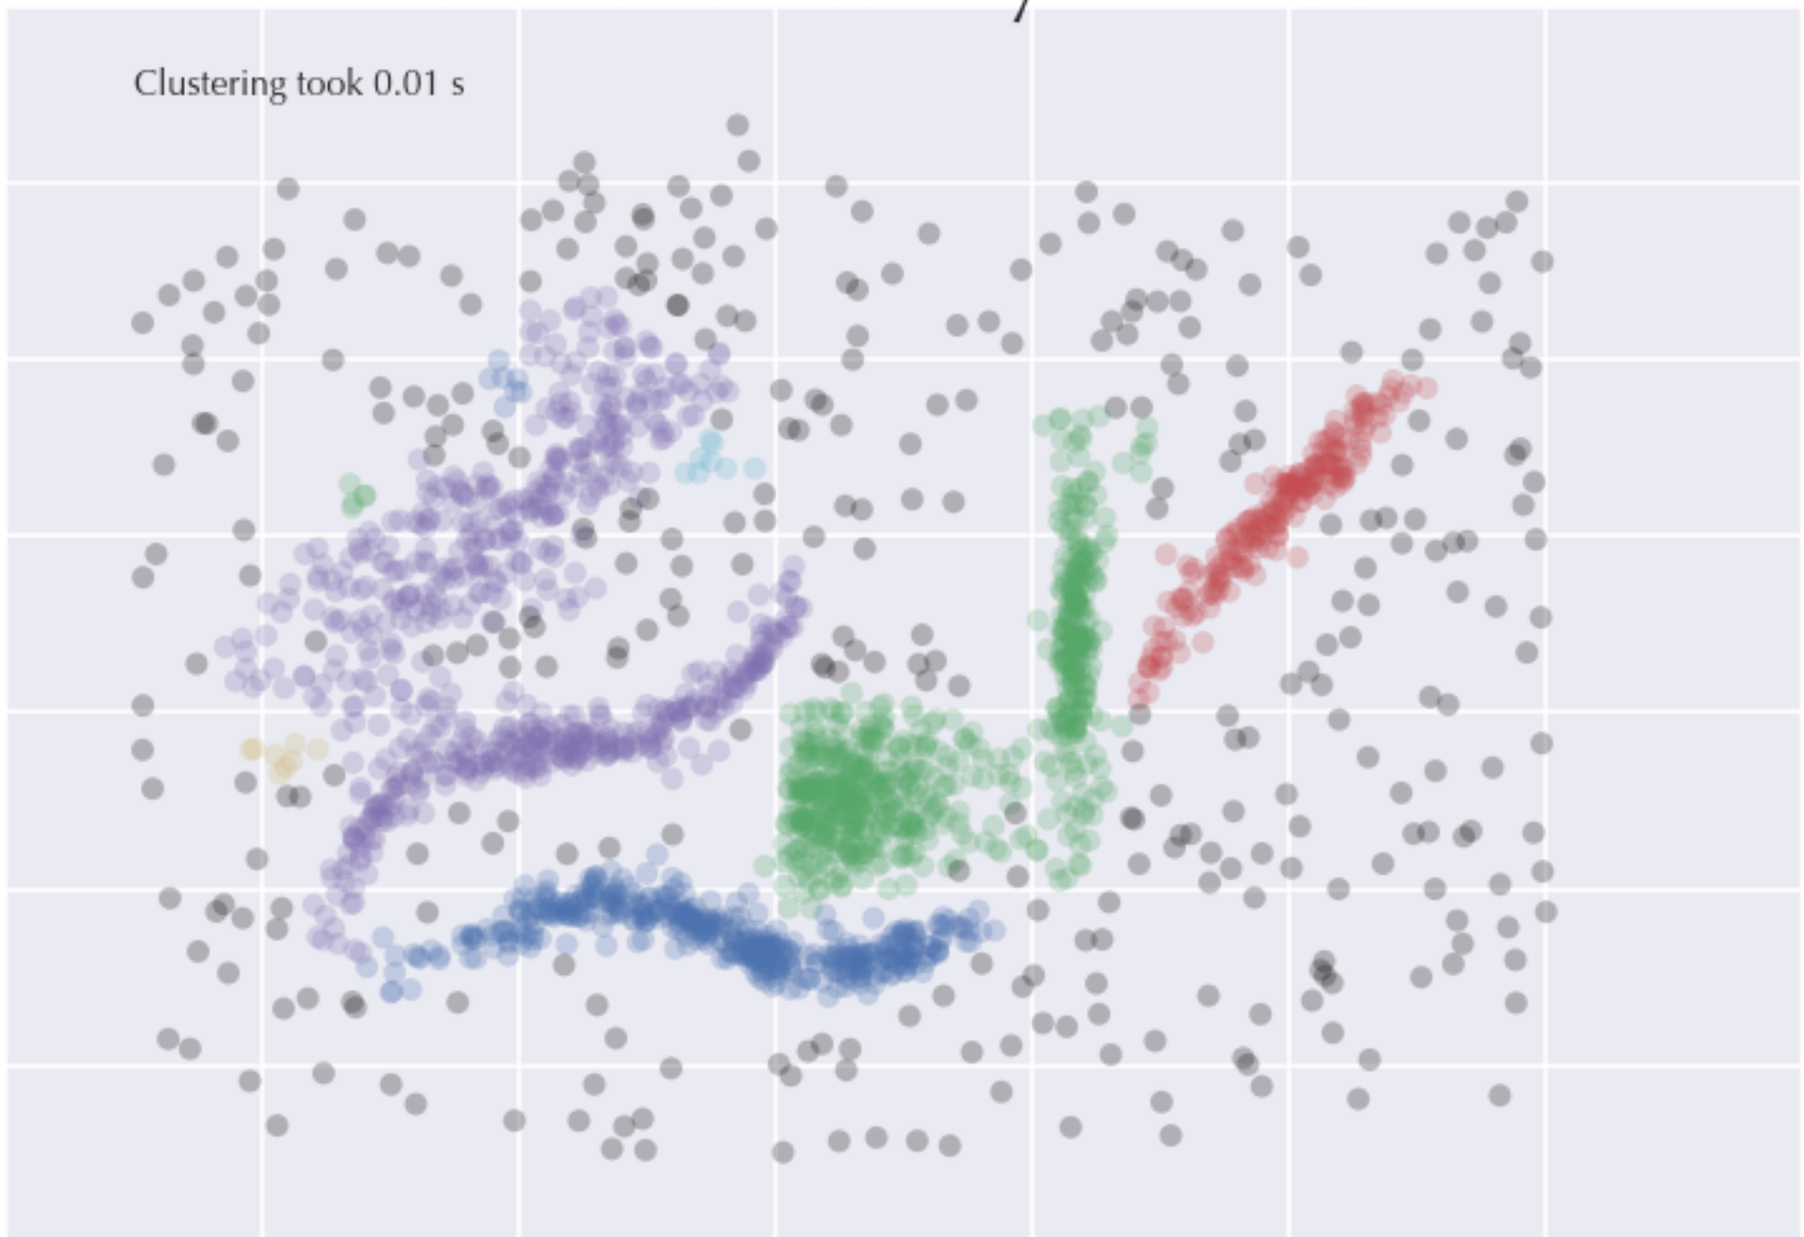

INSTITUT TUTTE INSTITUTE

Clusters are dense areas?

Separated by less dense areas?

INSTITUT TUTTE INSTITUTE

That's what I mean when I talk about clustering

INSTITUT TUTTE INSTITUTE

You should think hard about what you mean before doing any clustering

INSTITUT TUTTE INSTITUTE

Can we be more specific about density-based clusters?

INSTITUT TUTTE INSTITUTE

A cluster is …

INSTITUT TUTTE INSTITUTE

A connected component of a level set of the probability density function of the underlying (and unknown) distribution from which our data samples are drawn.

INSTITUT TUTTE INSTITUTE

What?

INSTITUT TUTTE INSTITUTE

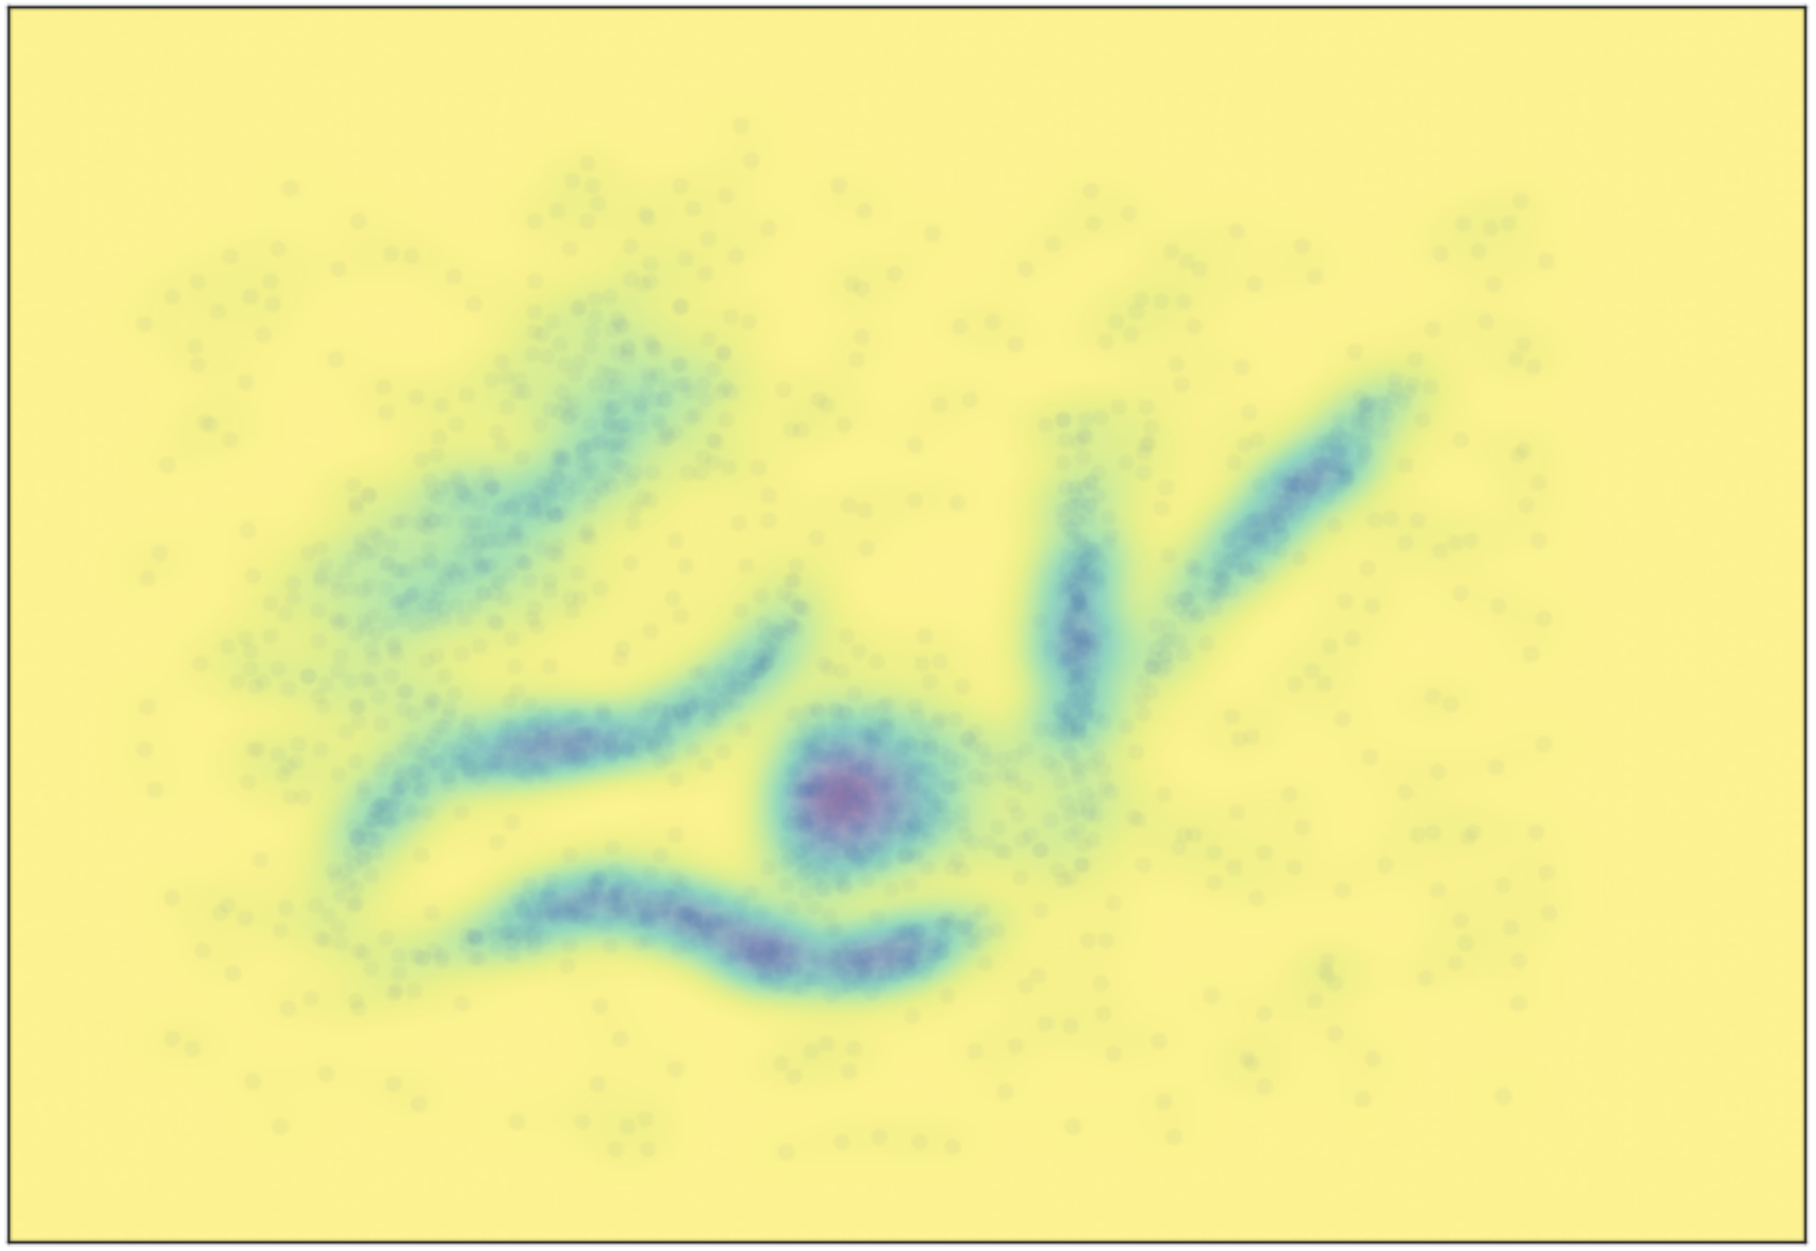

INSTITUT TUTTE INSTITUTE

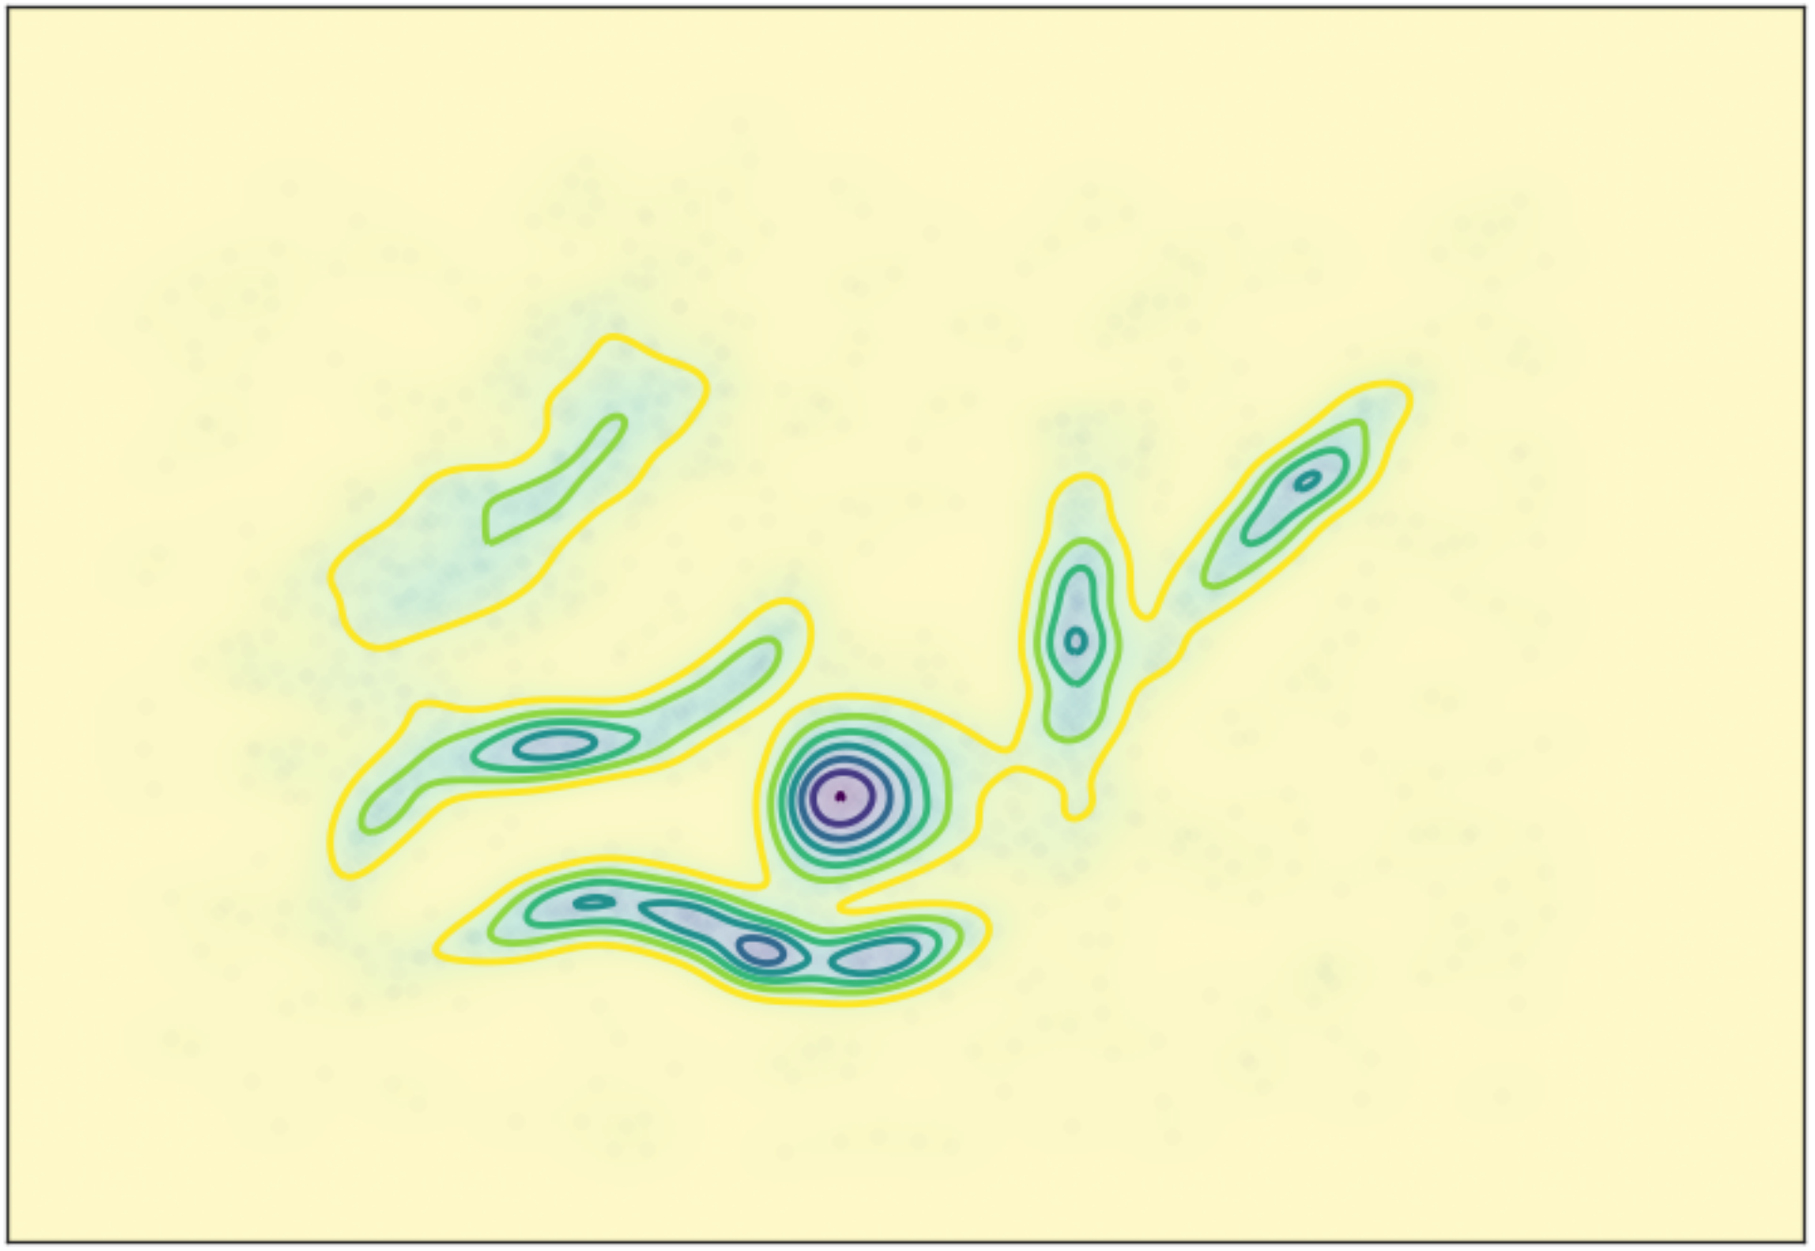

INSTITUT TUTTE INSTITUTE

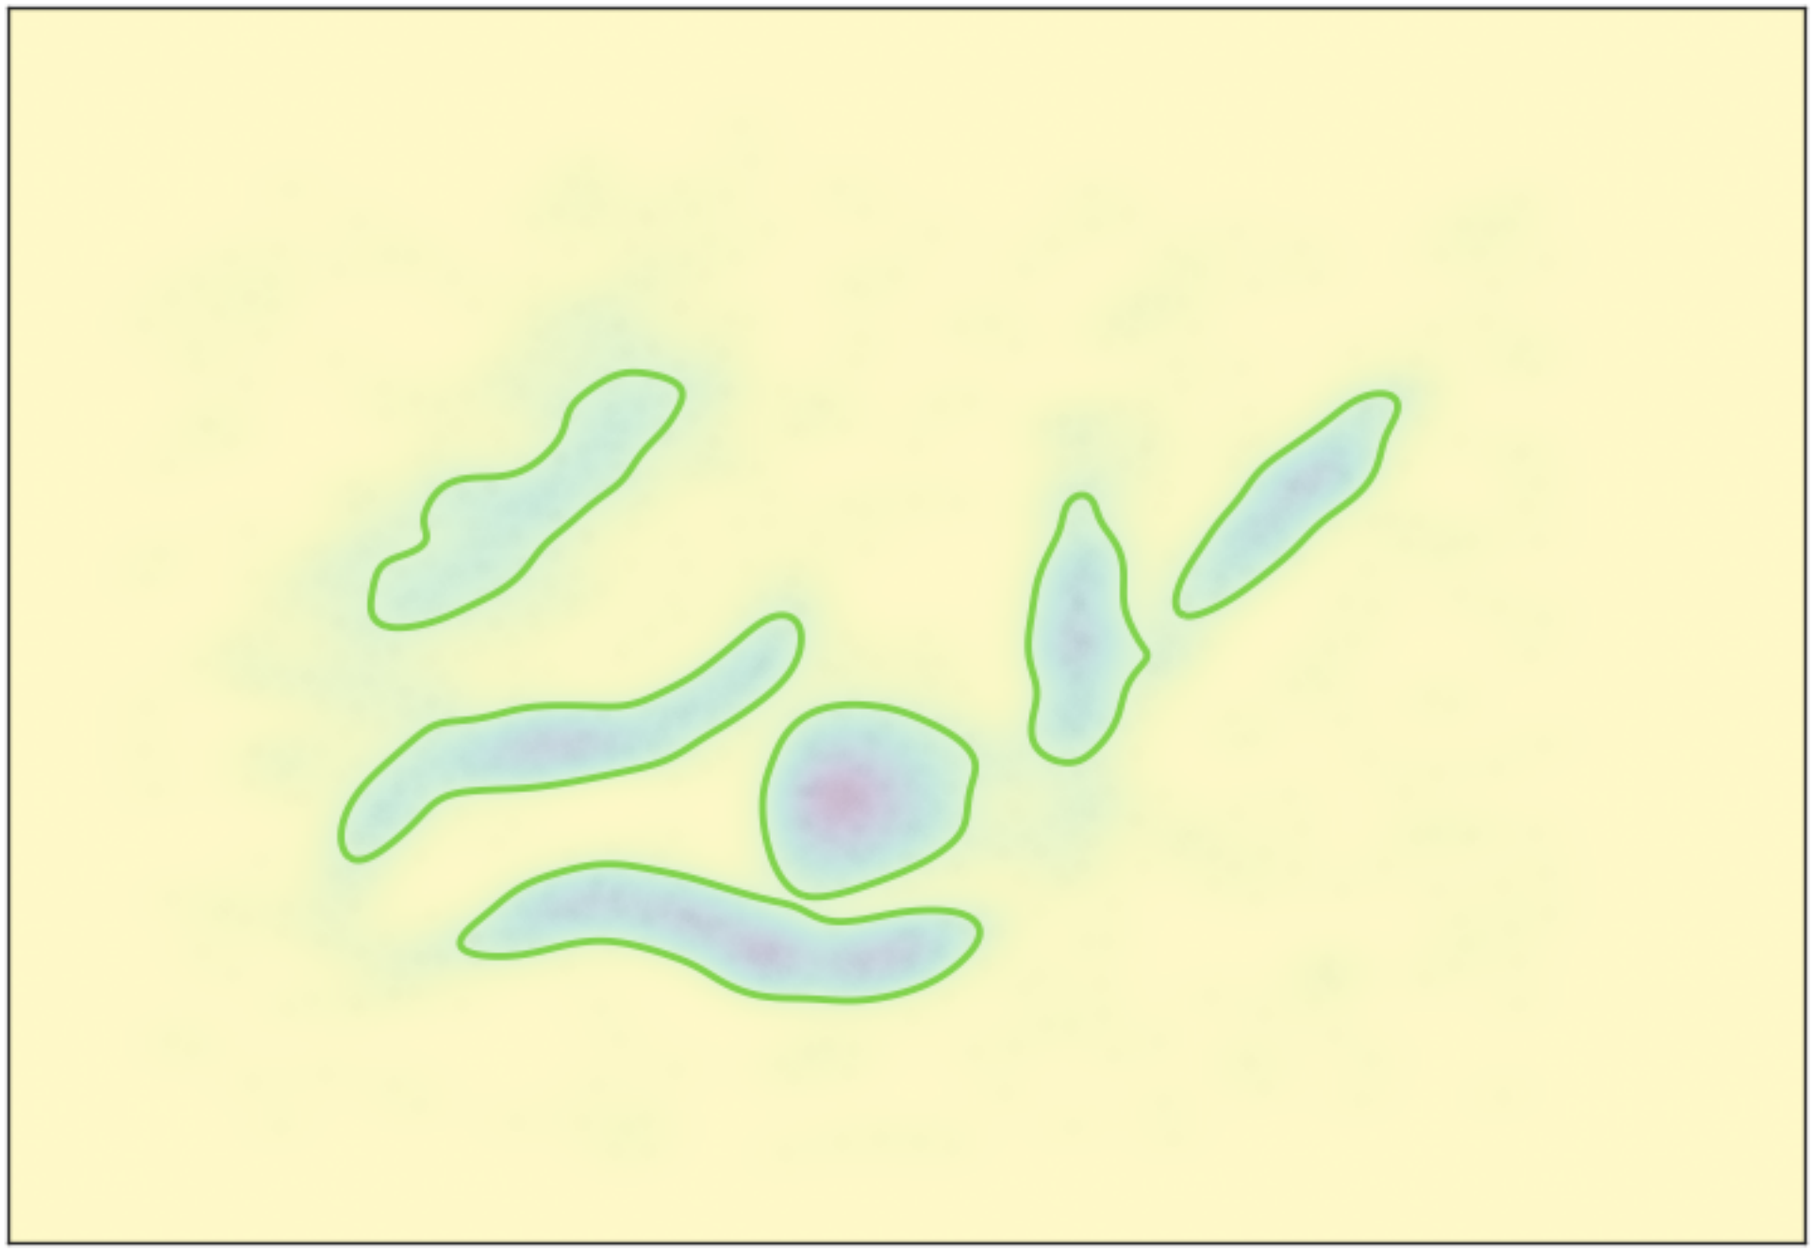

INSTITUT TUTTE INSTITUTE

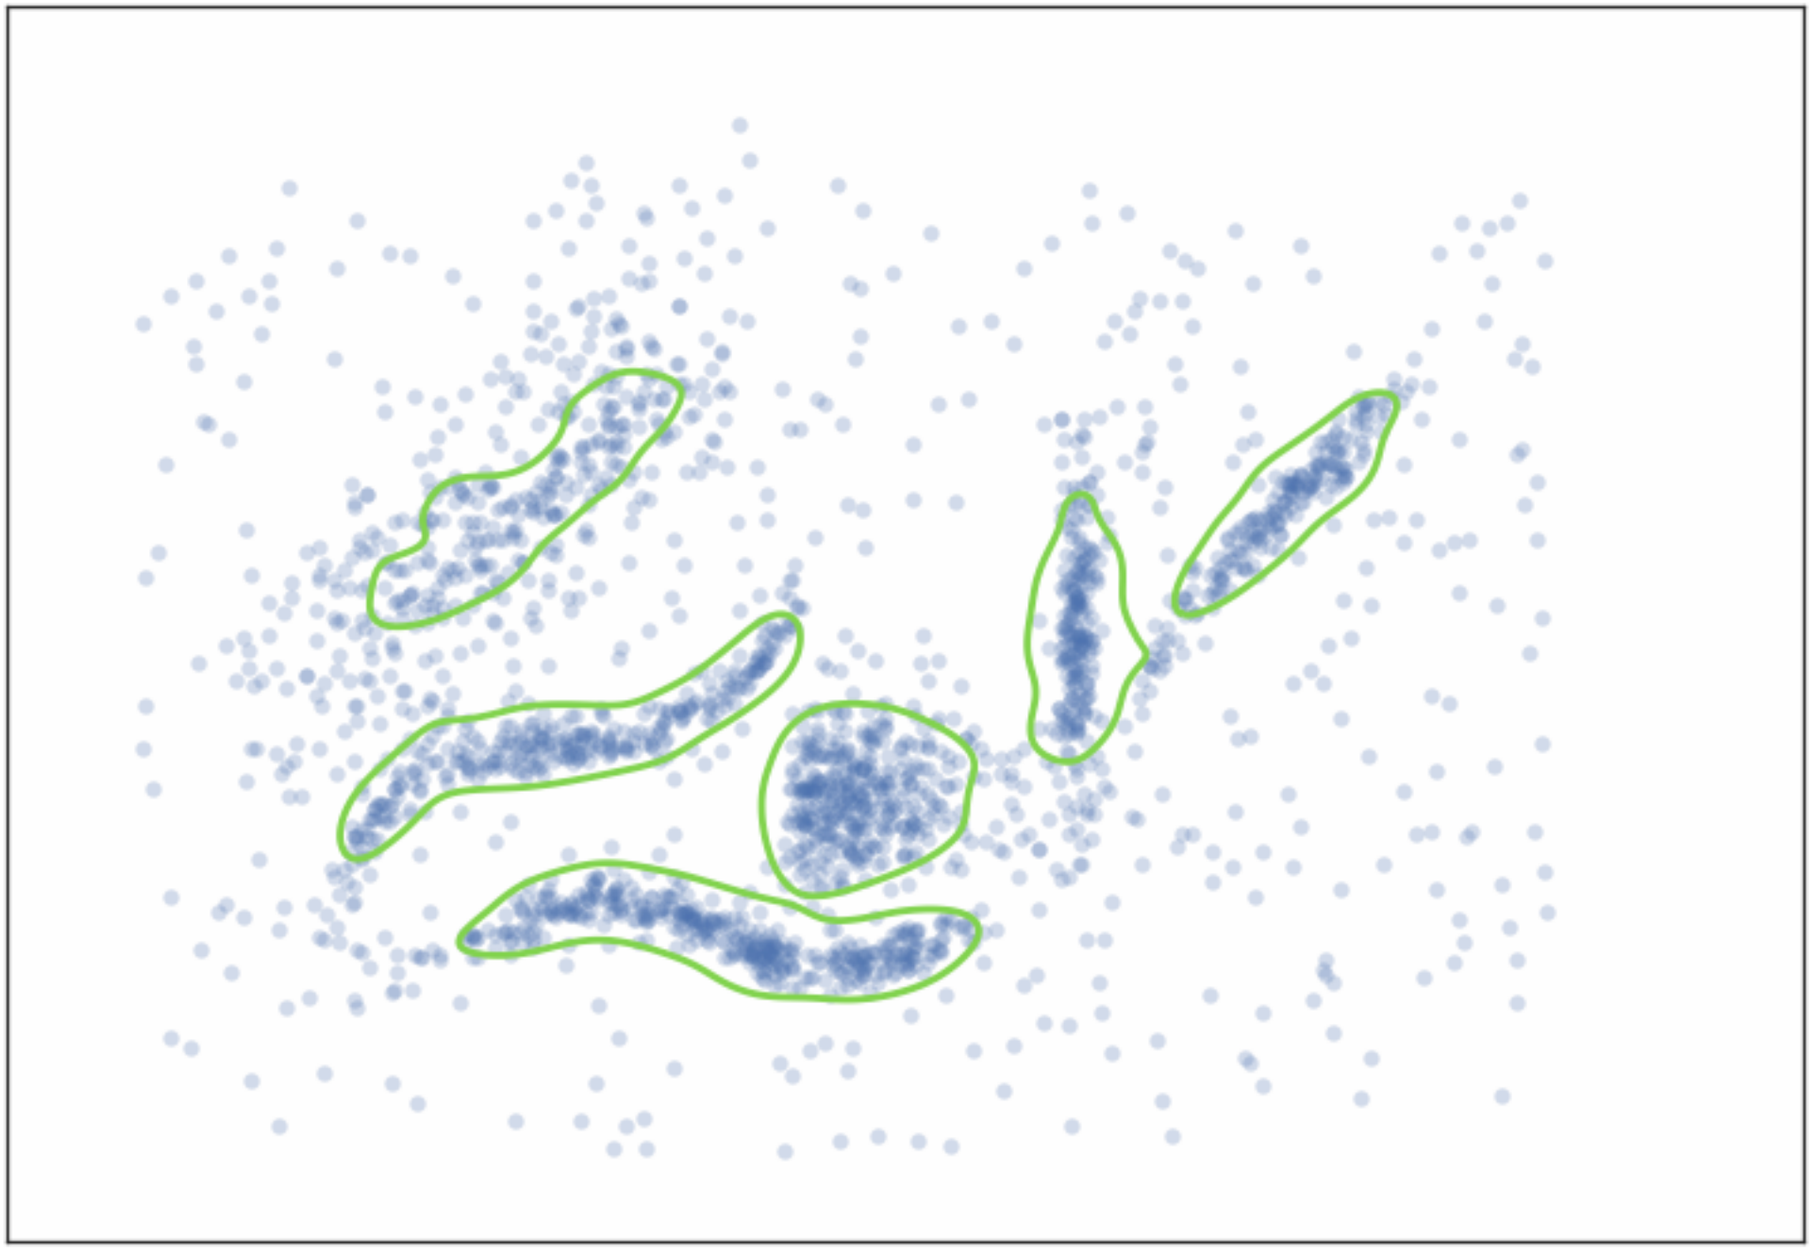

INSTITUT TUTTE INSTITUTE

Is that process a clustering algorithm?

1. We don’t know the PDF
2. Which level sets to choose?
3. Computational complexity

What can we do?

INSTITUT TUTTE INSTITUTE

Locally approximate the density

INSTITUT TUTTE INSTITUTE

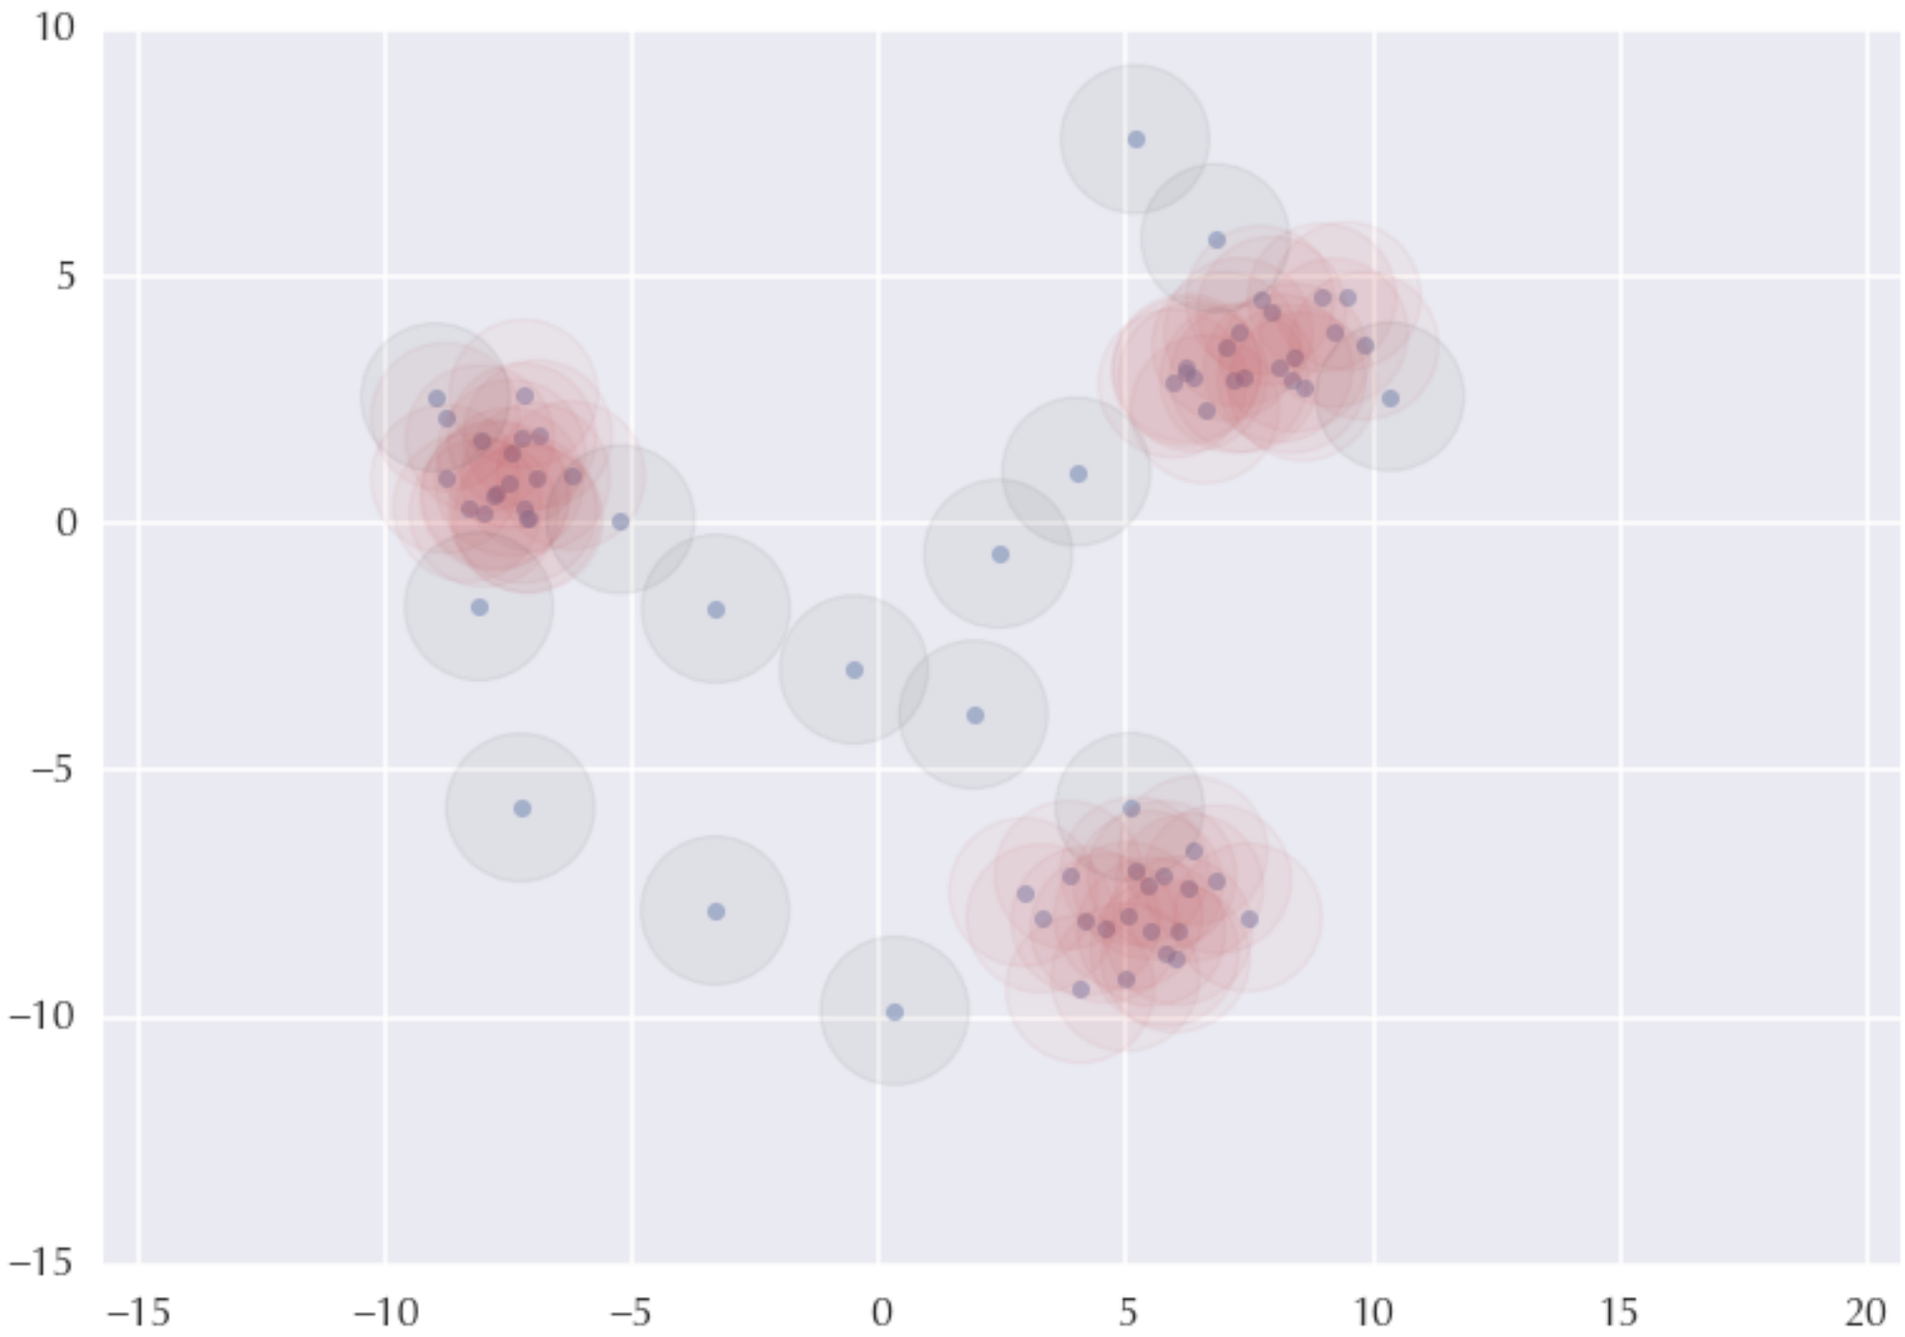

INSTITUT TUTTE INSTITUTE

INSTITUT TUTTE INSTITUTE

INSTITUT TUTTE INSTITUTE

The connected components of level sets form a tree

INSTITUT TUTTE INSTITUTE

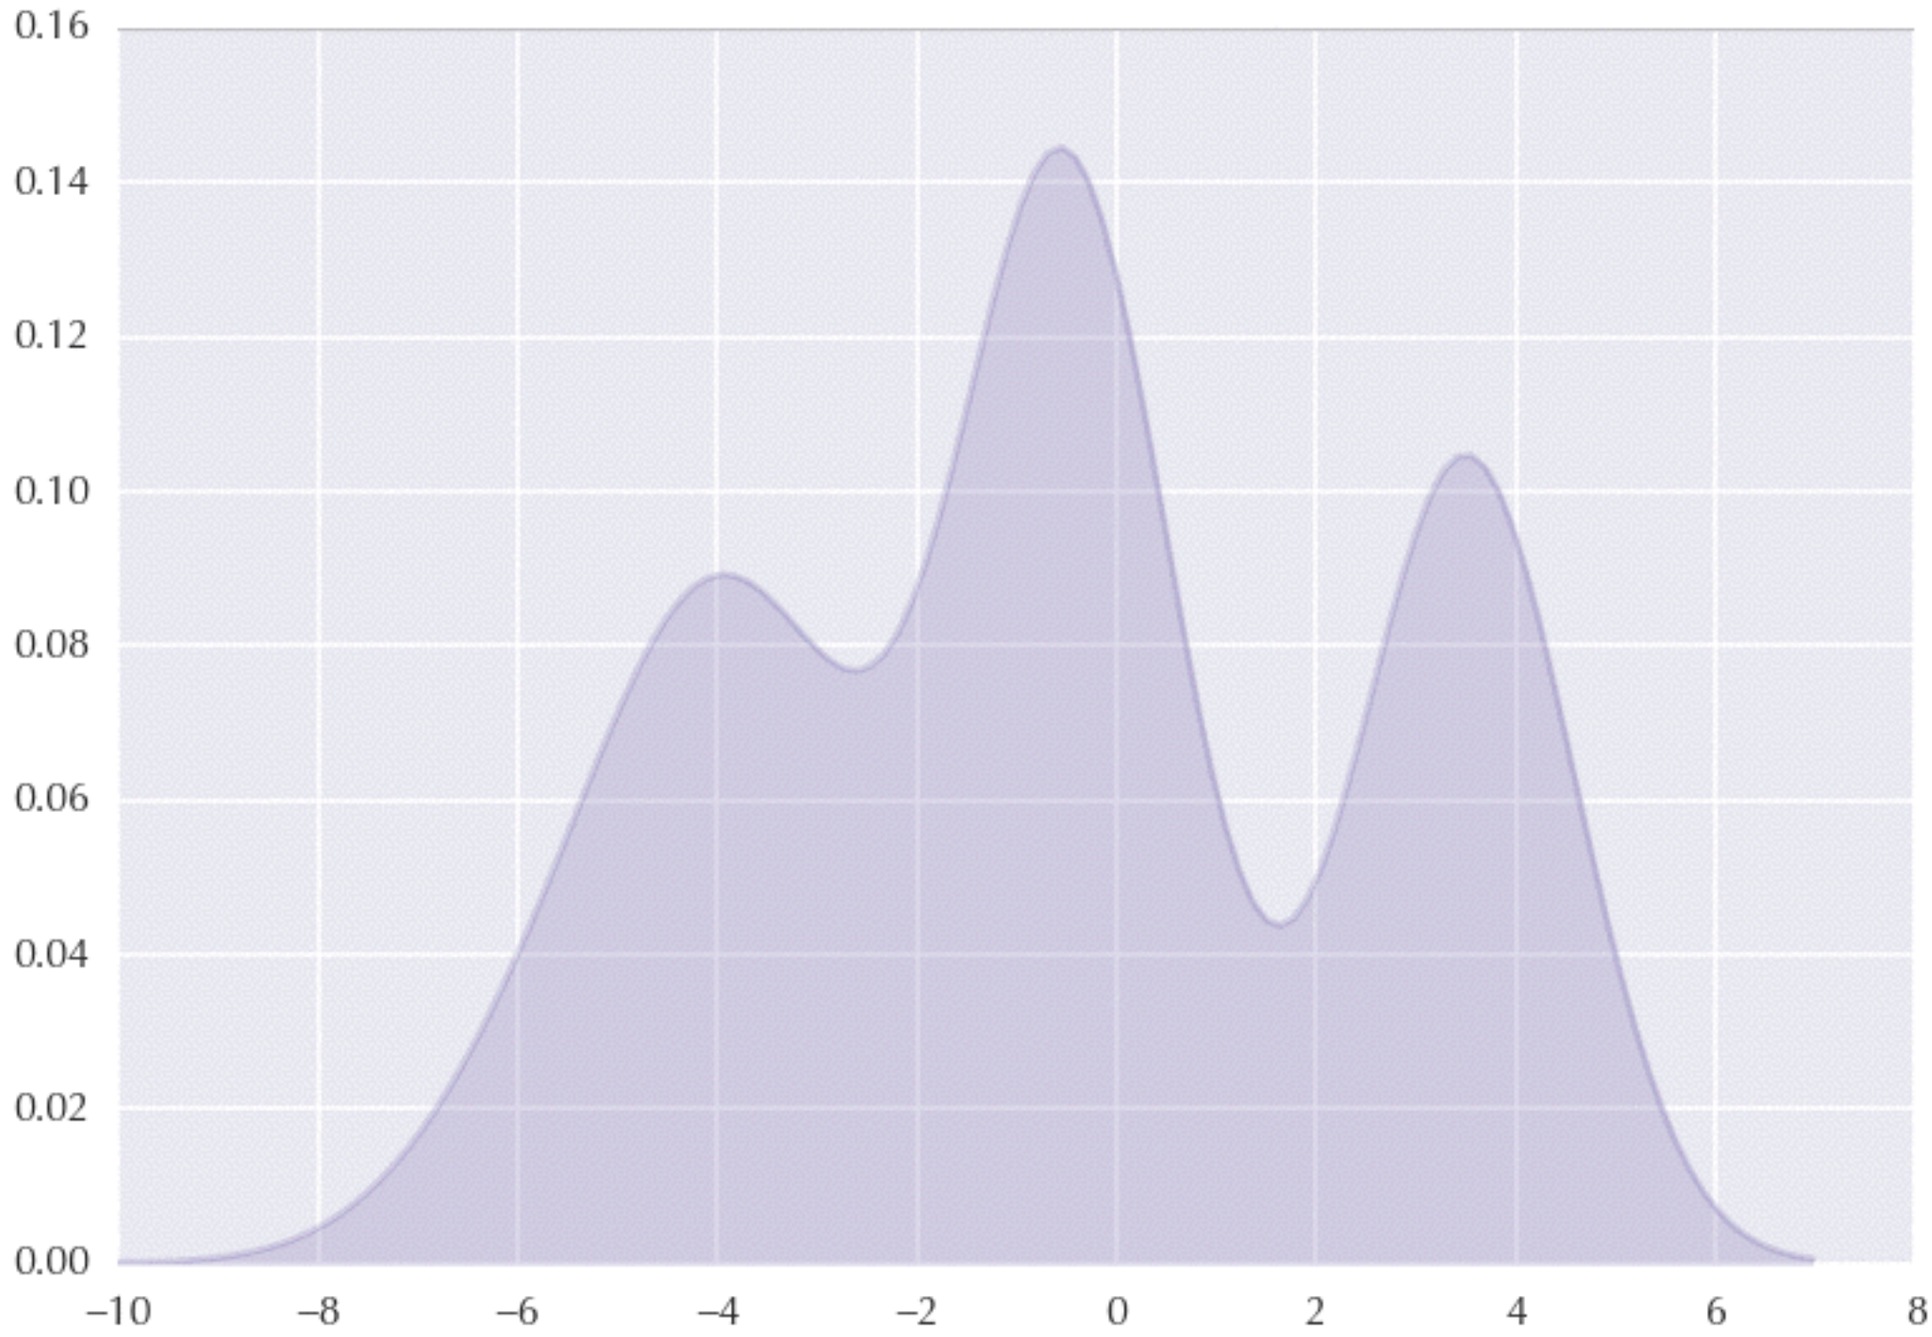

INSTITUT TUTTE INSTITUTE

INSTITUT TUTTE INSTITUTE

Approximate the level set tree, and use ‘Excess of Mass’ to select clusters

INSTITUT TUTTE INSTITUTE

What does this look like in practice?

INSTITUT TUTTE INSTITUTE

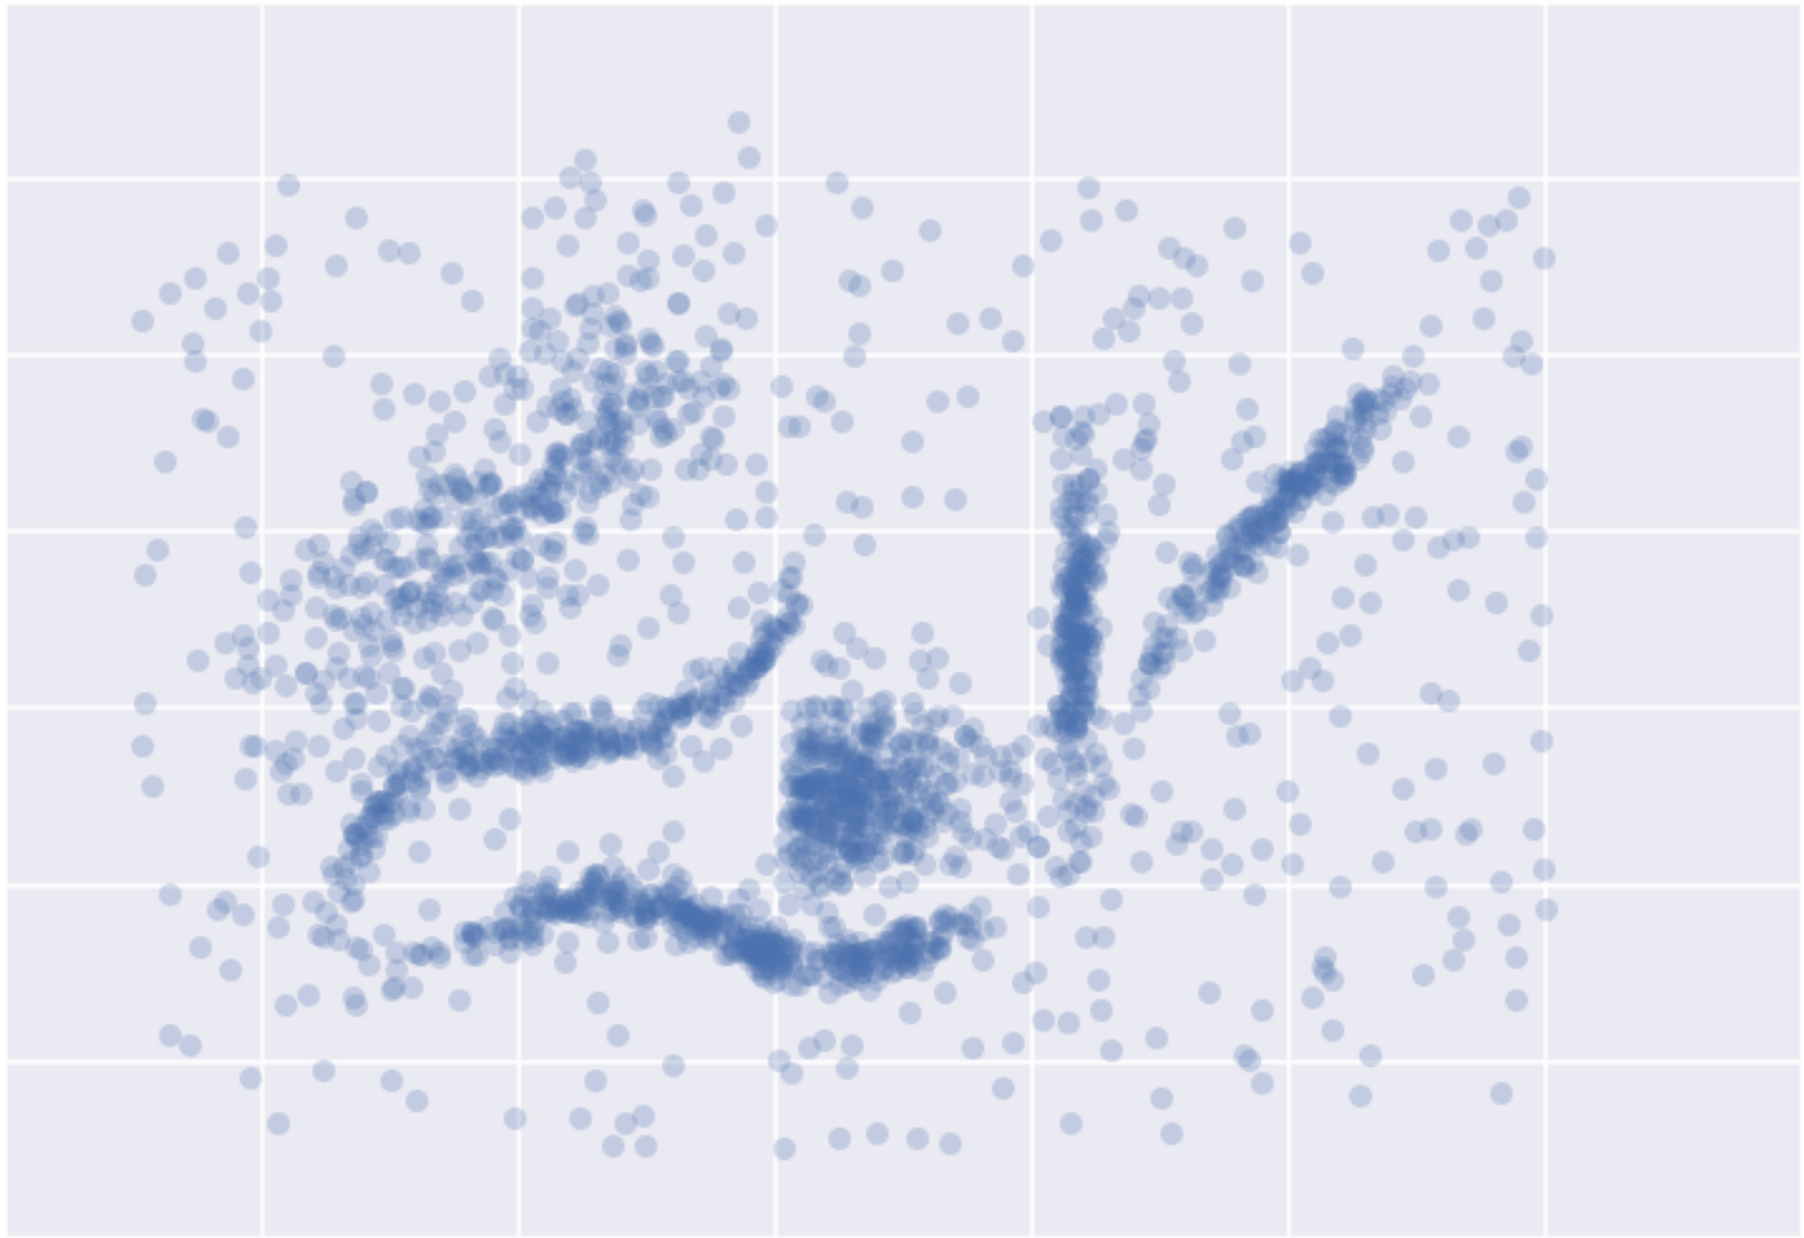

INSTITUT TUTTE INSTITUTE

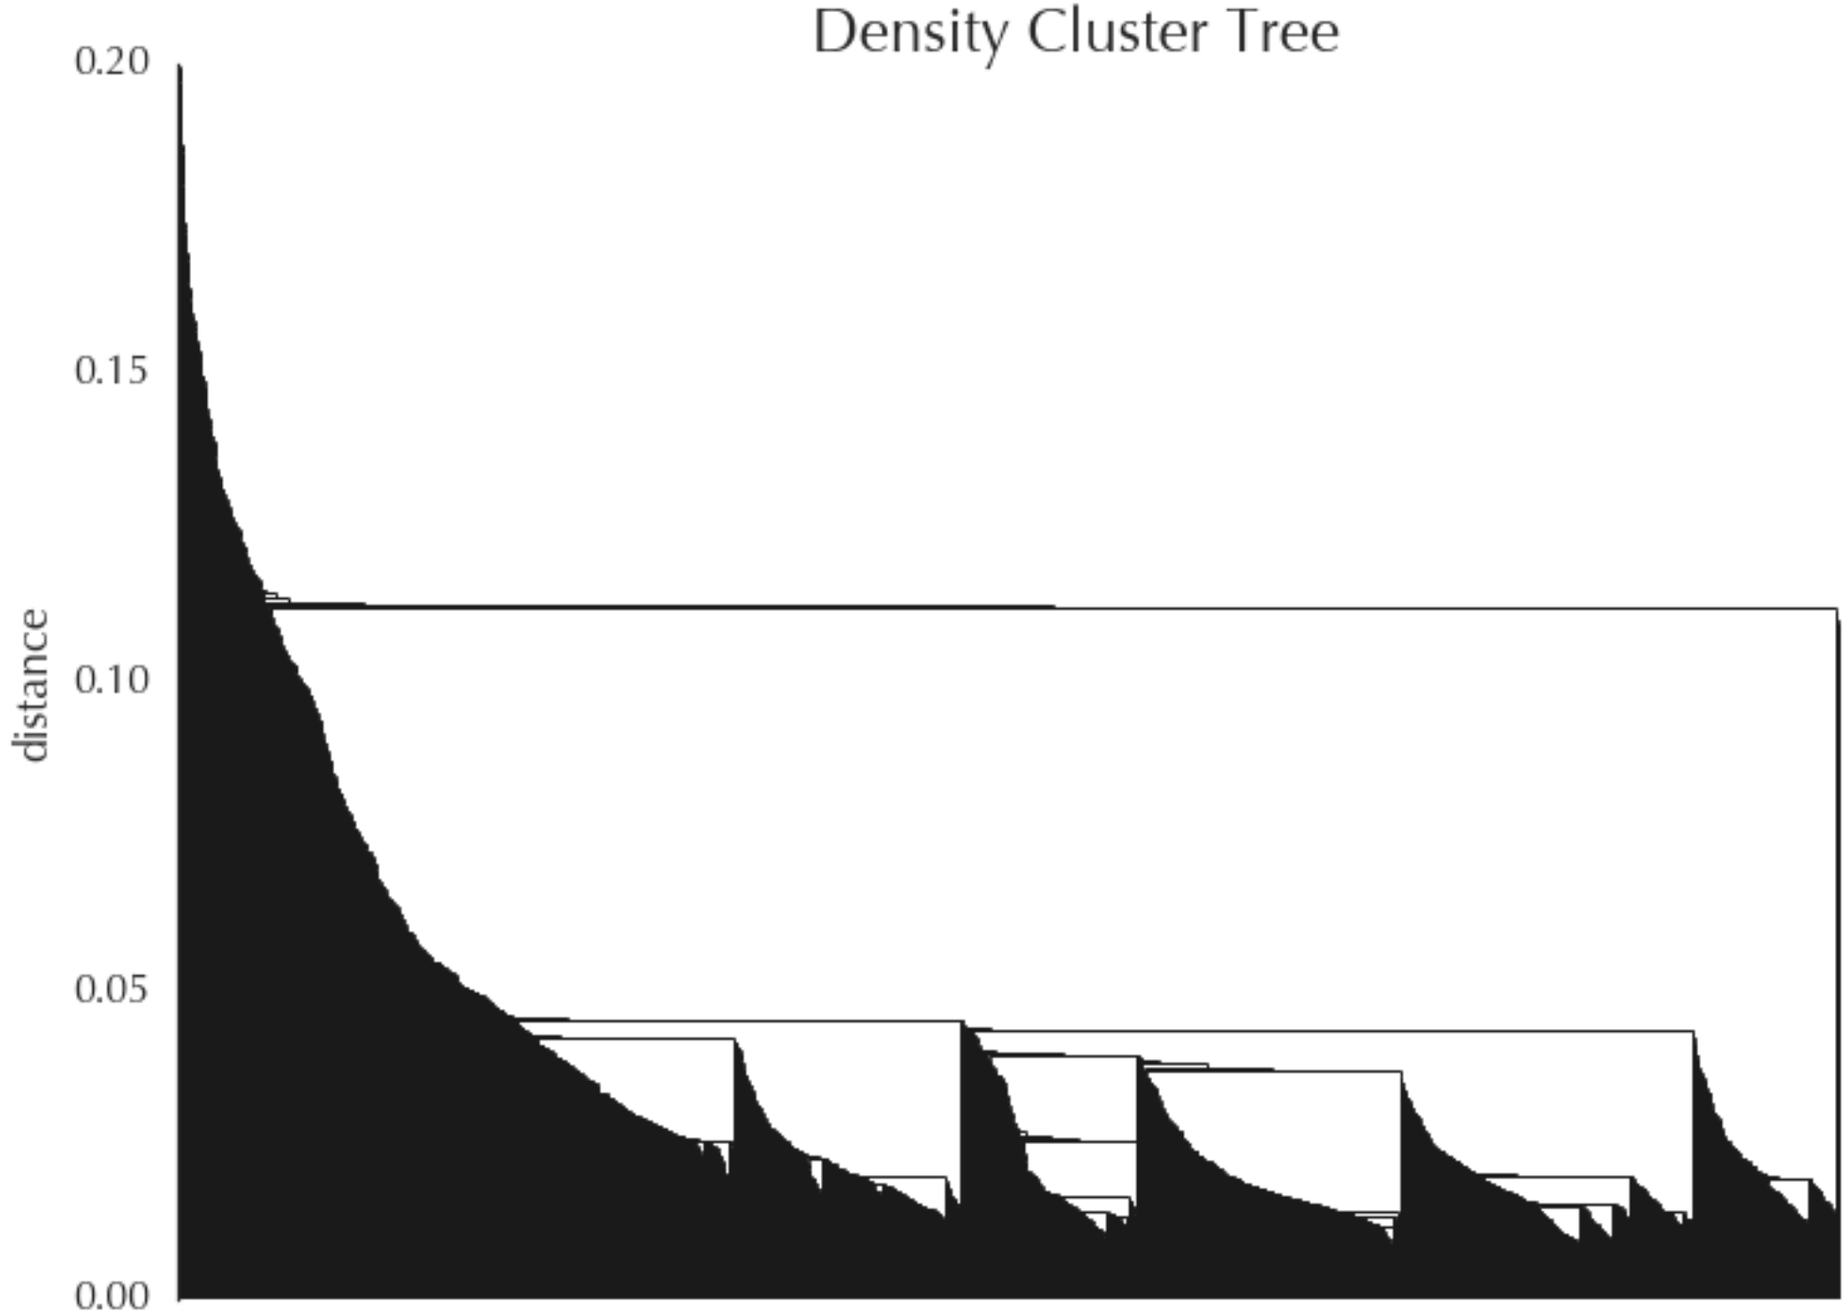

INSTITUT TUTTE INSTITUTE

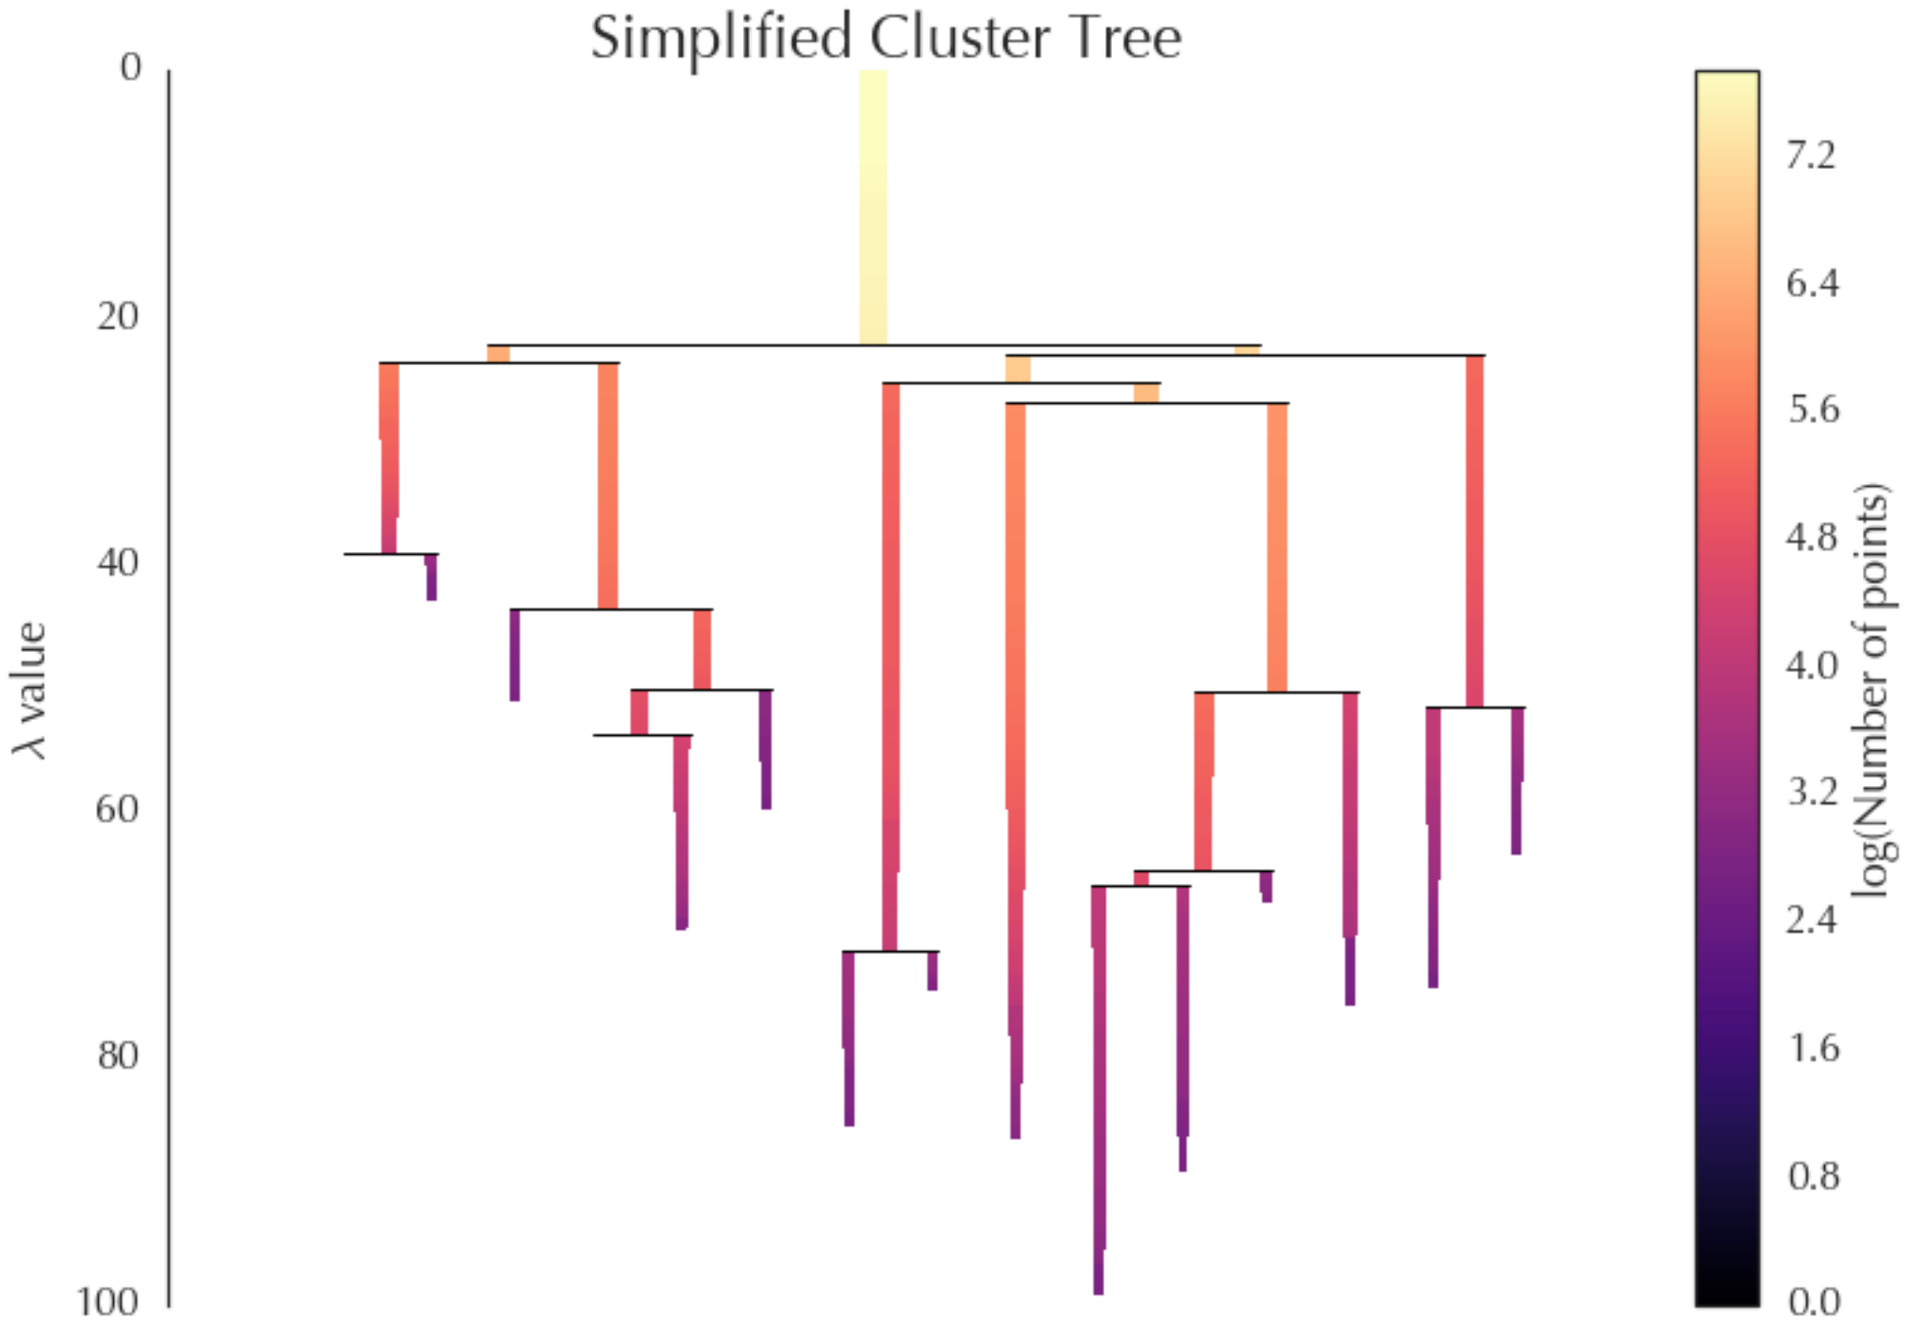

INSTITUT TUTTE INSTITUTE

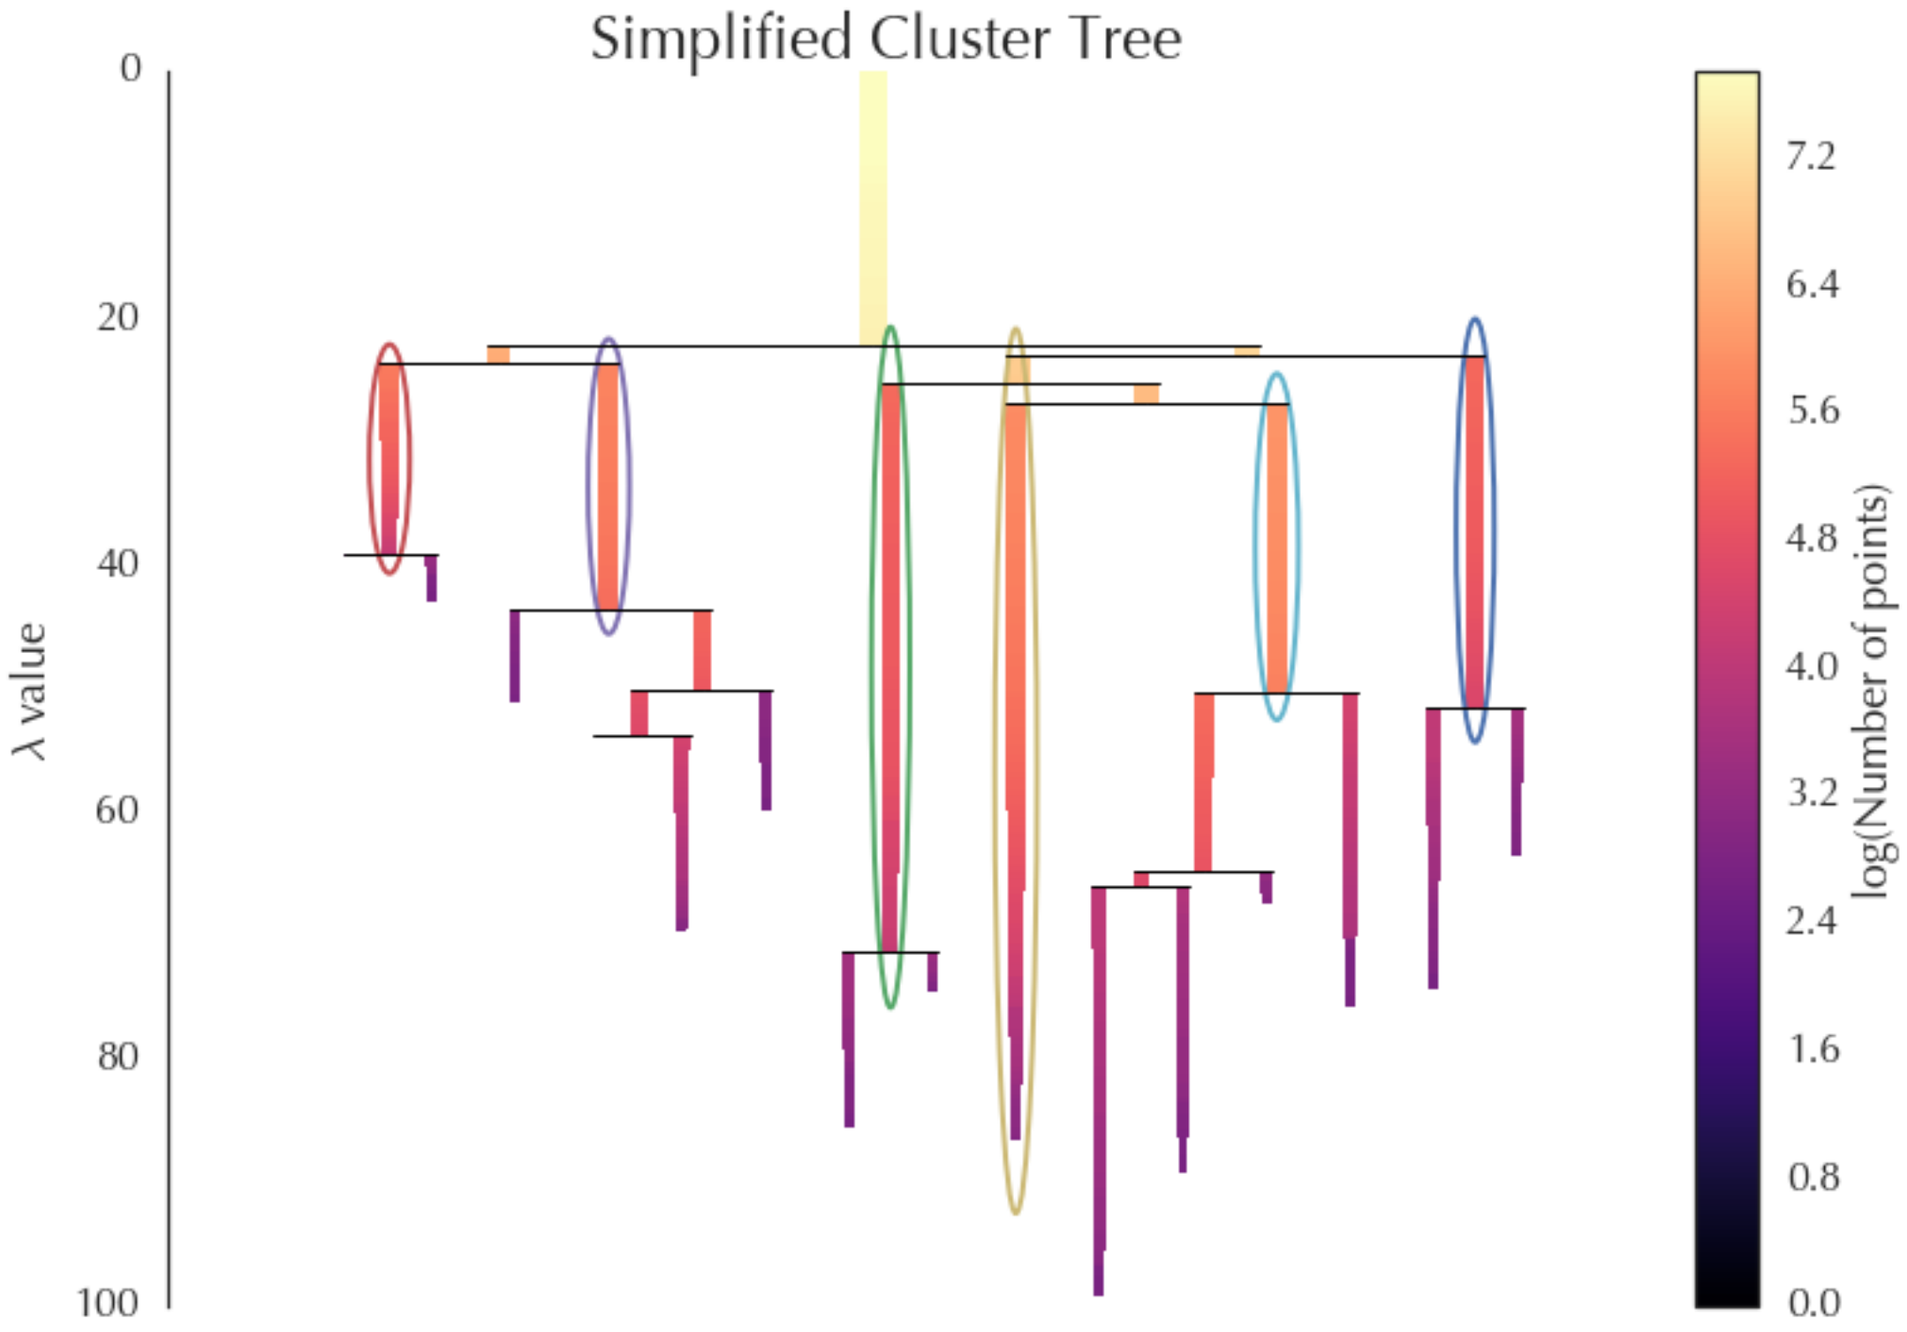

INSTITUT TUTTE INSTITUTE

And the resulting clusters…

INSTITUT TUTTE INSTITUTE

Clusters found by HDBSCAN

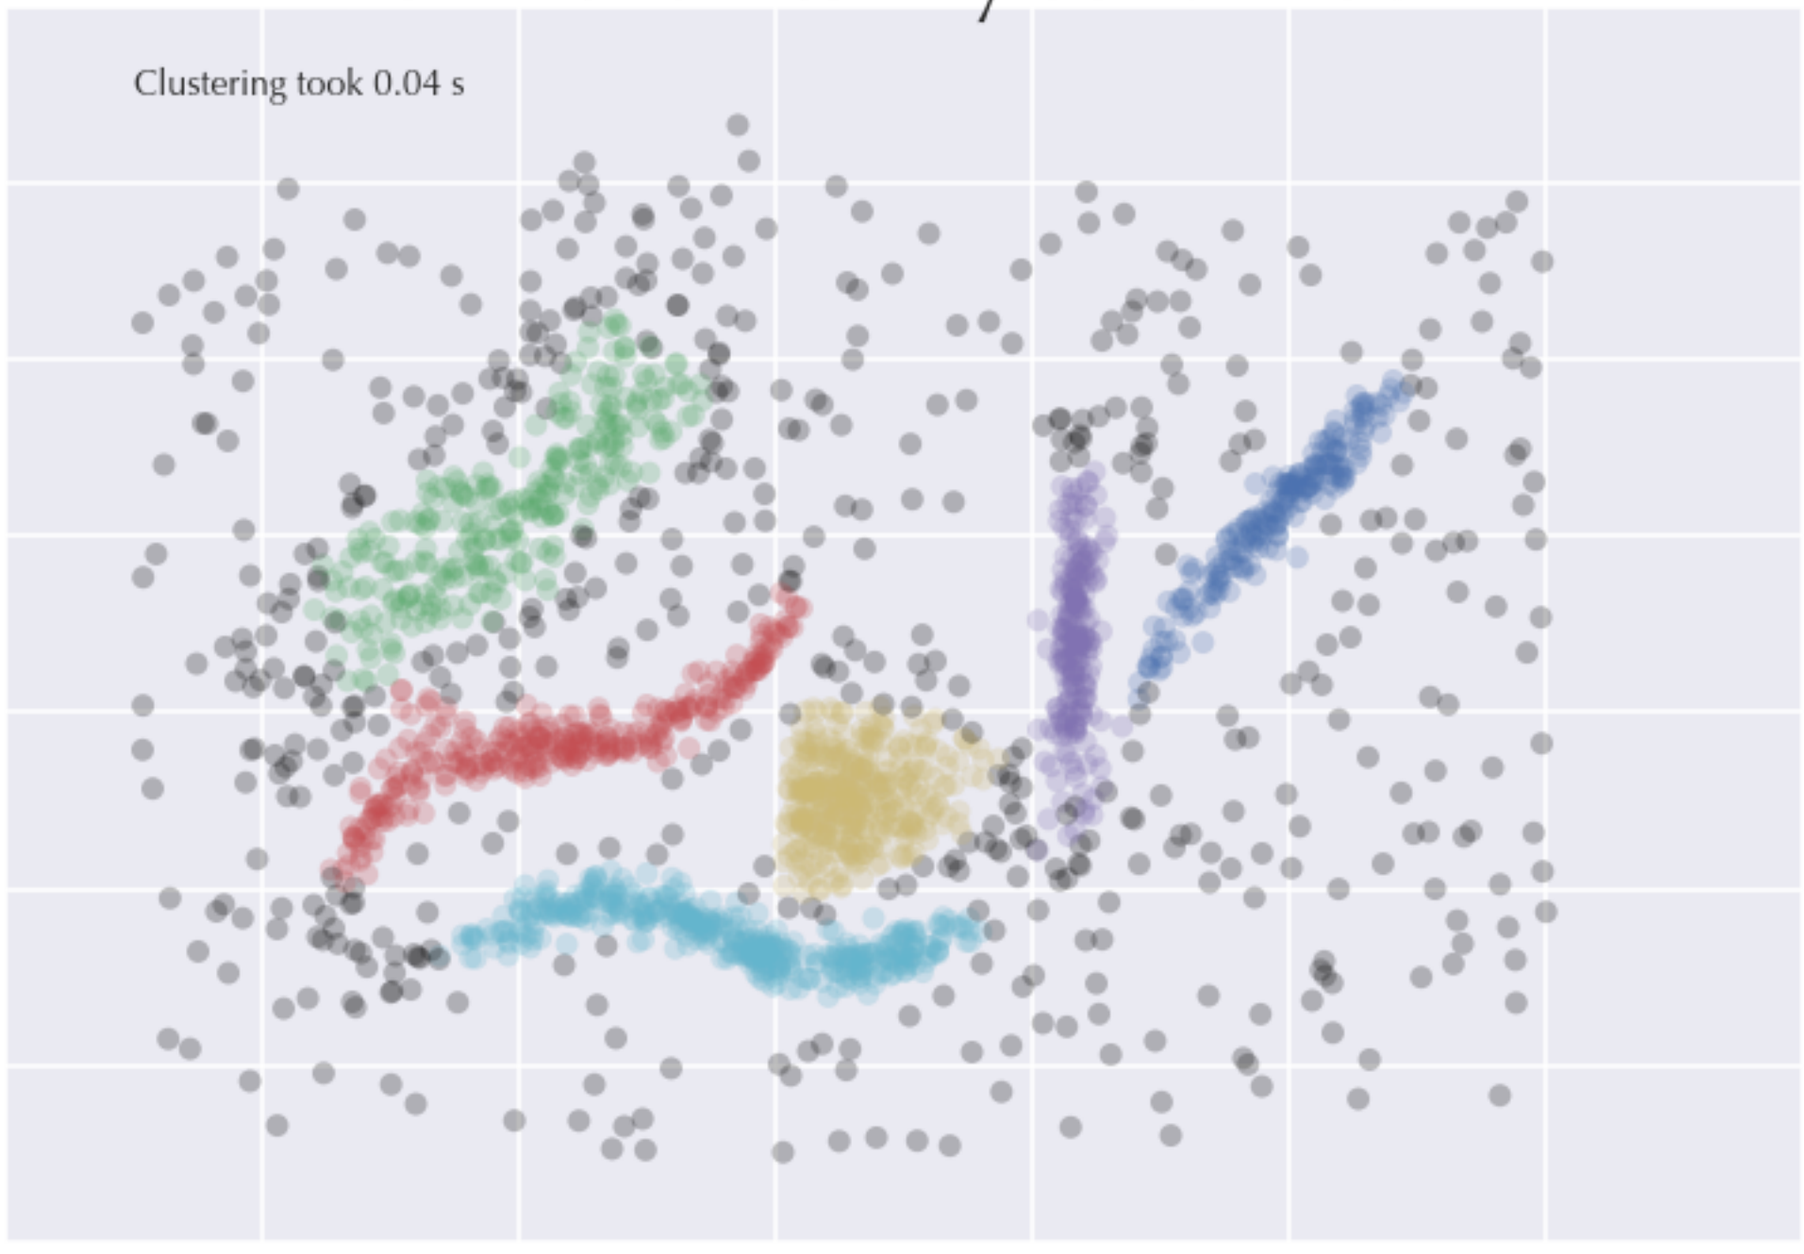

INSTITUT TUTTE INSTITUTE

This is the HDBSCAN* clustering algorithm

Campello, Moulavi, Sander 2013
Campello, Moulavi, Sander, Zimek 2015

INSTITUT TUTTE INSTITUTE

But what about performance?

INSTITUT TUTTE INSTITUTE

We don’t want to run connected components for every possible epsilon!

INSTITUT TUTTE INSTITUTE

Minimum spanning trees!

INSTITUT TUTTE INSTITUTE

The weighted graph is complete!

INSTITUT TUTTE INSTITUTE

We can use spatial indexing to compute fewer distances

INSTITUT TUTTE INSTITUTE

Spatial indexing is great for neighbour queries

INSTITUT TUTTE INSTITUTE

If we use Boruvka's algorithm we can reduce minimum spanning tree computations to repeated neighbour queries

INSTITUT TUTTE INSTITUTE

A modified Dual Tree Boruvka algorithm provides O(n log n) performance

March, Ram, Gray 2010

INSTITUT TUTTE INSTITUTE

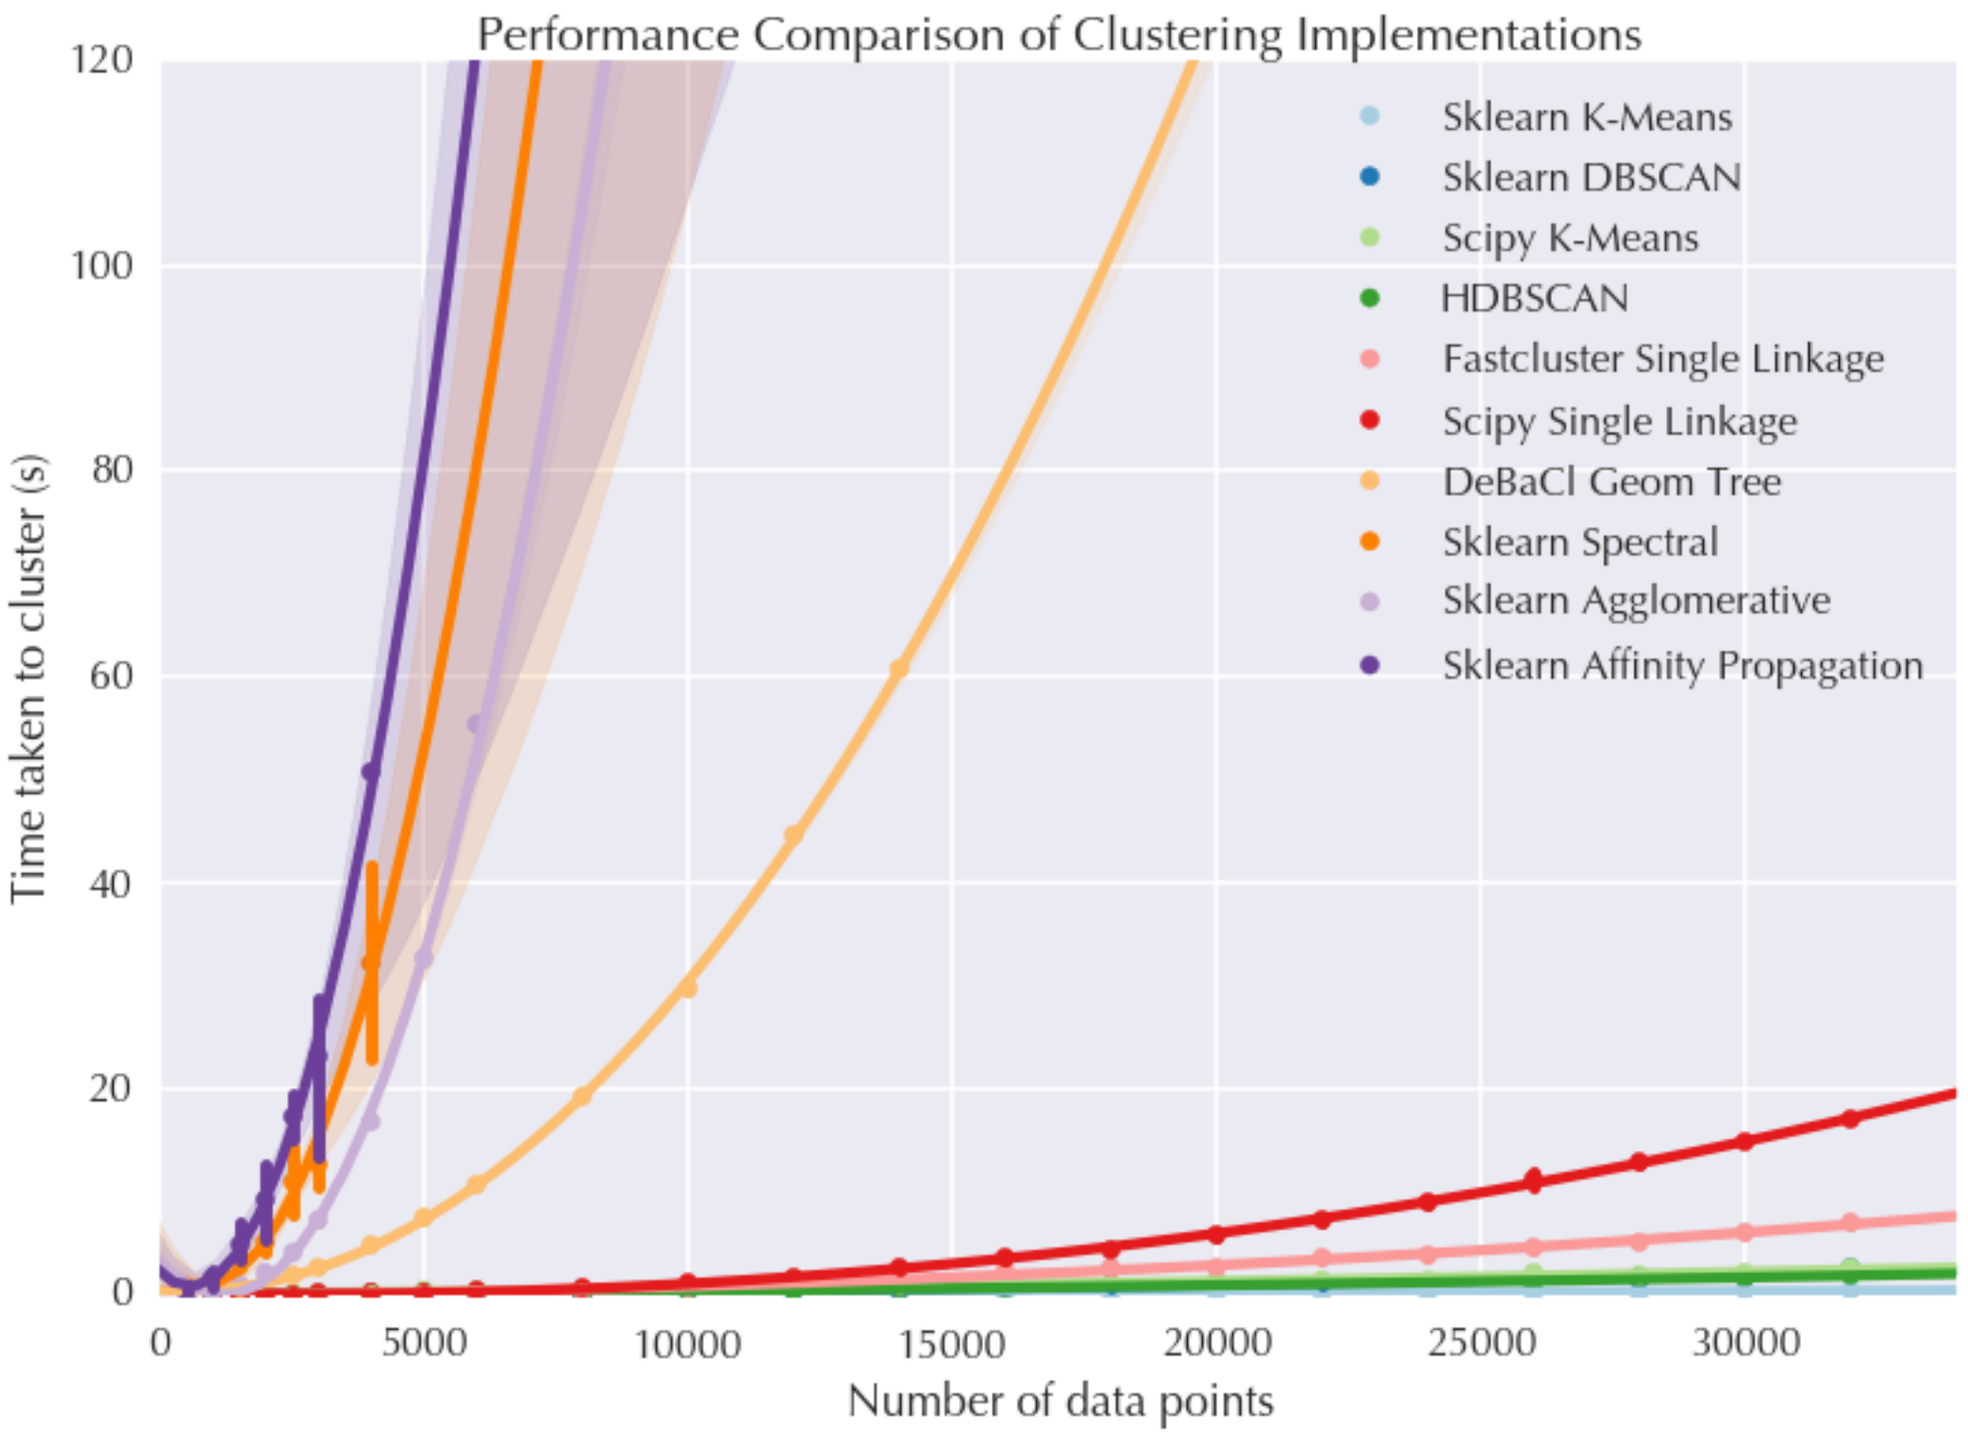

INSTITUT TUTTE INSTITUTE

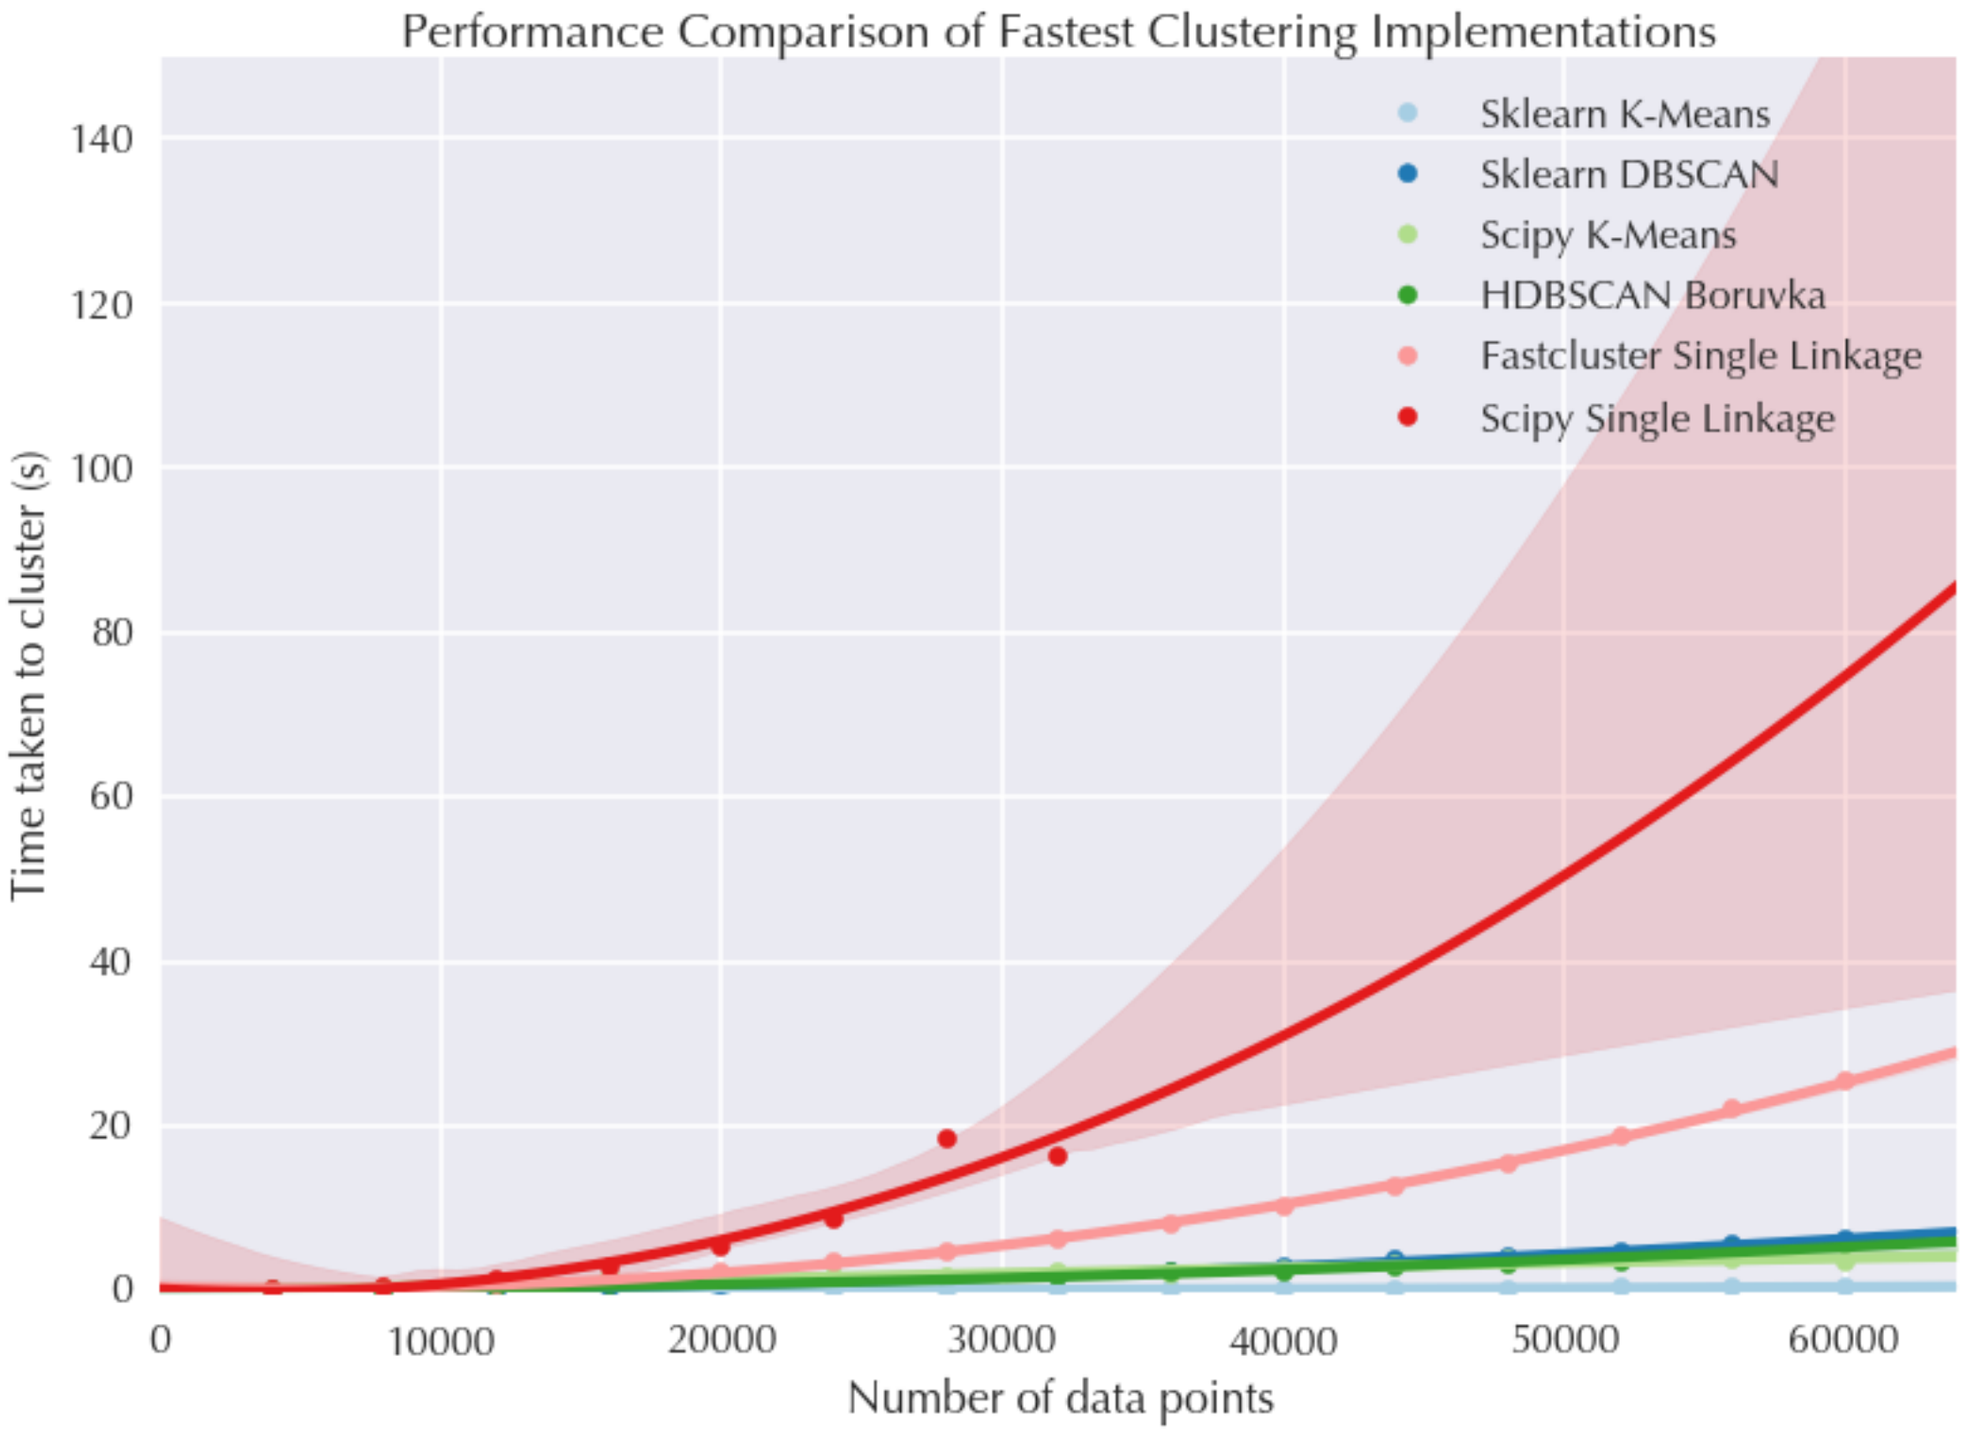

INSTITUT TUTTE INSTITUTE

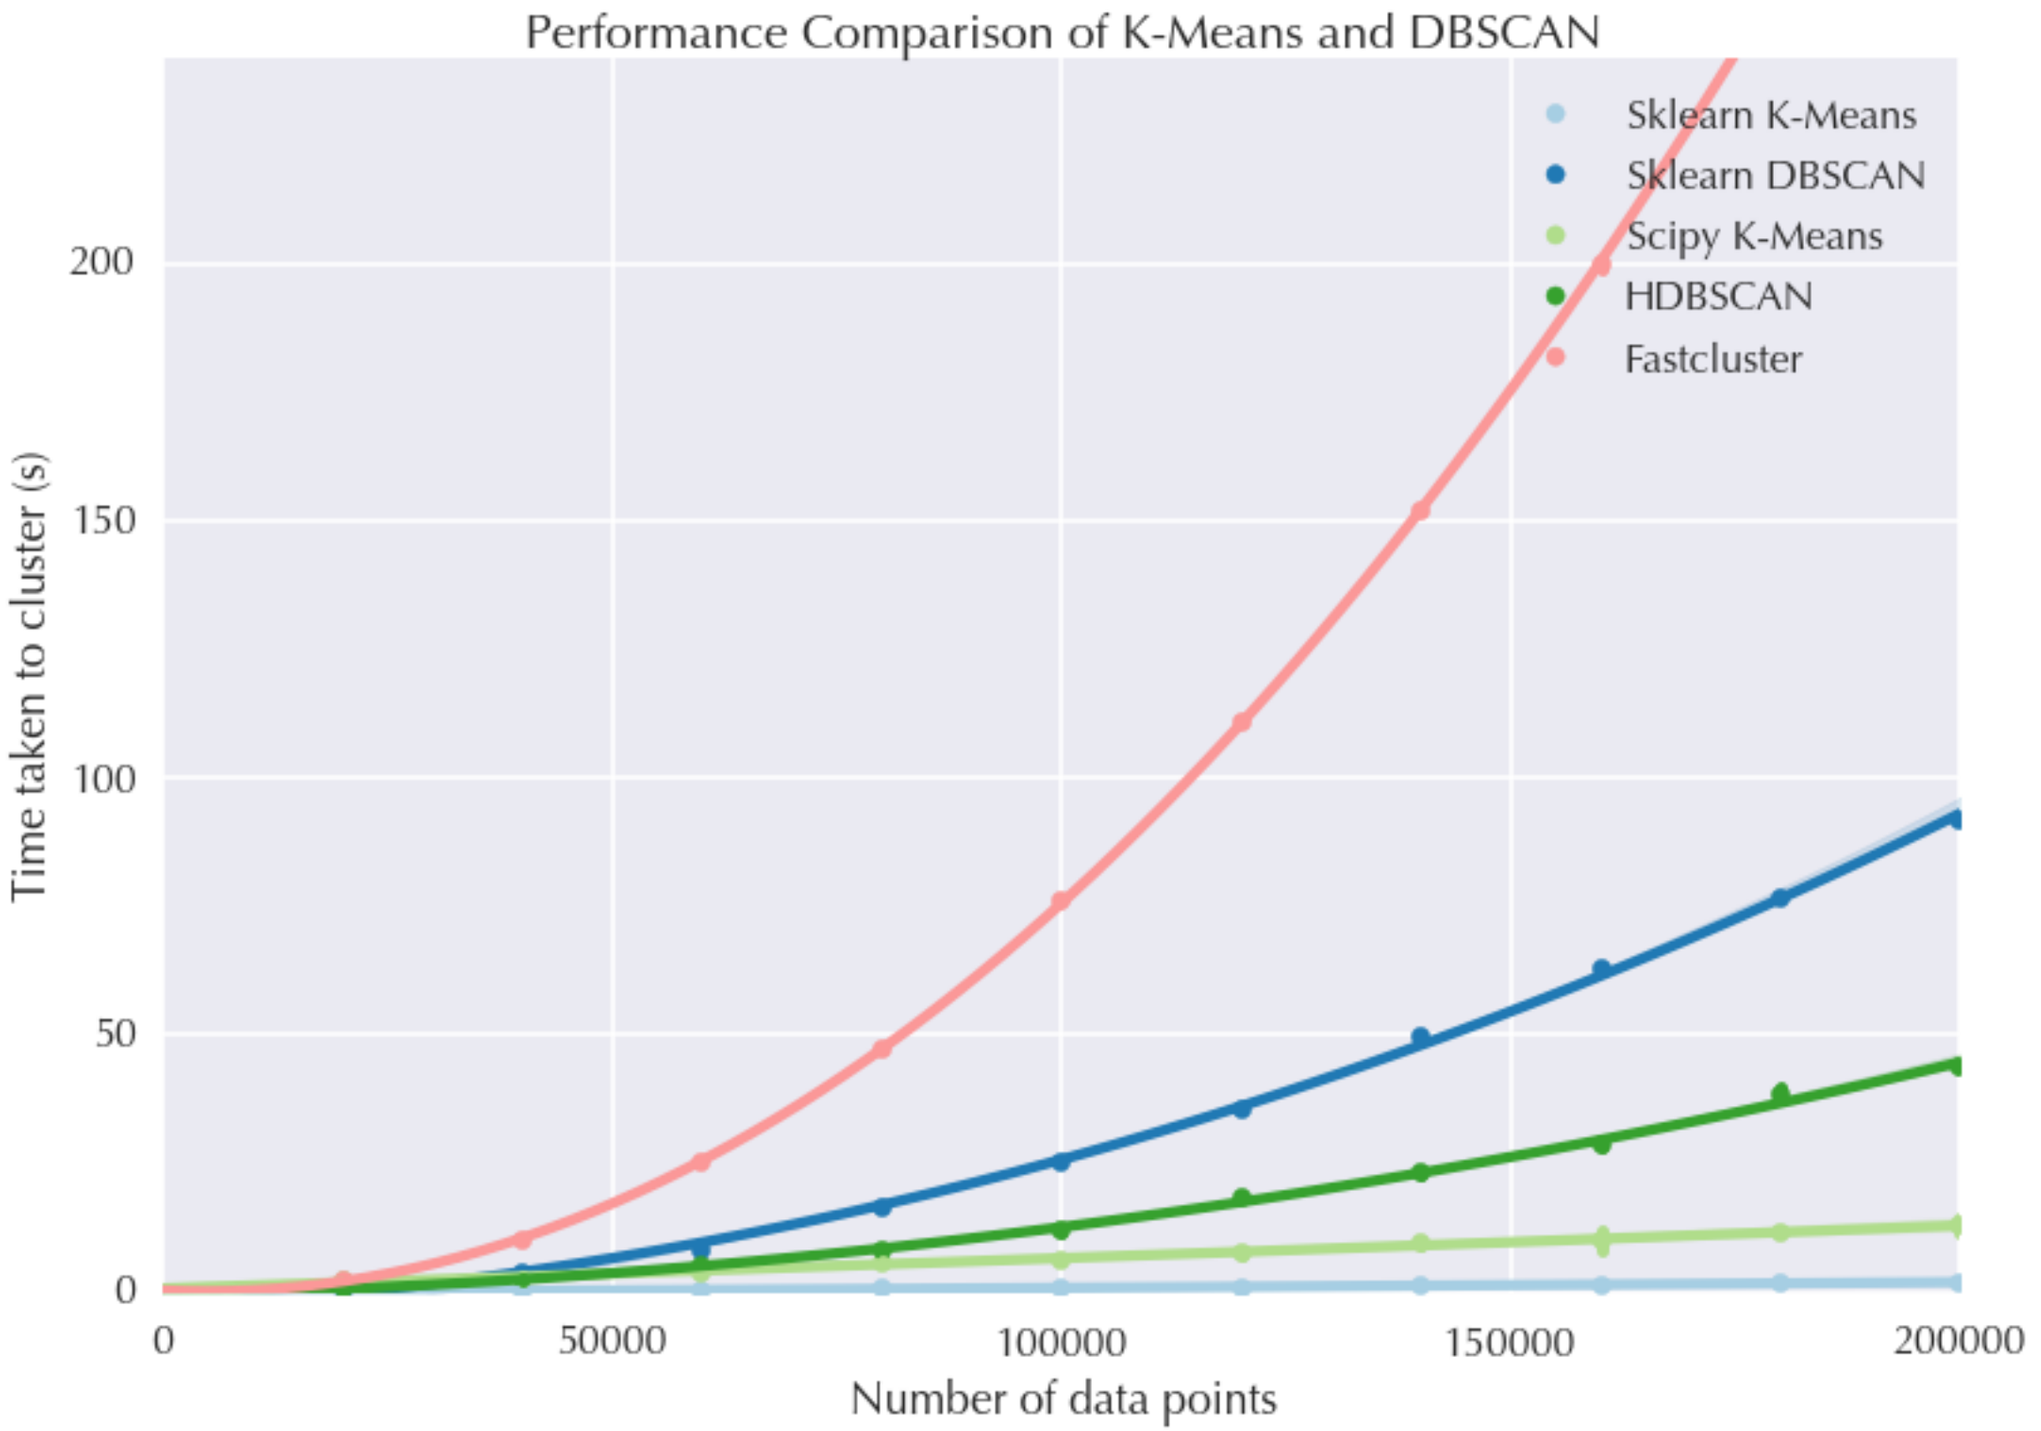

INSTITUT TUTTE INSTITUTE

	Number of data points
Interactive	100,000
Over coffee	500,000
Over lunch	1,000,000
Over night	5,000,000

INSTITUT TUTTE INSTITUTE

Better Clustering?

INSTITUT TUTTE INSTITUTE

We still have to choose a "number of points" value for the density estimate

Can’t we vary that the way we varied epsilon?

INSTITUT TUTTE INSTITUTE

Topology to the rescue!

INSTITUT TUTTE INSTITUTE

Building the initial HDBSCAN* tree can be described in terms of persistent homology

INSTITUT TUTTE INSTITUTE

Multi-dimensional persistent homology allows for persistence over multiple variables

INSTITUT TUTTE INSTITUTE

Unfortunately the result is not a tree

INSTITUT TUTTE INSTITUTE

We can re-interpret HDBSCAN* using sheaves instead of trees

INSTITUT TUTTE INSTITUTE

The resulting algorithm can be generalized to the multi-dimensional case!

INSTITUT TUTTE INSTITUTE

This (will be) Persistent Density Clustering

INSTITUT TUTTE INSTITUTE

More robust
Fewer parameters
Similar performance

INSTITUT TUTTE INSTITUTE

Conclusions

INSTITUT TUTTE INSTITUTE

K-Means shouldn't be your first choice

INSTITUT TUTTE INSTITUTE

K-Means probably shouldn't be your second choice either

INSTITUT TUTTE INSTITUTE

Think hard about what "cluster" means for your application

INSTITUT TUTTE INSTITUTE

You may as well run HDBSCAN* while you're thinking!

https://github.com/scikit-learn-contrib/hdbscan

conda install -c conda-forge hdbscan

pip install hdbscan